

# Fast Choreography of Cross-DevOps Reconfiguration with Ballet

Multi-Site OpenStack Case Study

---

Jolan Philippe, Hélène Couillon, Antoine Omond, Charles Prud'Homme, Issam Raïs

June 4th, 2024 - Journées GDR GPL

STACK team, IMT Atlantique - SeMaFoR project

# Hello world



Jolan Philippe, Postdoc in the SeMaFoR project  
IMT Atlantique, STACK Team (Inria)



Work with Hélène Couillon and Charles Prud'Homme

jolan.philippe@imt-atlantique.fr

Previous work on

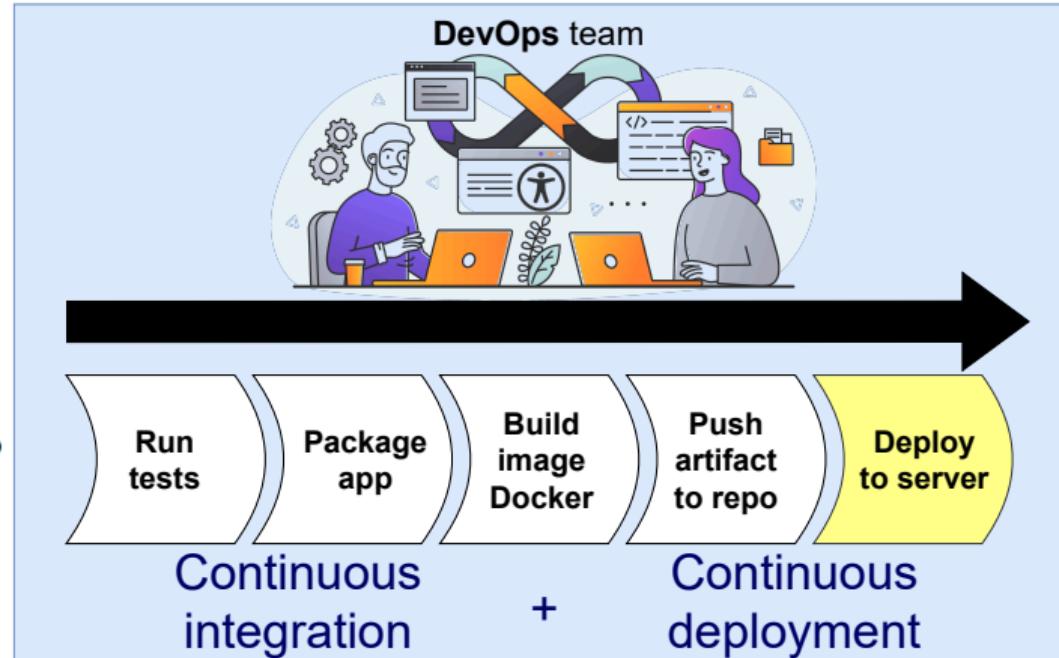
- Geometric approach for clustering, and cryptography
- Algorithmic Skeletons on Lists and Trees
- Study design-choice of distributed transformation engines

More details on: <https://jolanphilippe.github.io/>

# DevOps deployment and reconfiguration



Responsible for development



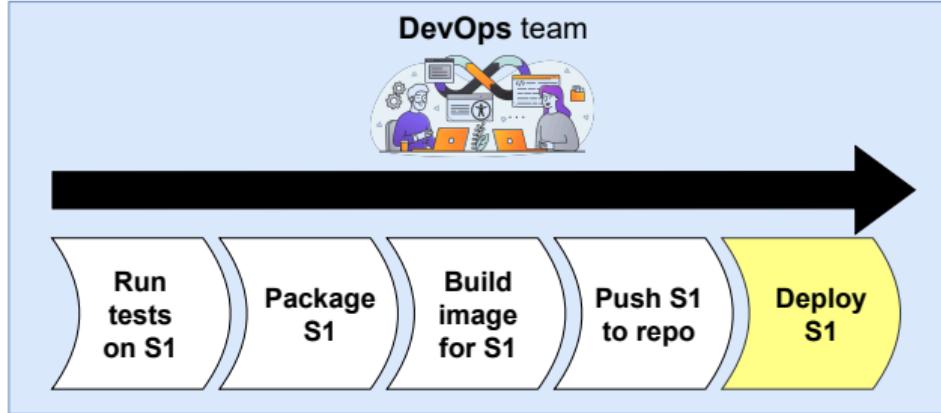
Responsible for operations

⇒ Continuous deployment then **reconfiguration**

# Cross-DevOps reconfiguration



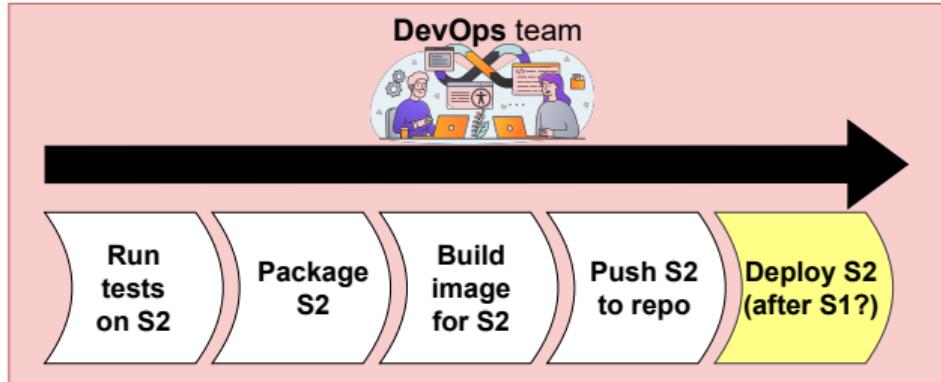
Responsible for development of S1



Responsible for operations on S1



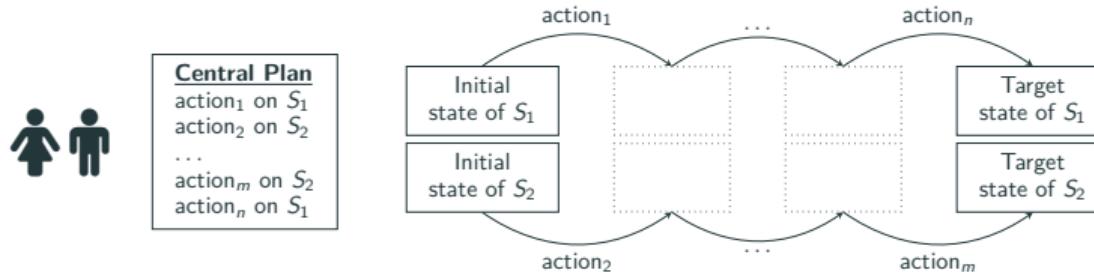
Responsible for development of S2 which uses S1



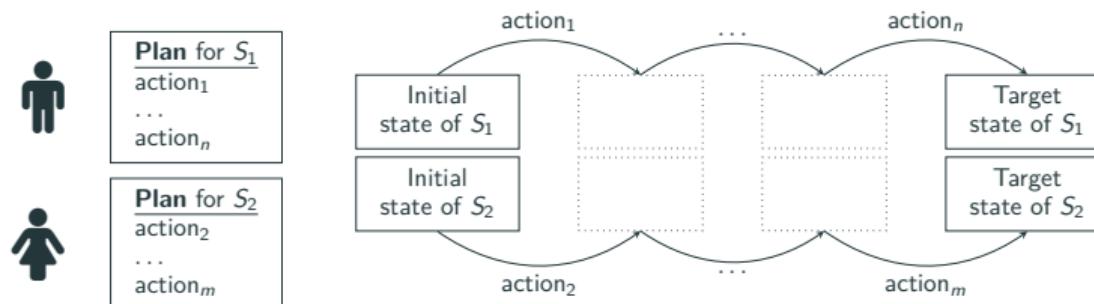
Responsible for operations on S2 (using S1)

# Decentralized reconfiguration

## Centralized approach



## Decentralized approach



## Example: Deploy distributed databases

Galera cluster:



### Database Master (DM) plan

1. Configure the service
2. Bootstrap the database
3. Start the service
4. Expose API

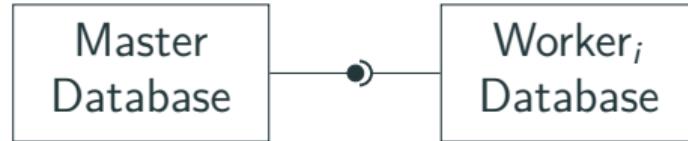
### Database Worker<sub>i</sub> (DW<sub>i</sub>) plan

1. Configure the service
2. Bootstrap the database
3. Start the service
4. Register to master
5. Expose API

- **Component granularity:** DM << DW<sub>i</sub>
- **Lifecycle granularity:** DM(4) << DW<sub>i</sub>(4) (partial order)

## Example: Update distributed databases

Galera cluster:



### Database Master (DM) plan

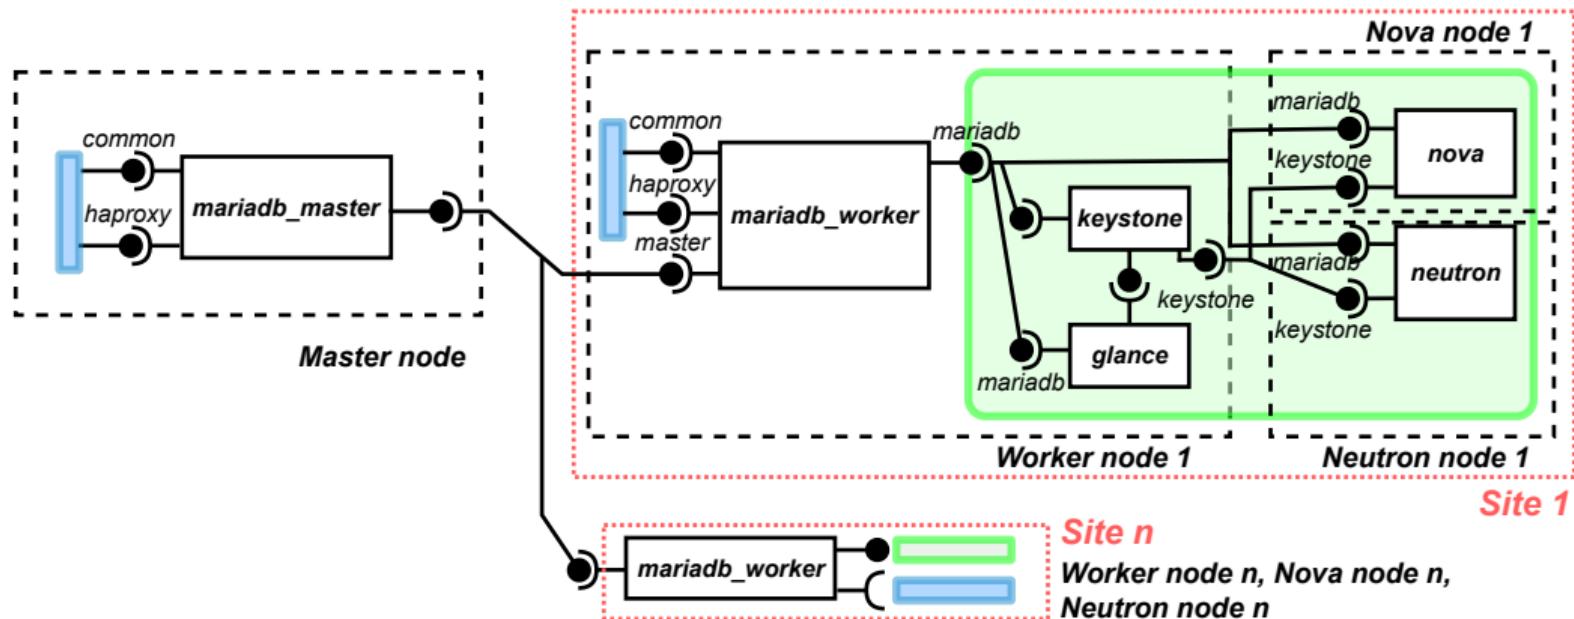
1. Interrupt the service
2. Make the update
3. Start the service
4. Expose API

### Database Worker<sub>i</sub> (DW<sub>i</sub>) plan

1. Interrupt the service
2. Make the update
3. Register to master
4. Start the service
5. Expose API

- **Component granularity?** Destroy DW<sub>i</sub> << Update DM << Deploy DW<sub>i</sub>
- **Lifecycle granularity:** DW<sub>i</sub>(1) << DM(1) & DM(4) << DW<sub>i</sub>(4)

# Case study: Deploy or update OpenStack with Galera cluster of MariaDB



**Figure 1:** Assembly of a multi-site OpenStack with a Galera cluster of distributed MariaDB databases.

# Decentralized reconfiguration

## Challenges

- Human communication between DevOps is error-prone
- Infer reconfiguration actions with a local view
- Execute concurrently reconfiguration actions

## Naive solution

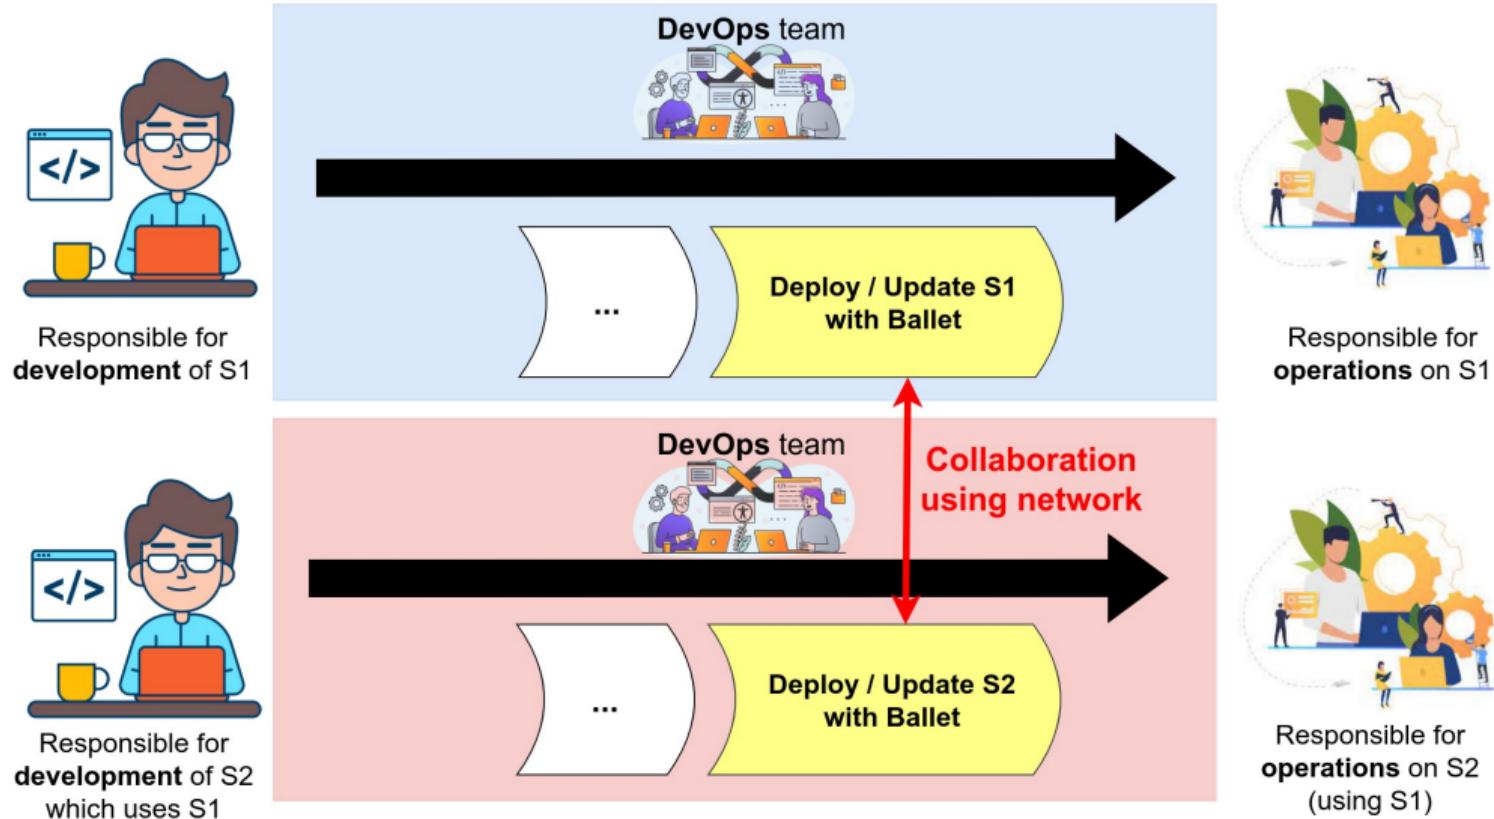
Using a centralized tool (e.g., Terraform, Kubernetes) on top of all DevOps teams is not suitable for scale and fault tolerance reasons.

## Decentralized solution

Make a plan for each DevOps team, and execute them concurrently.

*Muse (Sokolowski et. al.) covers cross-DevOps decentralized reconfiguration with planning, but inefficient because of the fixed life cycles (i.e., on-off mode for resources).*

# Ballet for DevOps



# Ballet's modules overview

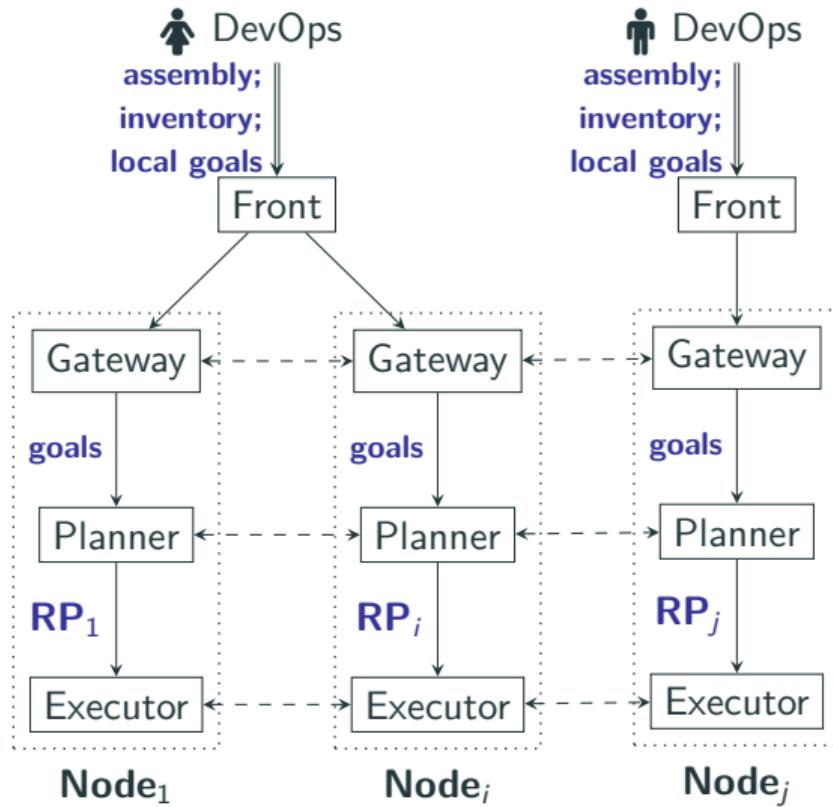


Figure 2: Ballet overview

- Decentralized tool (one instance of Ballet on each node)
- Declarative input
- Reconfiguration with automatic planning and efficient execution

## Gateway

Global knowledge building of reconfiguration goals

## Planner

Decentralized inference of reconfiguration plans (RPs)

## Executor

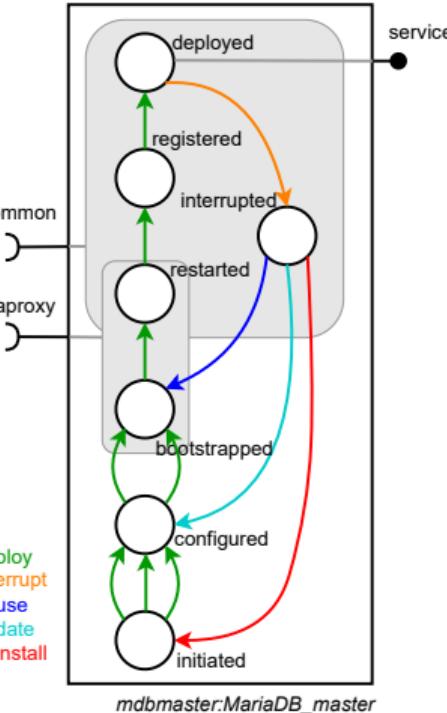
Coordinated execution of RP

# Developers' concern

## Life-cycle and dependencies

Simple language to define component

- **Places:** milestones of the reconfiguration
- **Behaviors:** interface of actions for the DevOps
- **Transitions:** concrete actions between places, associated to behaviors
- **Ports:** Provide (resp. use) information to (resp. from) external components
  - Ports are bounded to places and transitions



**Figure 3:** Visual representation of a component for MariaDB

# Developer's concern

**Listing 1:** Control component MariaDB master in PYTHON

```
1 class MariaDB_Master(Component):
2     def create(self):
3         self.places = [ "initiated", "configured", "bootstrapped", "restarted",
4                         "registered", "deployed", "interrupted"]
5         self.transitions = {
6             "configure0": ("initiated", "configured", "deploy", self.configure0),
7             "configure1": ("initiated", "configured", "deploy", self.configure1),
8             "configure2": ("initiated", "configured", "deploy", self.configure2),
9             ...
10        }
11        self.dependencies = {
12            "service": (DepType.PROVIDE, ["deployed"]),
13            "haproxy": (DepType.USE, ["bootstrapped", "restarted"]),
14            ...
15        }
16        self.initial_place = 'initiated'
17        self.running_place = 'deployed'
18
19    def configure0(self):
20        # concrete actions
```

# DevOps' concern

## Target assembly (YAML)

- A list of components to appear
- How components are connected

## Reconfiguration goals

Declarative language for defining reconfiguration goals

- **Behavior goal:** Specify a behavior that must be executed
- **Port goal:** Specify a port status (active, inactive)
- **State goal:** Specify a component state (specific, running, initial)

**Listing 2:** Language to define reconfiguration goals for DevOps usage

```
<goals> ::= behaviors: <bhvr_list>
           ports: <port_list>
           components: <comp_list>
<bhvr_list> ::= ...
<bhvr_item> ::= - forall: <bhvr_name>
                  | - component: <comp_name>
                    behavior: <bhvr_name>
<port_list> ::= ...
<port_item> ::= - forall: <port_status>
                  | - component: <comp_name>
                    port: <port_name>
                    status: <port_status>
<comp_list> ::= ...
<comp_item> ::= - forall: <comp_status>
                  | - component: <comp_name>
                    status: <comp_status>
```

# Example - OpenStack assembly, update MariaDB master

## Target assembly:

```
components:  
  mdbmaster:  
    type: 'MariaDB_master'  
  facts:  
    type: 'Facts'  
  common:  
    type: 'Common'  
  haproxy:  
    type: 'HAProxy'  
  ...  
  
connections:  
  - facts, service, common, facts-service  
  - facts, service, haproxy, facts-service  
  - ...  
  ${{ each site in range(0,sites) }}:  
    - mdbmaster, service, mdbworker${{ site }}, master-service
```

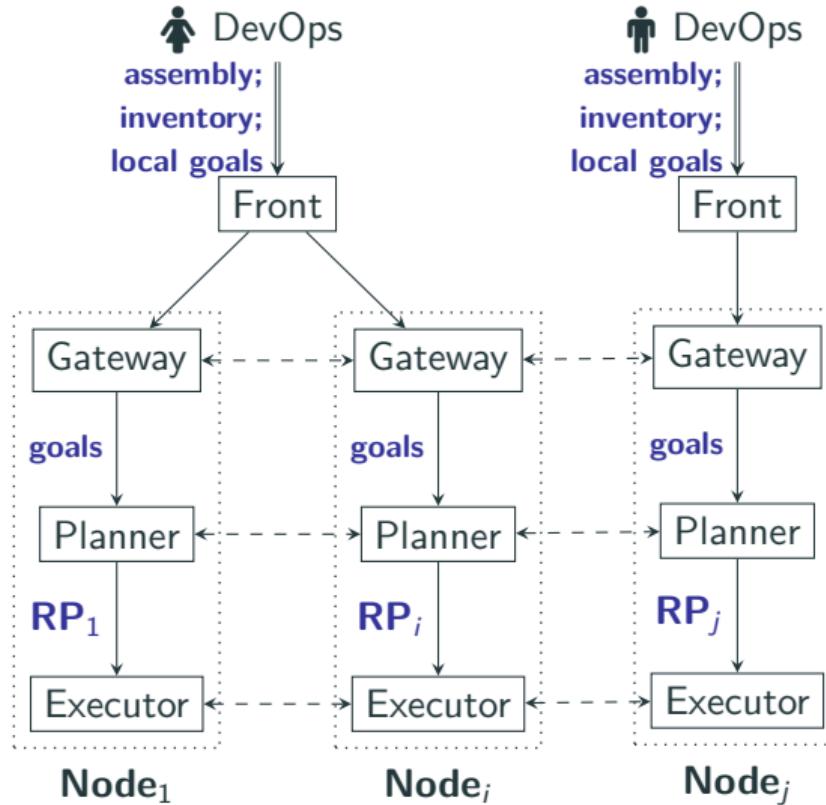
## Goals for deployment

```
components:  
  - forall: running
```

## Goals for update database

```
behaviors:  
  - component: mariadb_master  
    behavior: update  
components:  
  - forall: initial
```

# Reconfiguring process with Ballet



## Outline, for clarity reasons

1. Skip gateway since there is no scientific challenge
2. Start with the execution, to understand planner's challenges
3. Followed by the planning

# Decentralized execution

## Concerto-D

- Based on Concerto model
  - As efficient as Concerto
  - Decentralized version
  - A model that includes communications inter-nodes
- One controller per node

## Reconfiguration program

Within such a model, plans can

1. Create assemblies of components (software system)
2. Make this assembly evolve at runtime
3. Interact with the life cycle of components

## Concerto-D language

The used language propose instructions for:

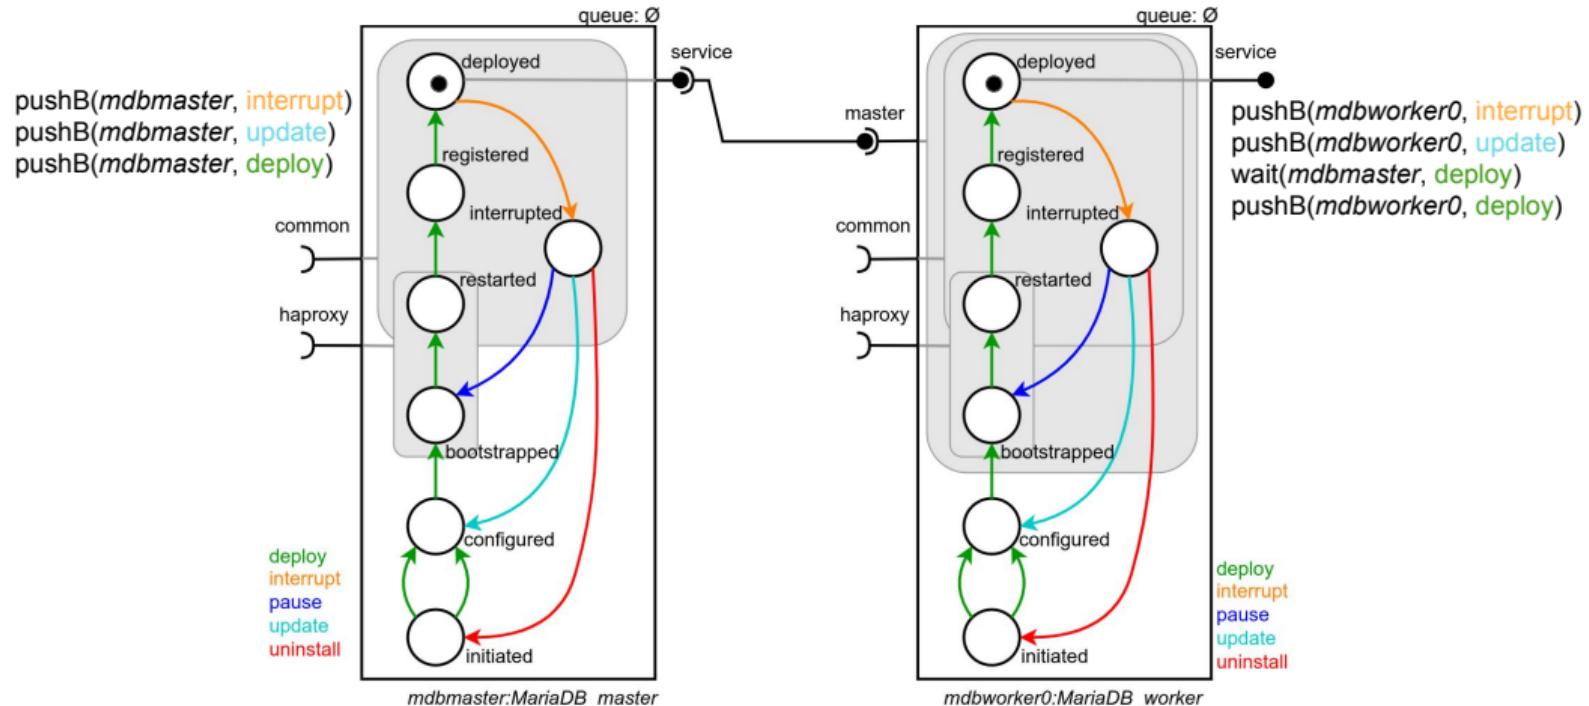
Add/remove a component instance to the current assembly

Connect/disconnect two component instances with compatible ports

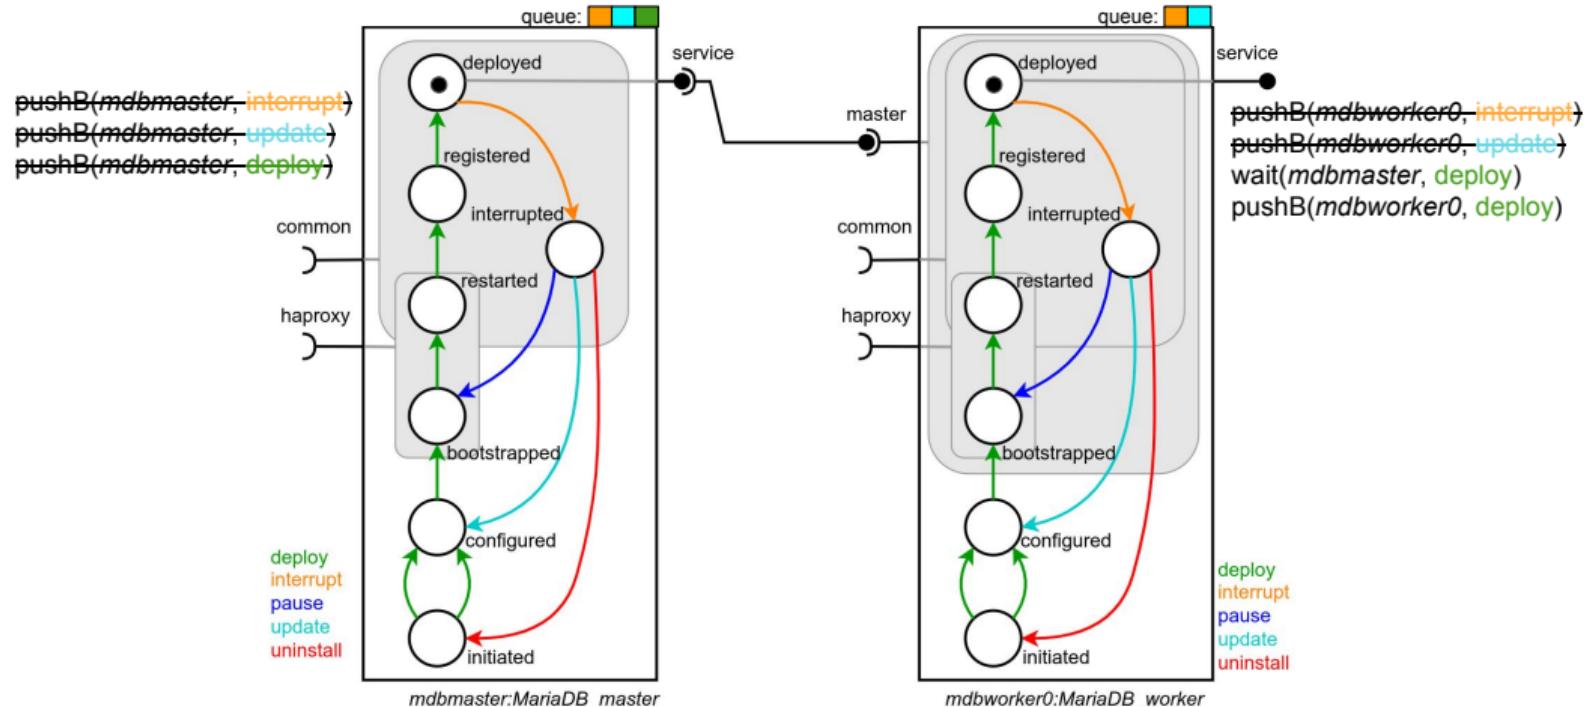
Push behavior to the behavior queue on a component instance

Wait for a given component instance to execute a behavior

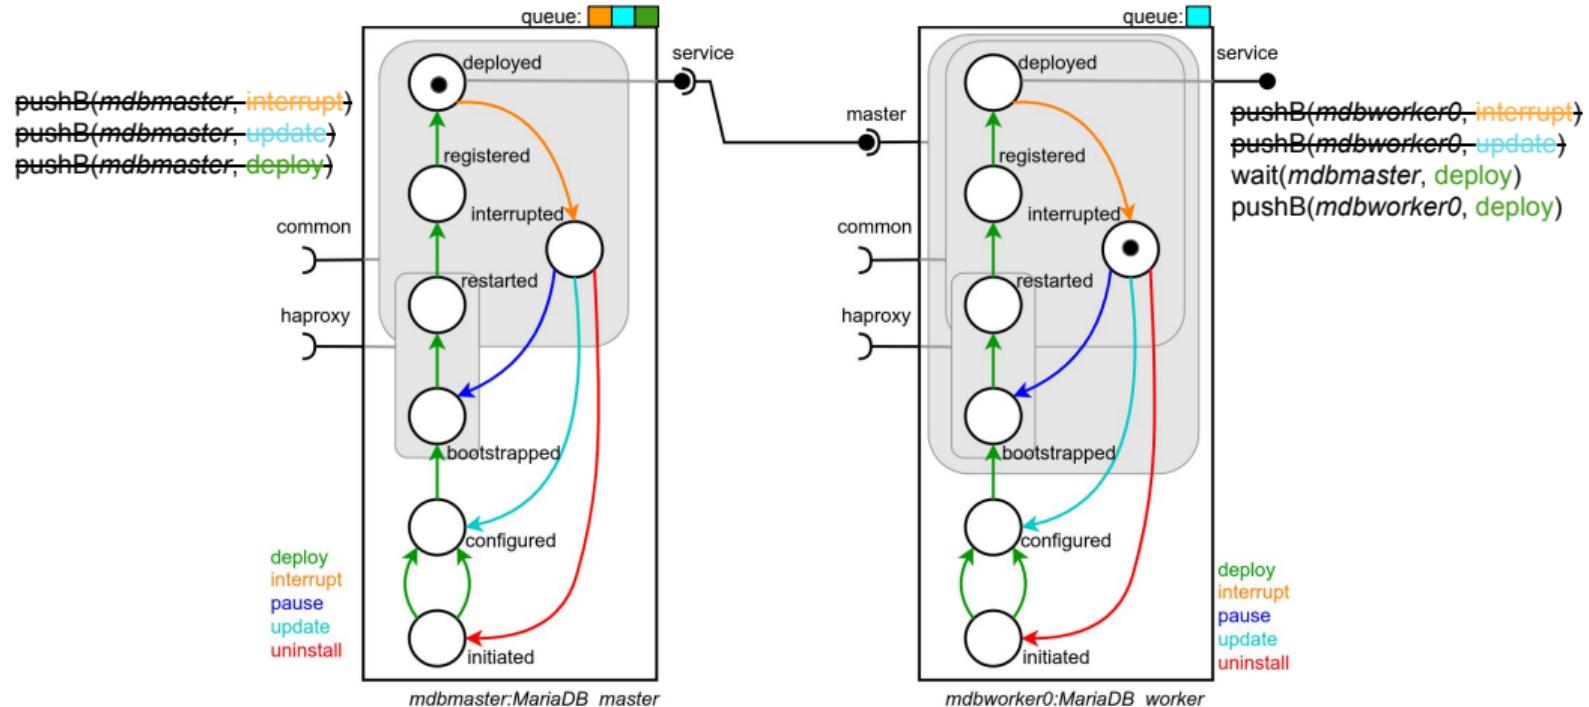
# Program example



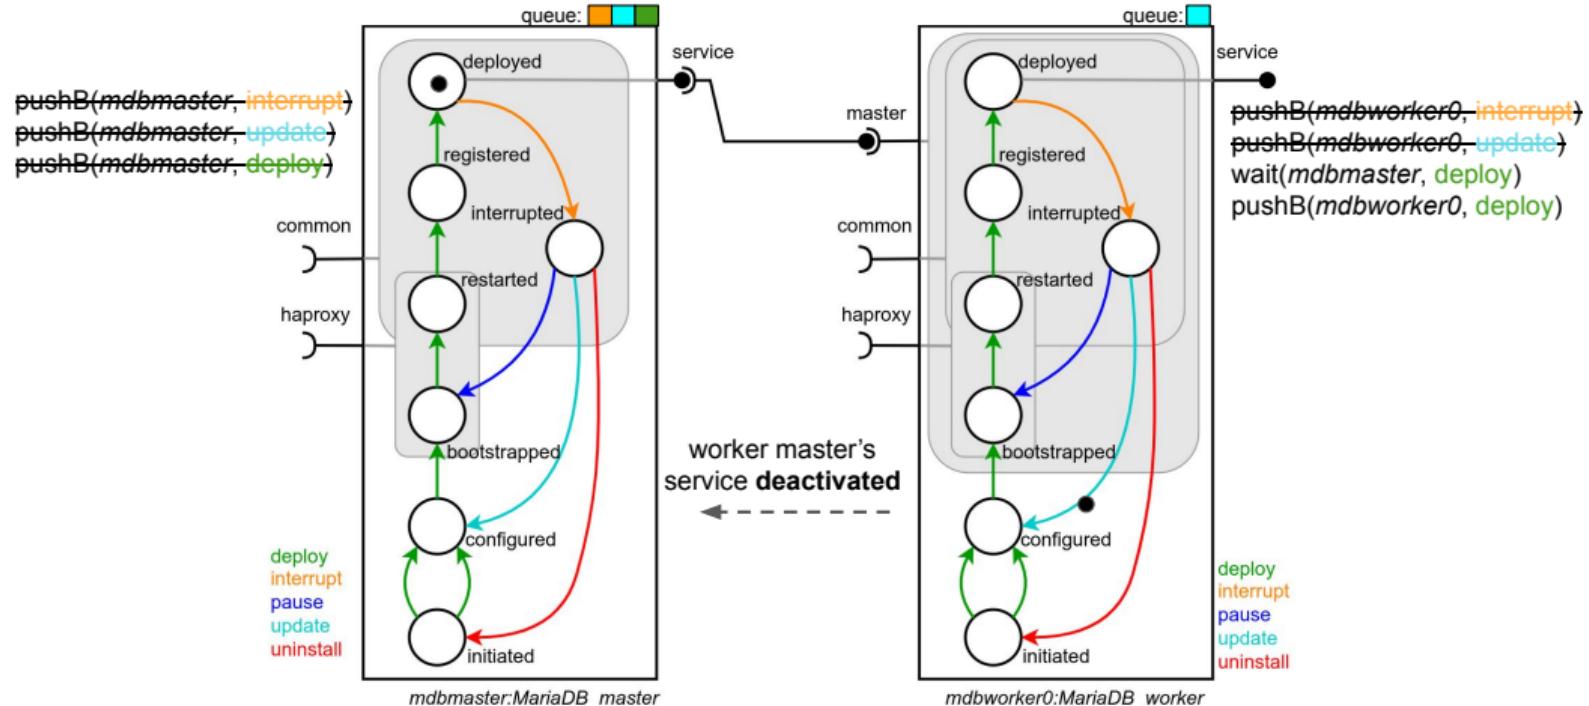
# Program example



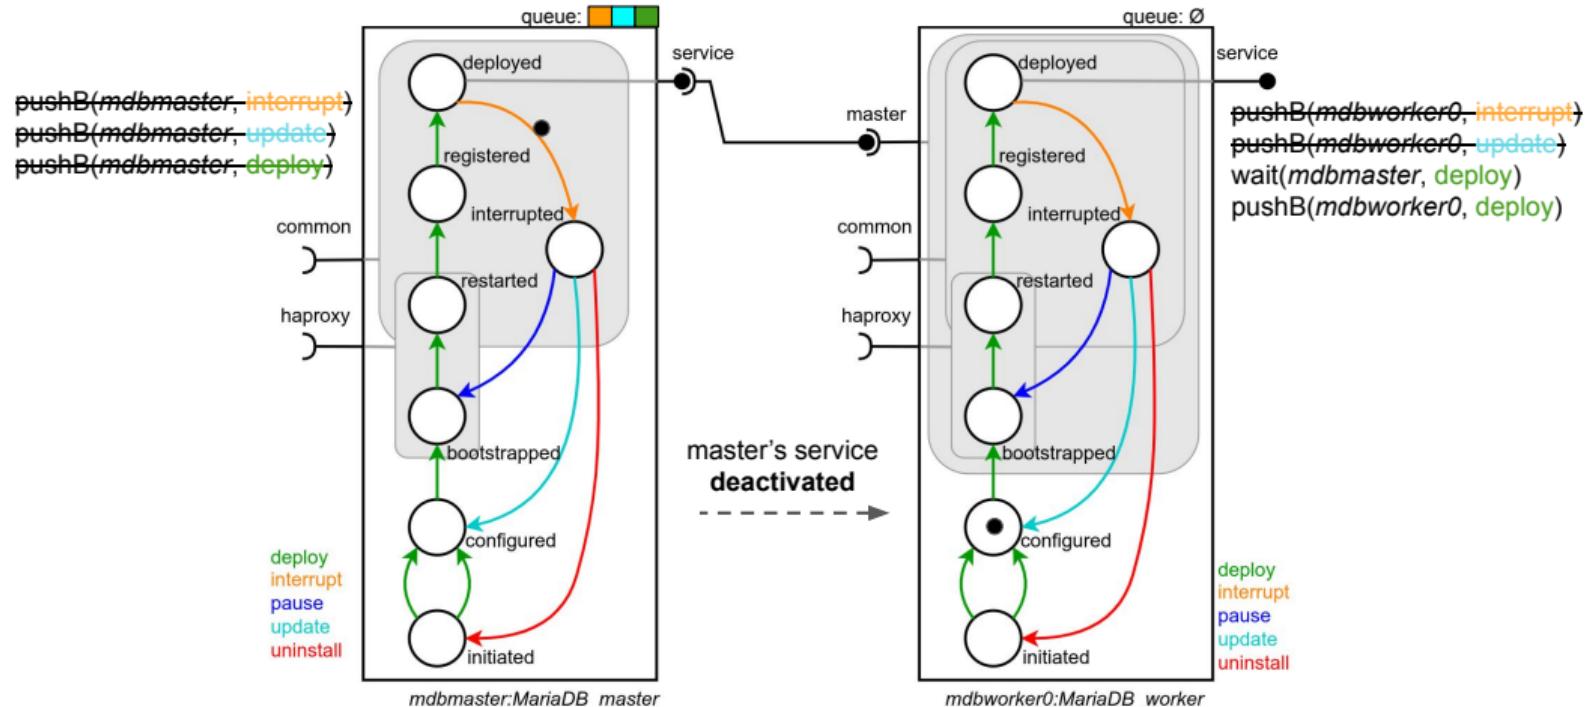
# Program example



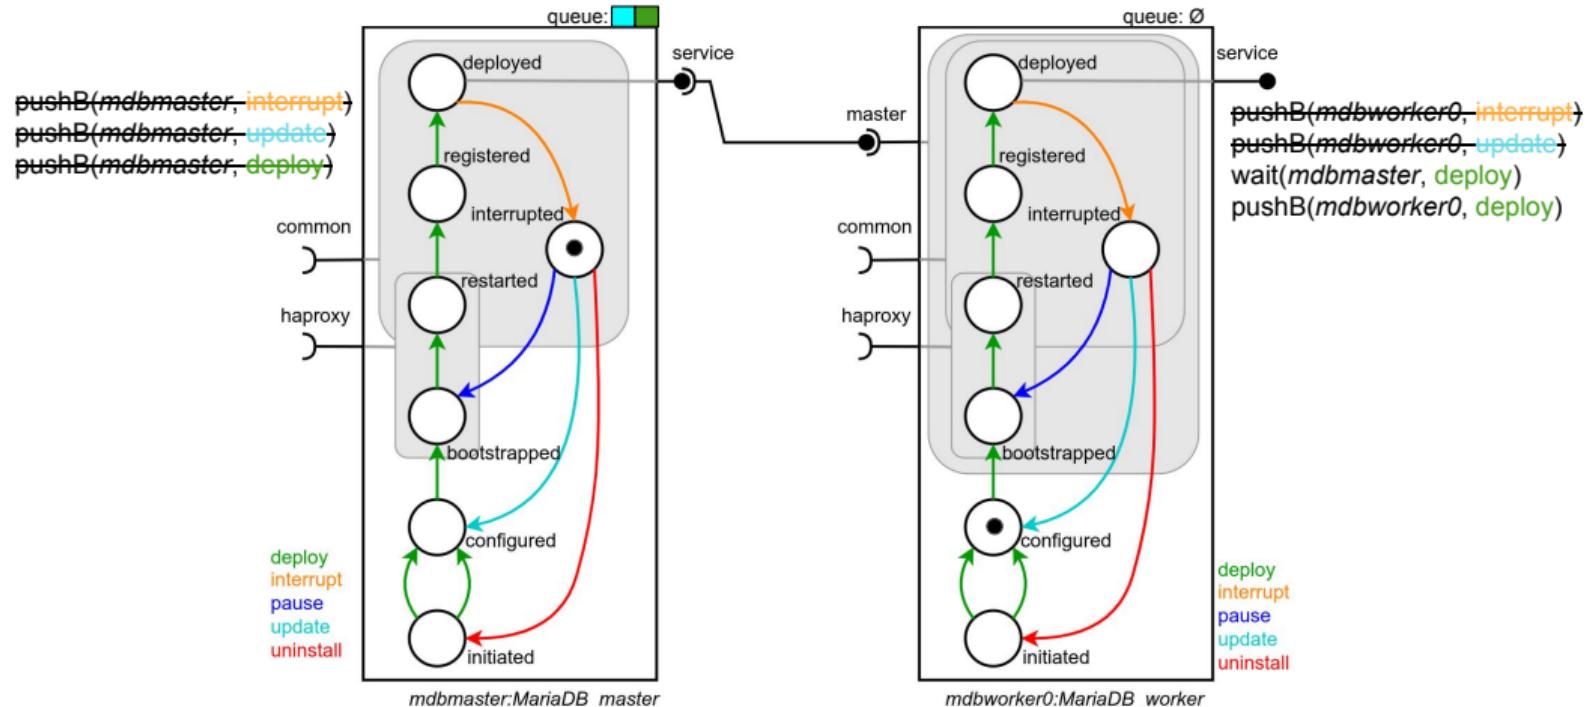
# Program example



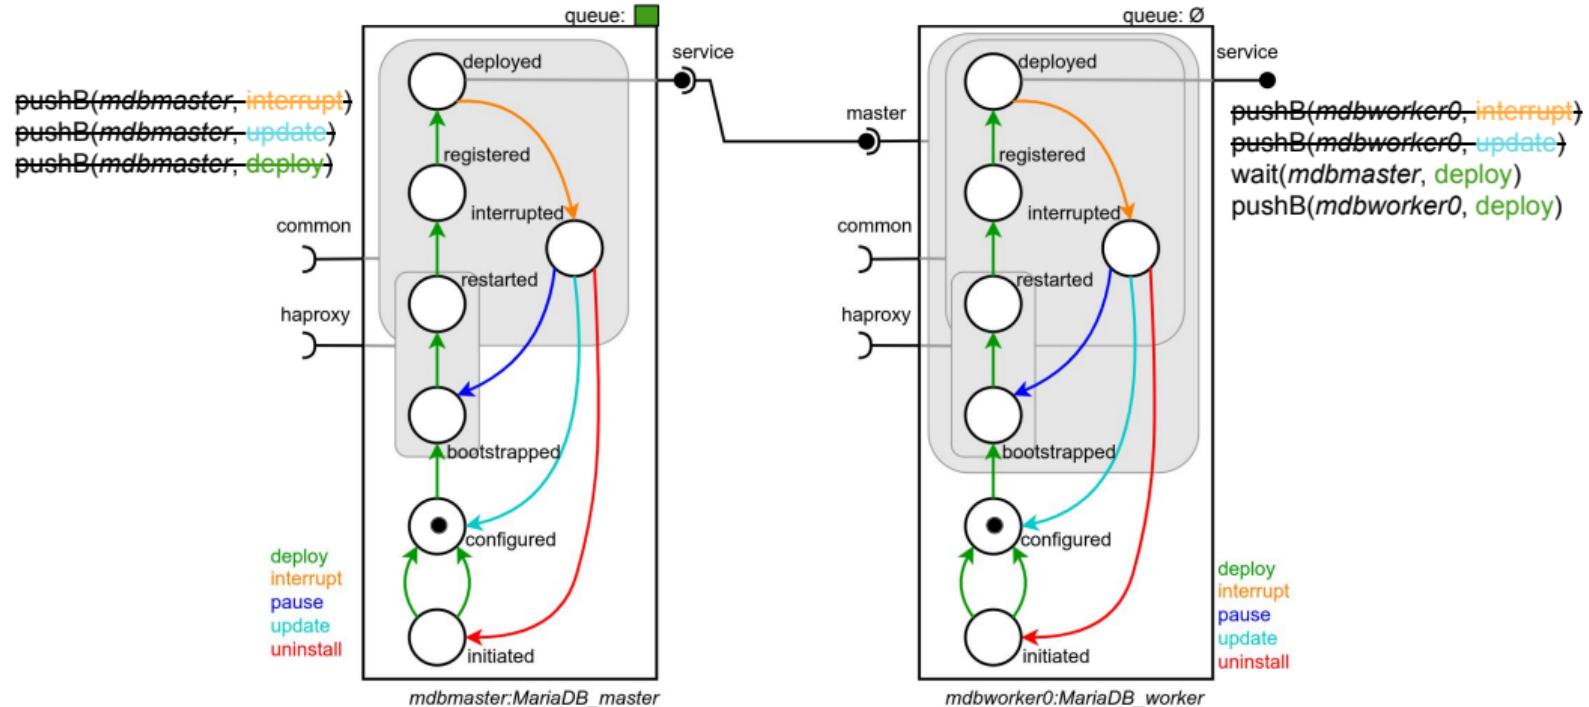
# Program example



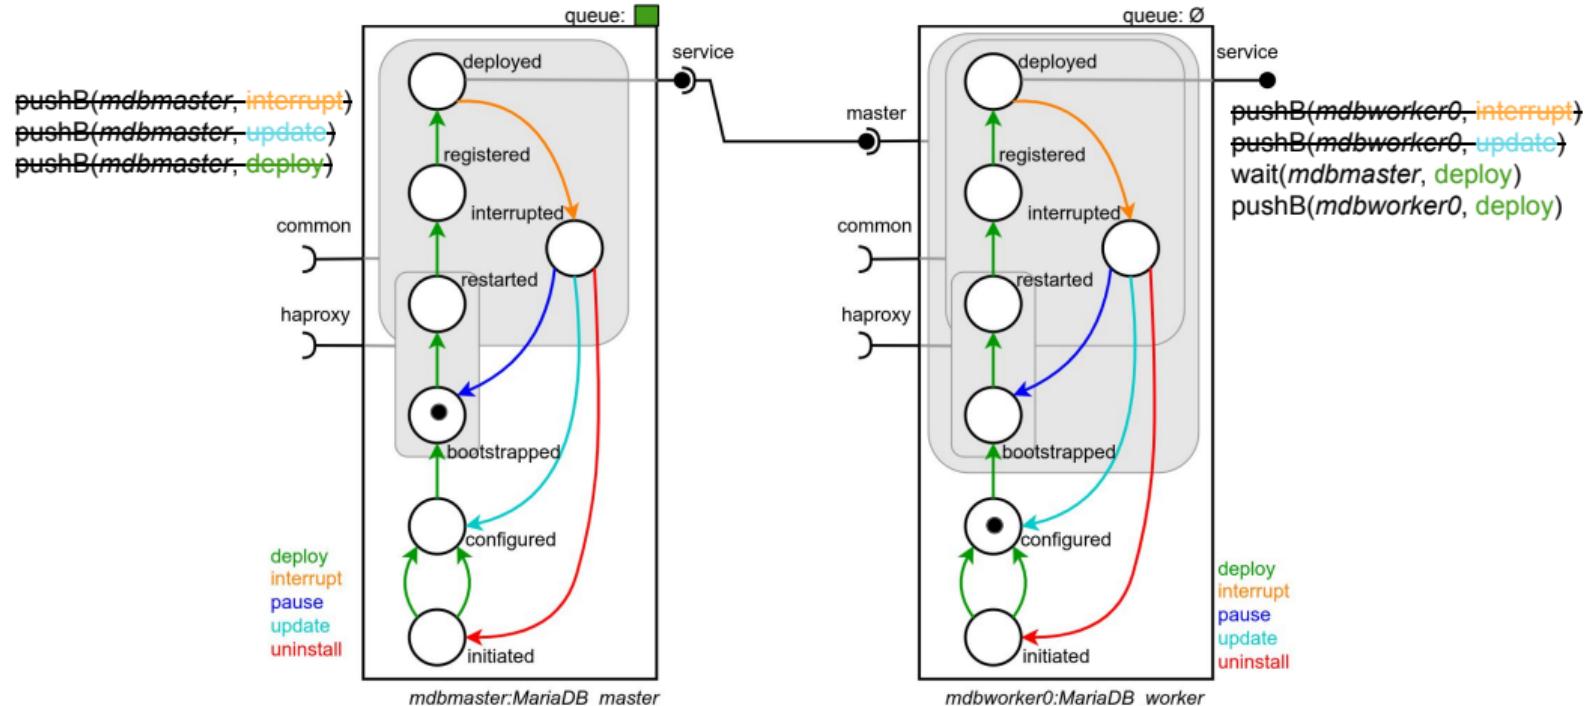
# Program example



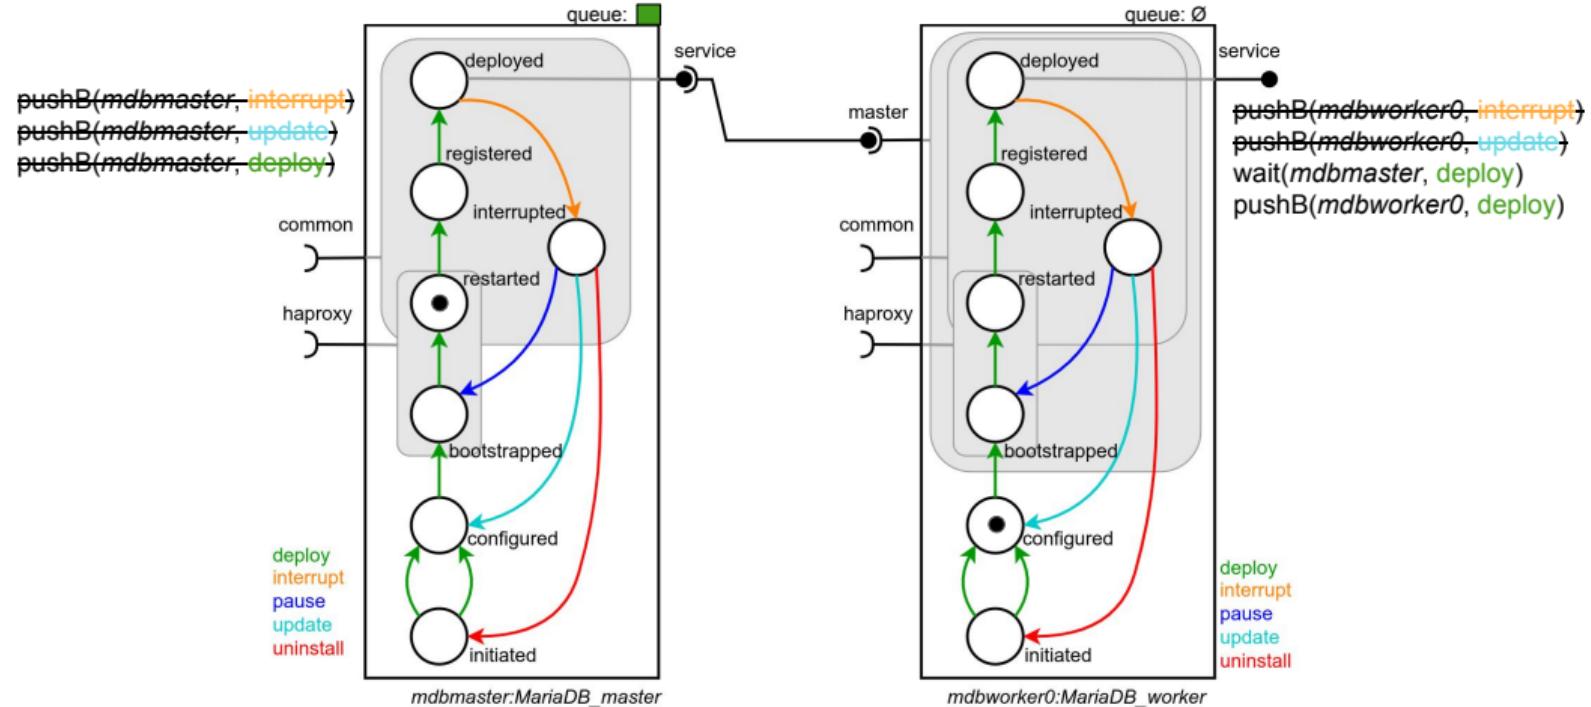
# Program example



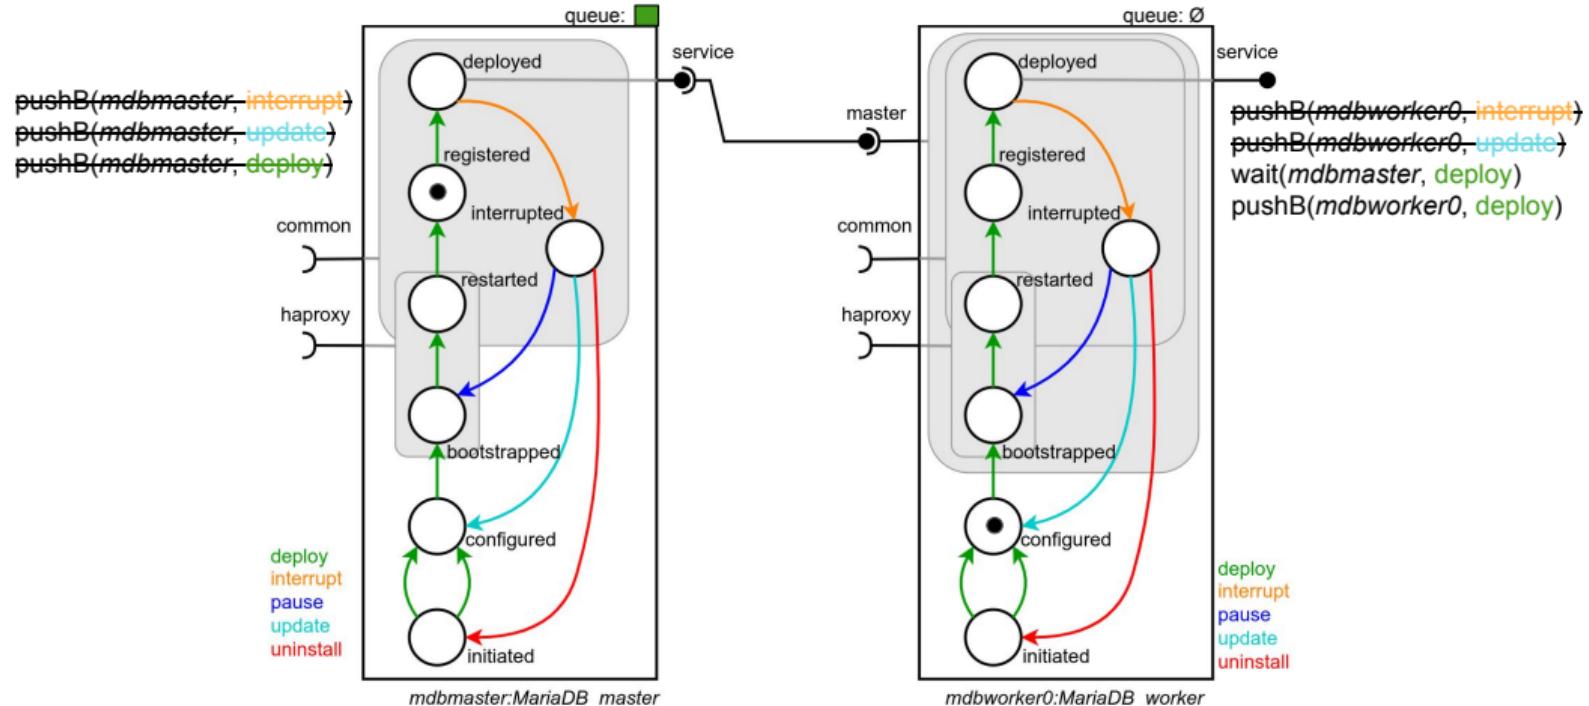
# Program example



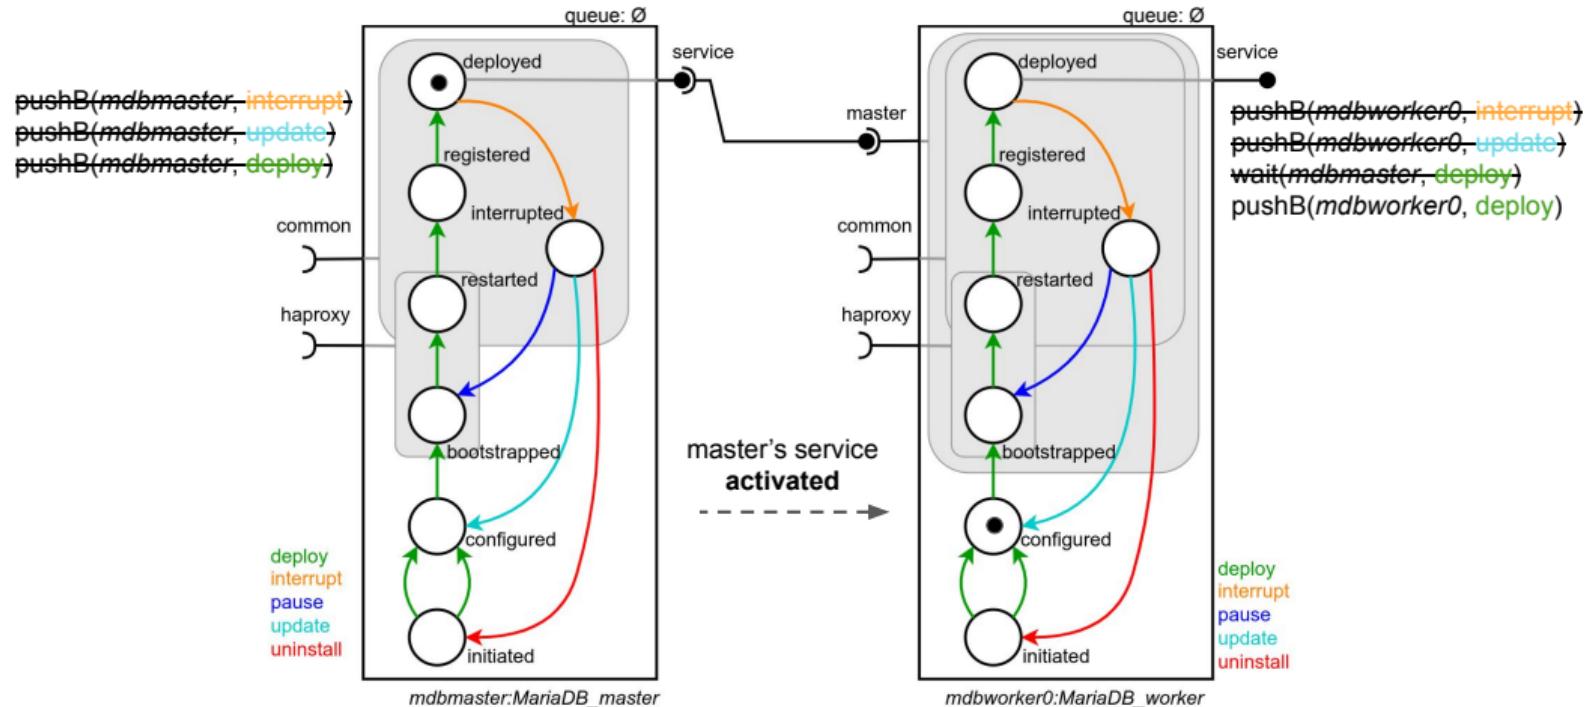
# Program example



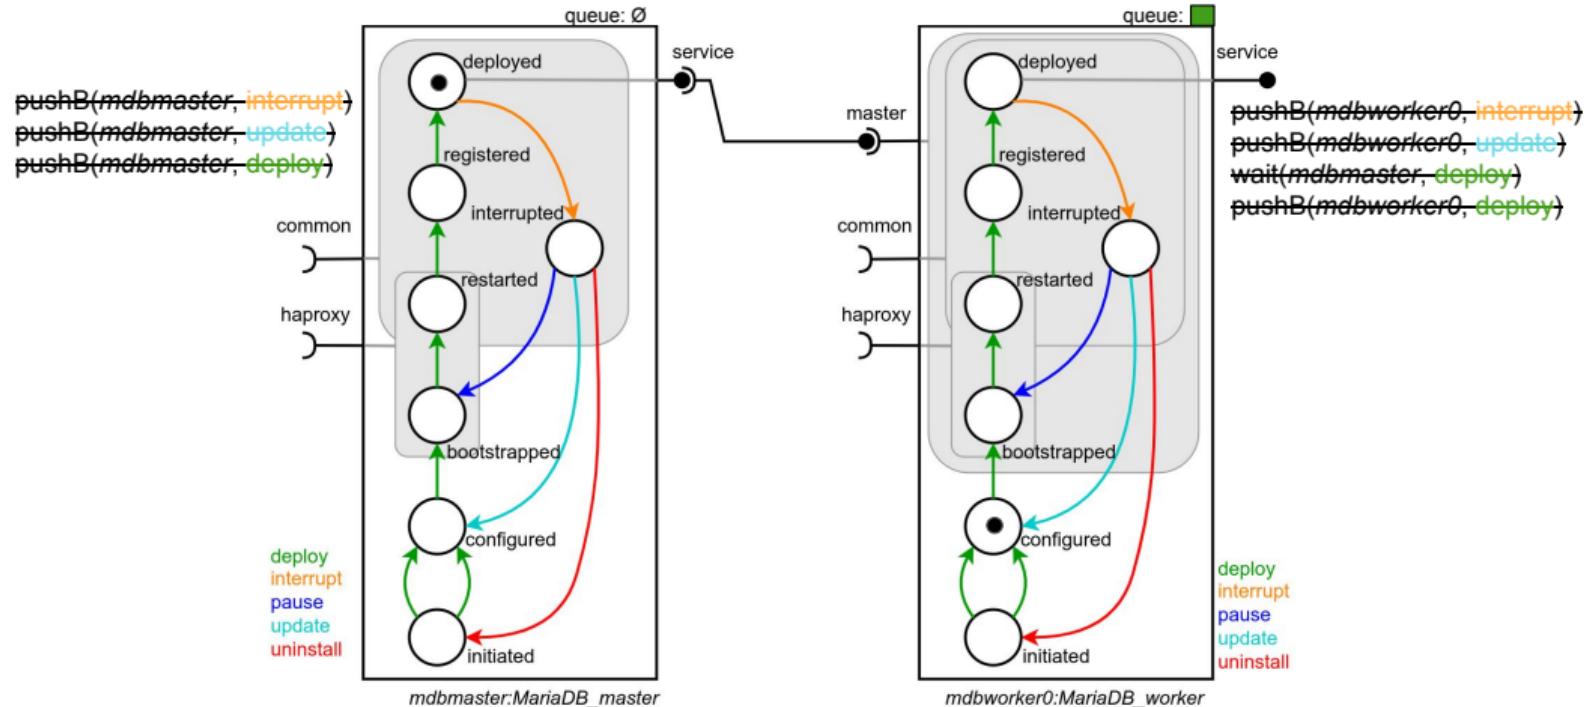
# Program example



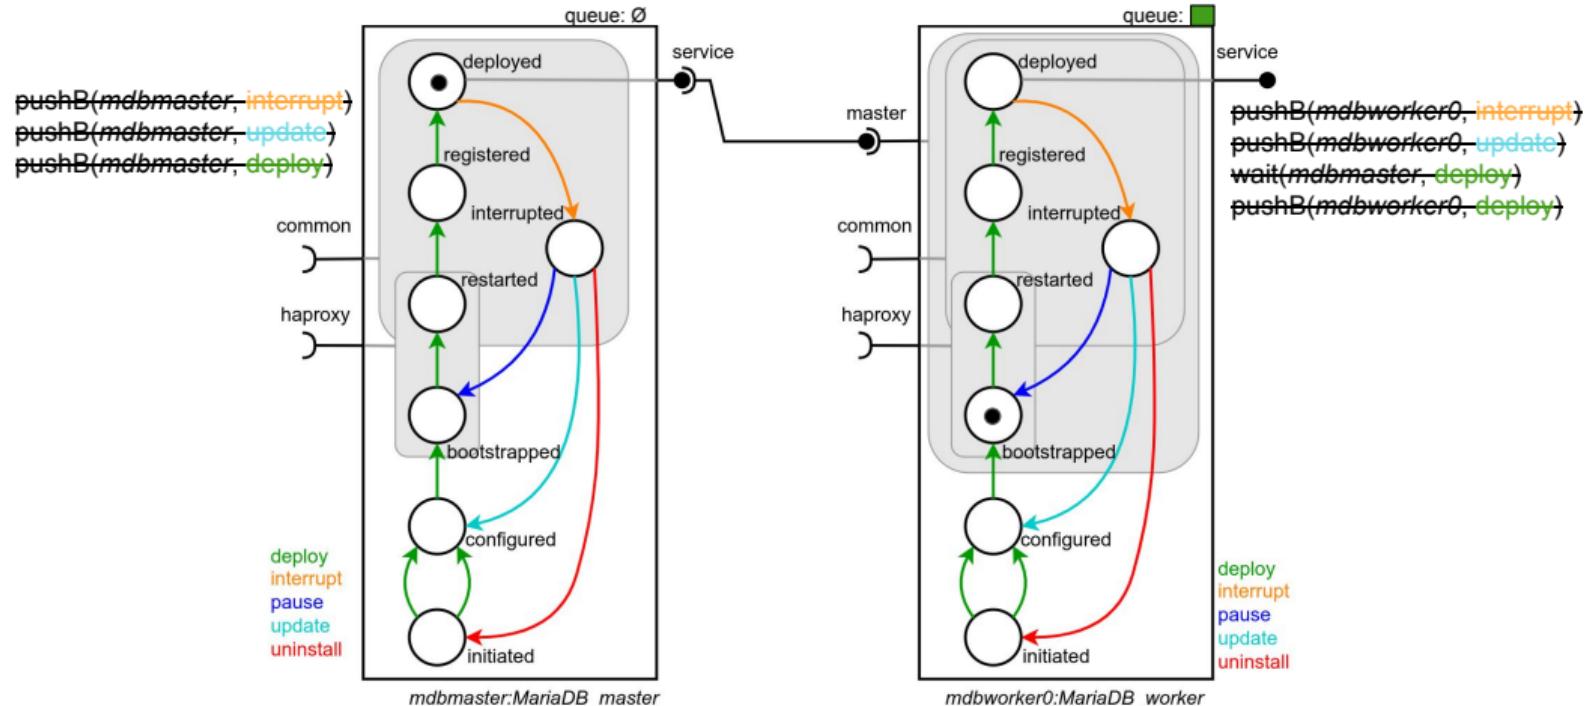
# Program example



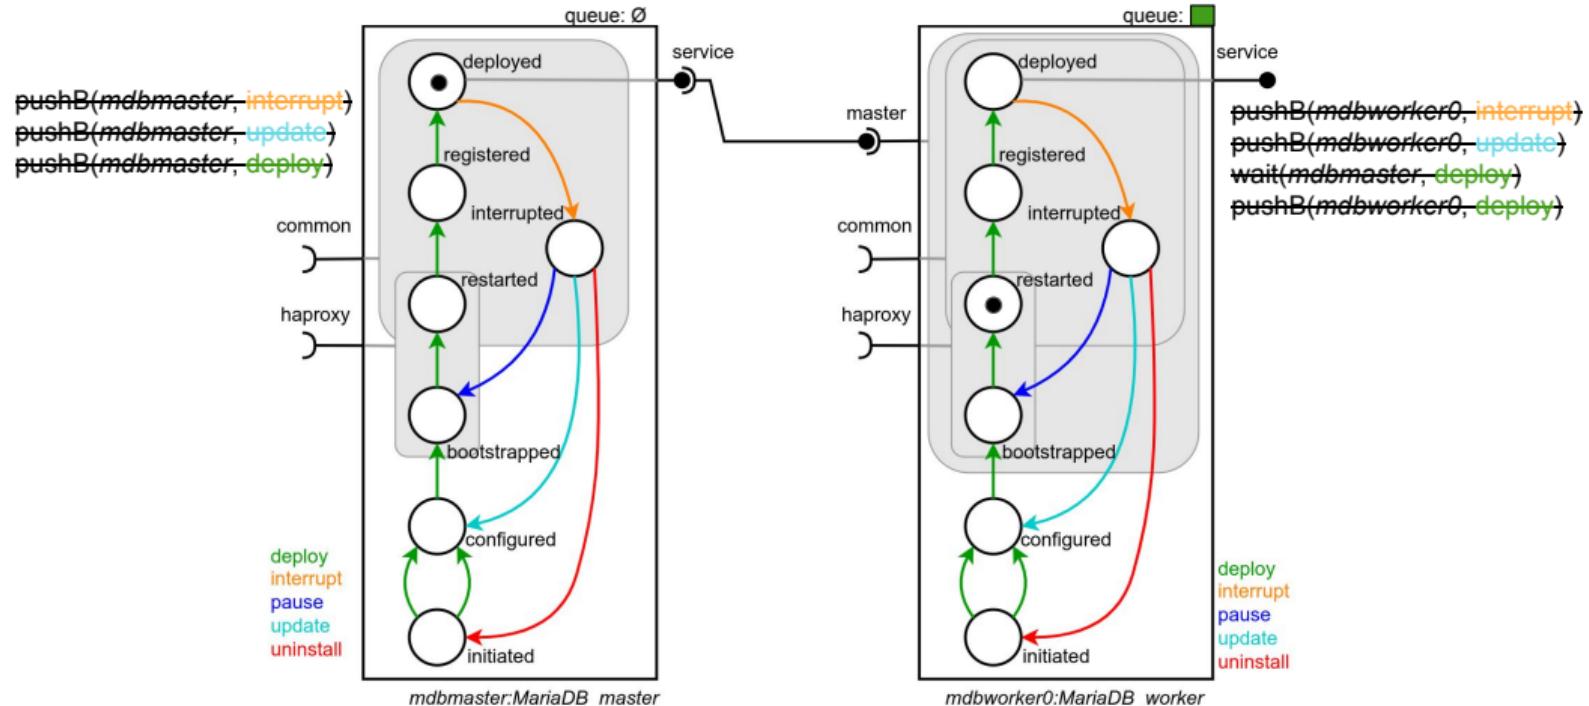
# Program example



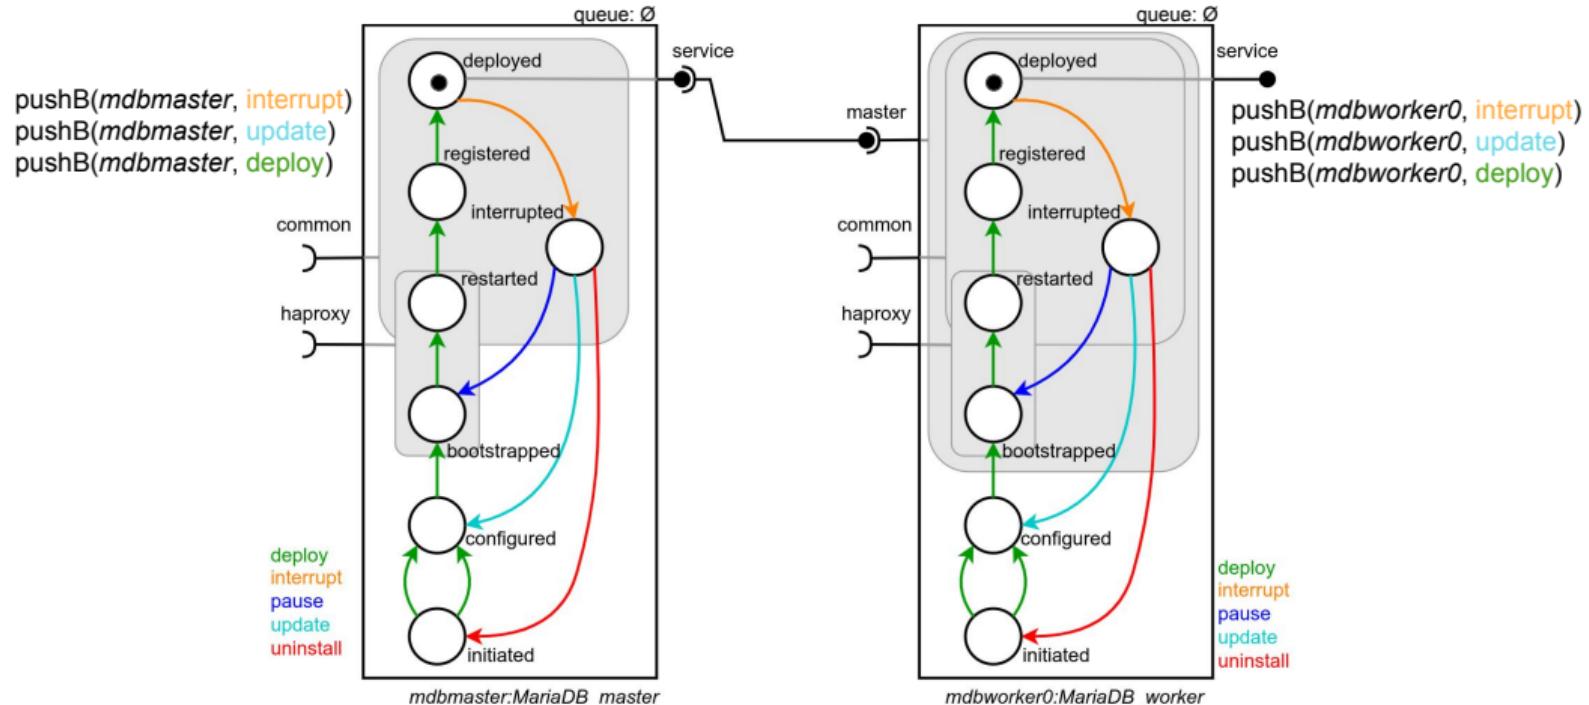
# Program example



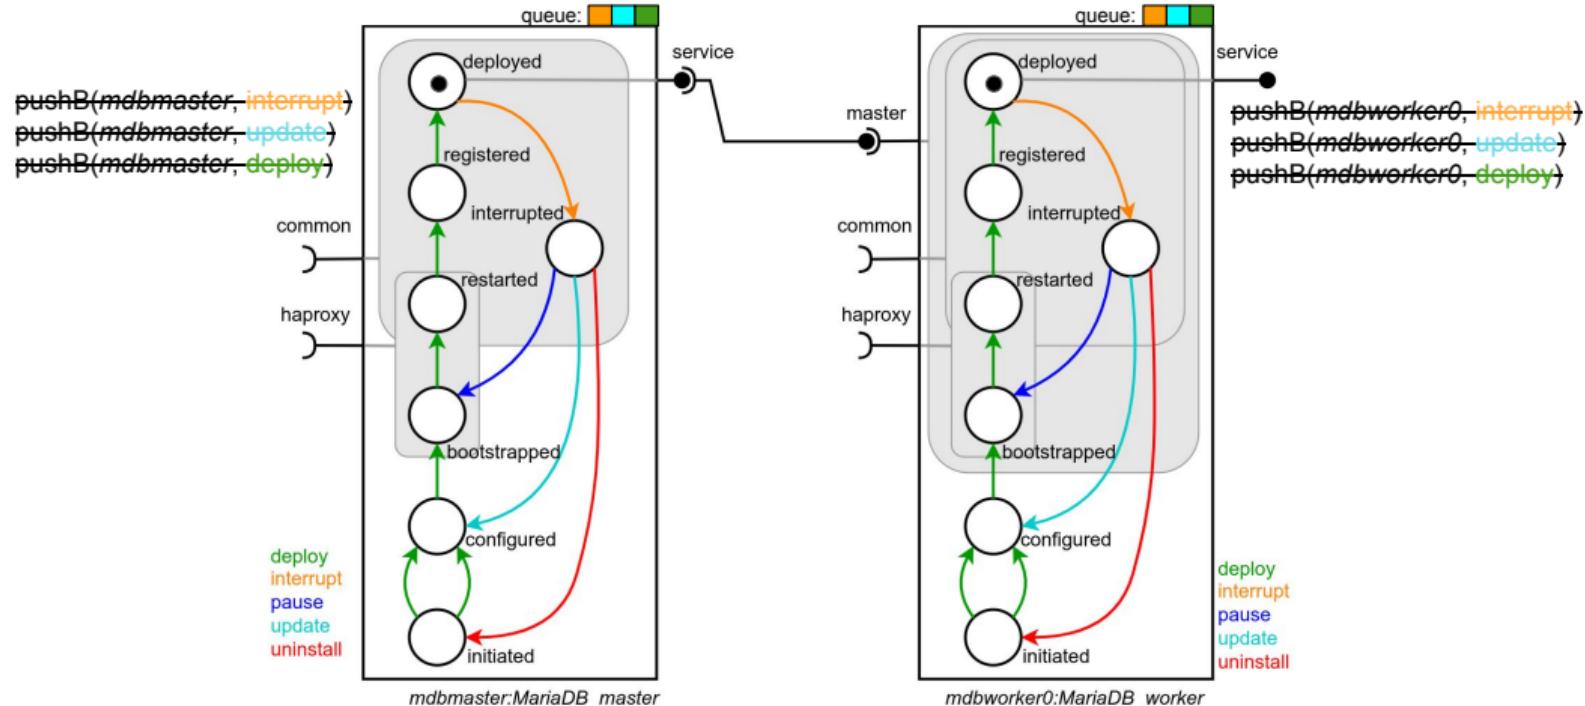
# Program example



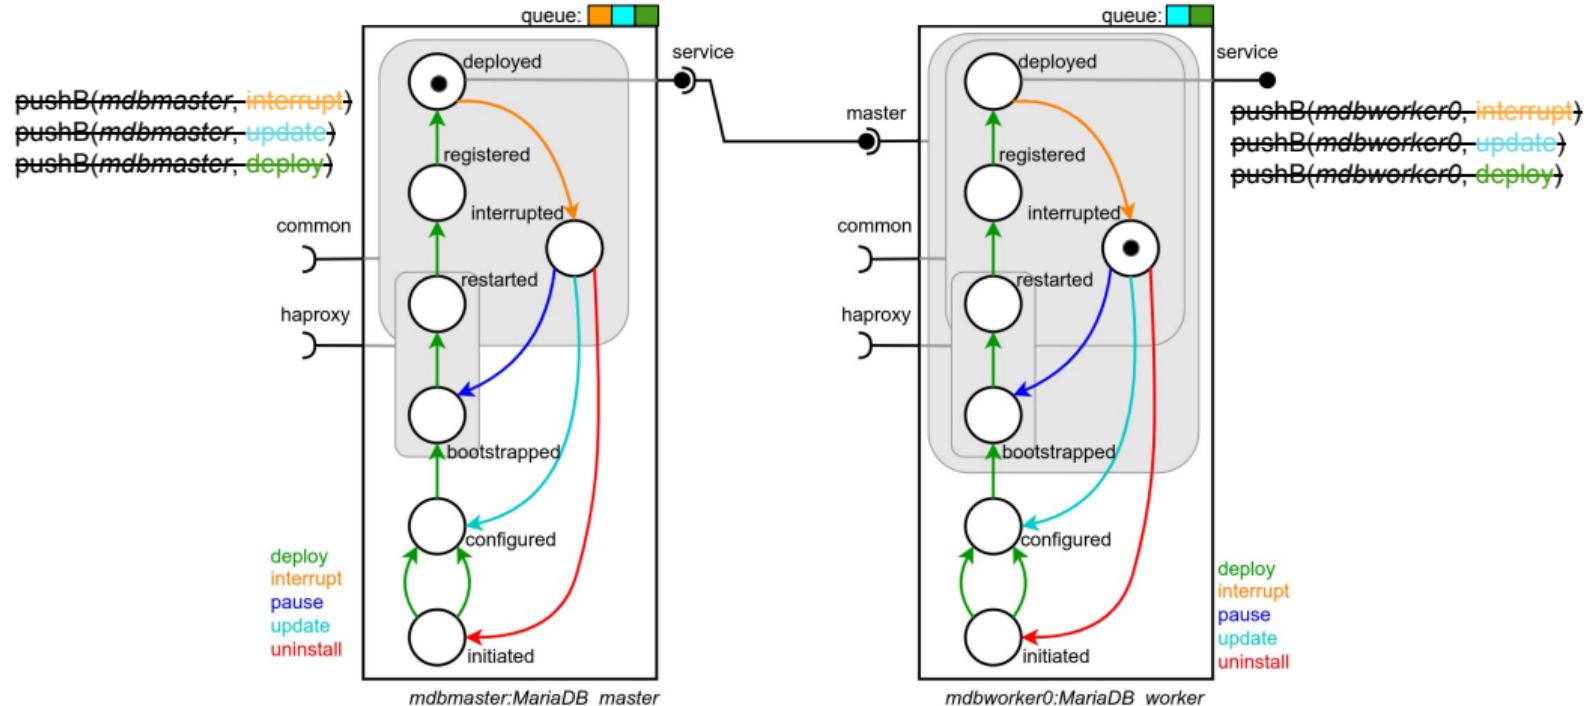
# Failing example



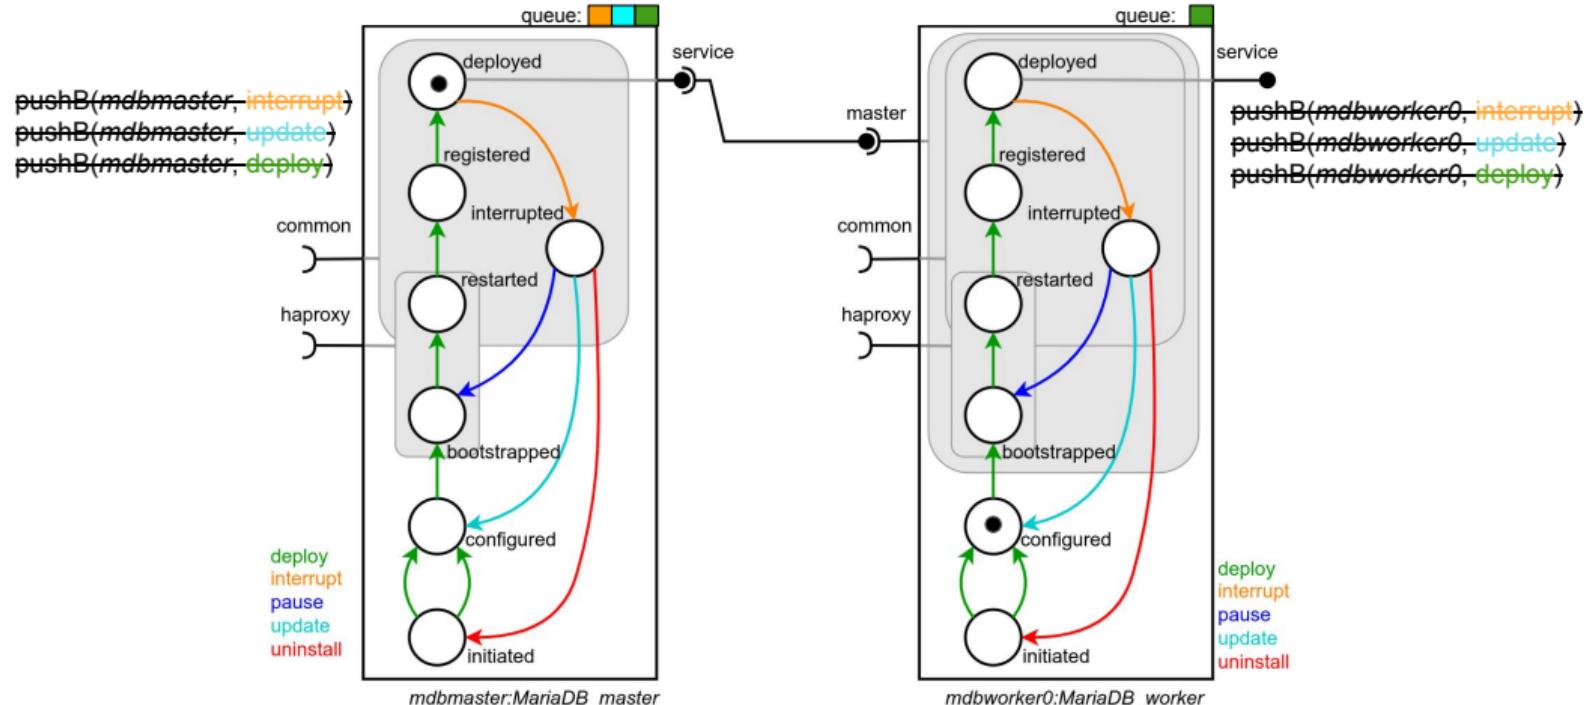
# Failing example



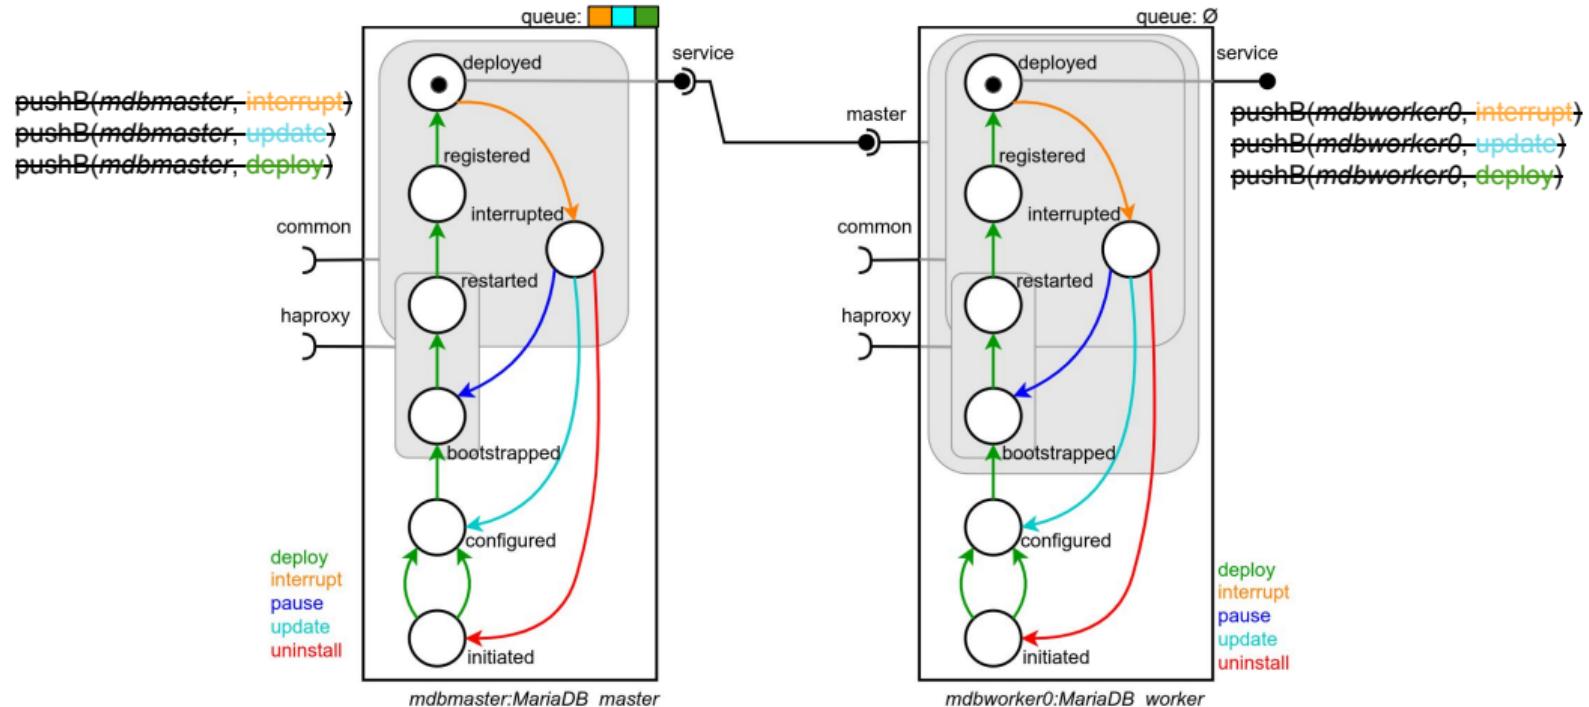
# Failing example



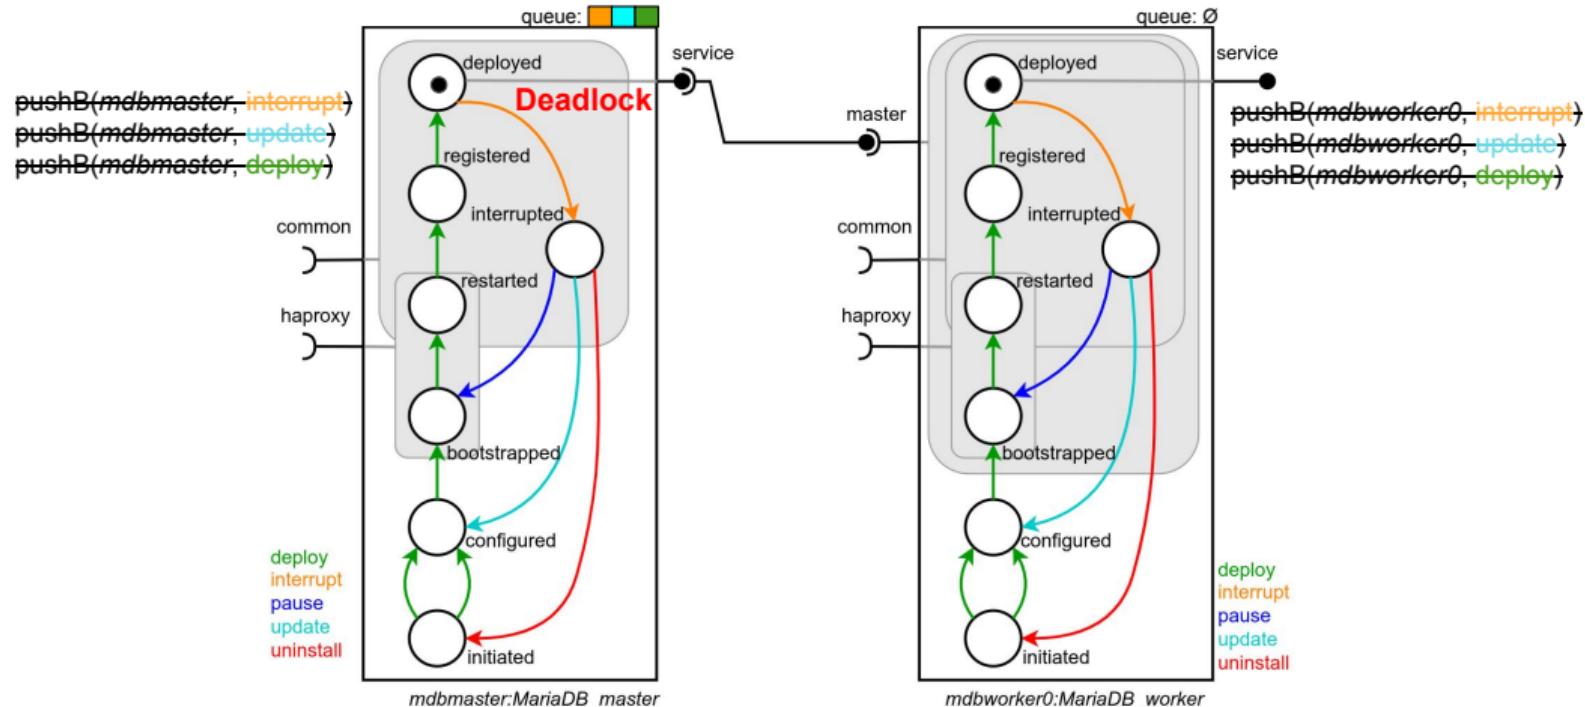
# Failing example



# Failing example



# Failing example



# Planner - Purpose and challenge

**Goal:** Infer local reconfiguration plan

## Assembly

- **Goal:** Infer add/connect and delete/disconnect instructions
- **Approach:** Make a diff between assembly.yaml and current state of the assembly
- **Challenge?**

## Local behaviors and sync. barriers

- **Goal:** Find a sequence of pushB and wait
- **Approach:** Constraint programming
- **Challenge:** Infer synchronization barriers (i.e., wait instructions)

# Global approach

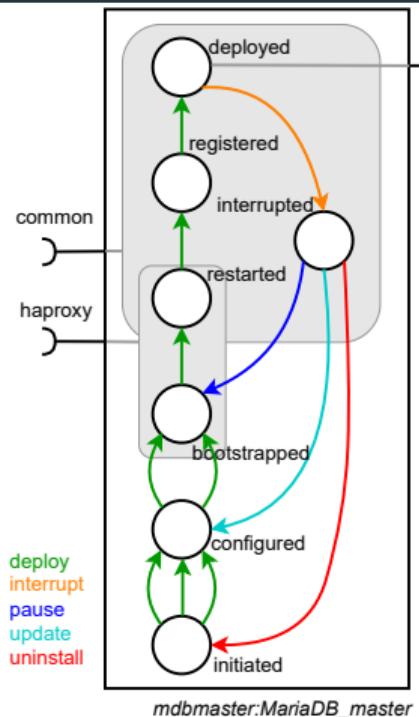
## Local resolution

- Model components as automata
- Find a word in the automata
- Add constraint on the word to meet with reconfiguration goals

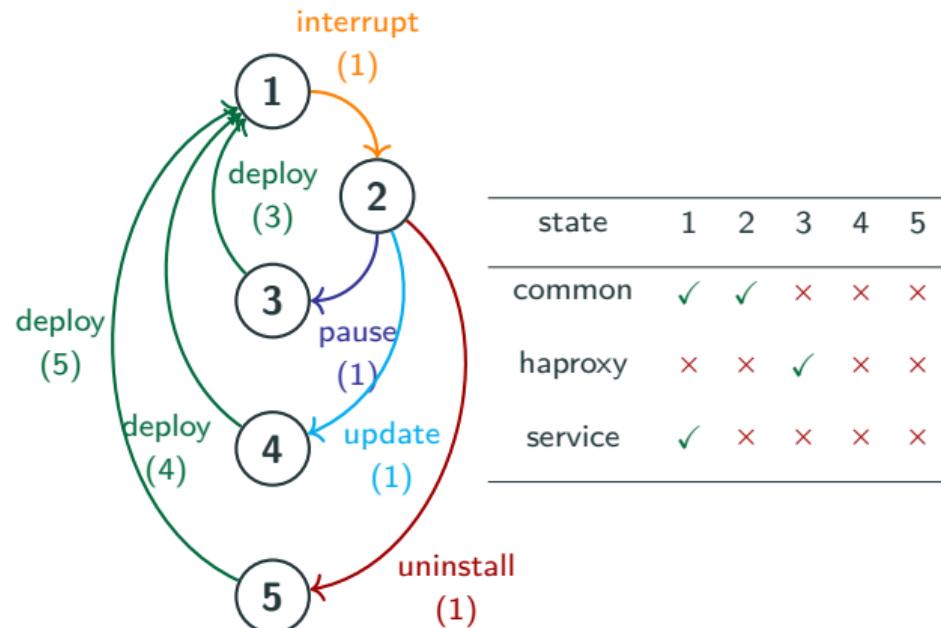
## Constraint diffusion

- Analyze what would be the impact on port
- Infer messages to send new constraints to neighbors
- Gossip approach

## Example: MariaDB master as an automaton



**Figure 4:** *MariaDB\_master* control component



**Figure 5:** Automaton representation of *Mariadb\_master* component's life cycle with its matrix for ports statuses.

# Find a sequence for MariaDB master

## Case study reconfiguration

### behaviors:

- **component**: mariadb\_master
- behavior**: update

### components:

- **forall**: running

## CP Result

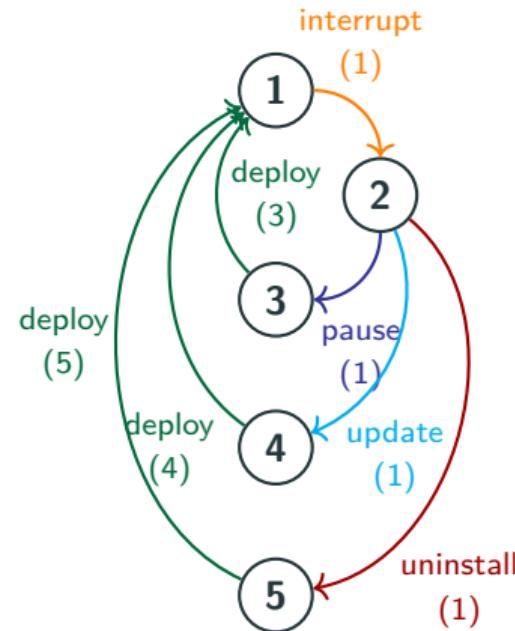
Sequence := [interrupt, update, deploy]

States := [1, 2, 4, 1]

common: [✓, ✓, ✗, ✓]

Ports statuses := haproxy: [✗, ✗, ✗, ✗]

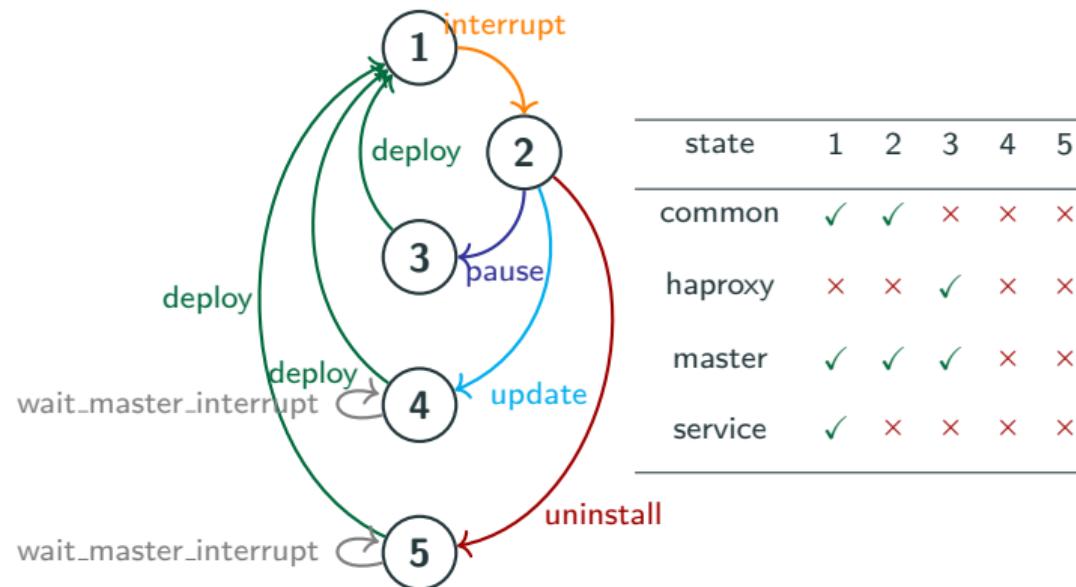
service: [✓, ✗, ✗, ✓]



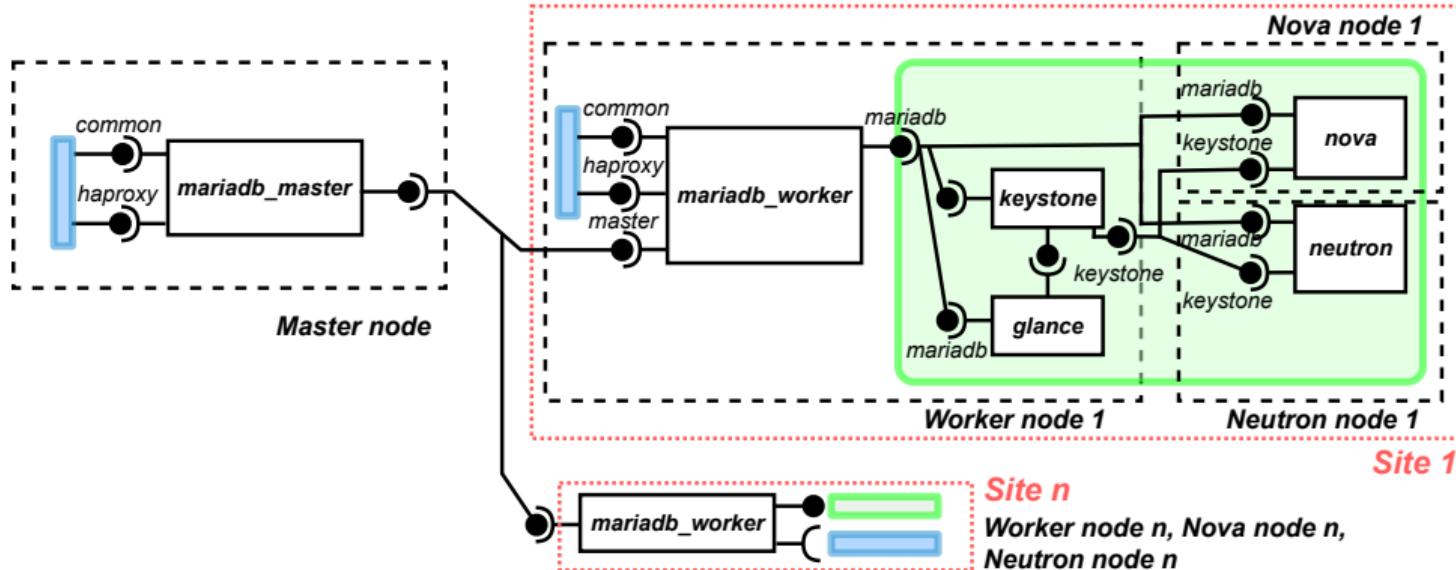
- “Components using **master’s service** must disconnect until **interrupt** ends”  
⇒ Message: (master, service, disconnect, interrupt)

# Enrich MariaDB workers model

- “Components using **master’s service** must disconnect until **interrupt** ends”
  - ⇒ Message: (master, service, disconnect, interrupt)
  - ⇒ Constraint: the use of master’s service must be turned off during the reconfiguration, and a wait instruction must be placed



# Constraint diffusion



Diffuse constraint from `mariadb_master` for master's service

- `mariadb_master`  $\Rightarrow$  `mariadb_worker`
- `mariadb_worker`  $\Rightarrow$  `keystone`; `glance`; `nova`; `neutron`
- `keystone`  $\Rightarrow$  `glance`; `nova`; `neutron`

## Implementation of Ballet

- Open source: <https://github.com/Concerto-D/Ballet>
- about 10000 LoC (Python)
- Communications are operated with gRPC
- Planner uses MiniZinc modeling language with C++ Gecode solver as backed

# Experiments

Deployment and update of OpenStack with Galera cluster of MariaDB with  $n \in [1, 2, 5, 10]$  sites, that is a total of  $7 + 11 * n$  components.

## Metric of interest

- For both the planner and the executor: Execution time
- For the planner: Inferred constraints, inferred actions, number of communications

## Setup

- Results on  $1 + 3 * n$  nodes Gros (Nancy) of Grid'5000
- Comparison to Muse (decentralized reconfiguration)
- Reproducible example on Grid'5000

## Experimental results - Performance

Sc.	# Sites	Ballet			Muse	Gain
		Planning	Execution	Total		
Deploy	1	1.69s	306.02s	307.71s	536.57s	42.7%
	2	1.78s	306.09s	307.86s	536.69s	42.6%
	5	1.77s	306.19s	307.97s	537.09s	42.7%
	10	2.02s	306.14s	308.19s	538.13s	42.7%
Update	1	3.36s	416.84s	420.20s	555.56s	24.4%
	2	4.39s	416.92s	421.31s	555.70s	24.2%
	5	6.05s	417.17s	423.22s	556.08s	24.0%
	10	5.97s	417.46s	423.43s	556.77s	24.0%

**Table 1:** Comparison of time for planning and executing a deployment and an update of the MariaDB\_master instance with Ballet and Muse.

## Experimental results - Inference

Sc.	#Sites	#Constraints	#Instructions	#Messages
Deploy	$n$	$7 + 11 * n$	$7 + 11 * n$	0
	1	18	18	0
	2	29	29	0
	5	62	62	0
	10	117	117	0
Update	$n$	$3 + 20 * n$	$8 + 11 * n$	$9 * n$
	1	23	19	9
	2	43	30	18
	5	103	63	45
	10	203	118	90

**Table 2:** Results of the planning phase for the *deploy* and *update* scenario when varying the number of Mariadb\_workers in a Galera cluster.

# Few more words

## Ongoing and future work

1. Study and management of unsat reconfiguration
  - Find the minimum Minimal Unsatisfiable Subsets of Constraints
  - Return conflicts to DevOps teams
2. Maude model of Ballet's executor (Concerto-D) for reasoning
3. Study safety properties of reconfiguration in modern IaC languages (Terraform, Pulumi)

## Publications

[SANER24] Jolan Philippe, Antoine Omond, Hélène Coullon, Charles Prud'Homme, and Issam Raïs. Fast Choreography of Cross-DevOps Reconfiguration with Ballet: A Multi-Site OpenStack Case Study.

[ICE24] Farid Arfi, Hélène Coullon, Frédéric Loulergue, Jolan Philippe, and Simon Robillard. A Maude Formalization of the Distributed Reconfiguration Language Concerto-D.

# Conclusion

## Contributions

- Ballet as a DevOps reconfiguration tool
- Code and benchmark: <https://zenodo.org/records/10472116>
- Infer reconfiguration actions
- Efficient execution of actions

## Target applications

- Multi-site OpenStack
- CPS with sensors
- Management of Fog area

# Backup

---

# CP Model

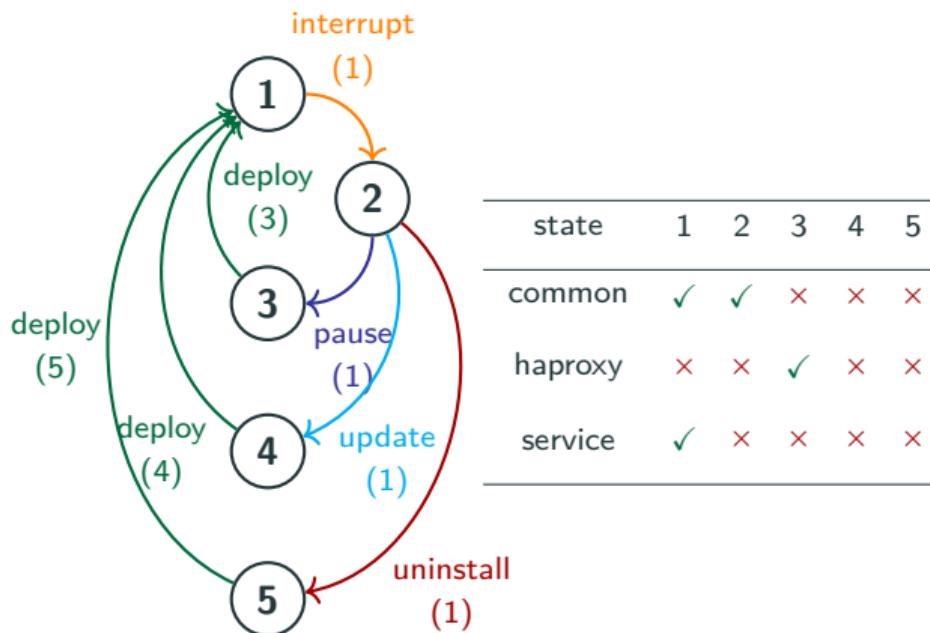


Figure 6: Automaton representation of *Mariadb\_master* component's life cycle with its matrix for ports statuses.

- COSTREGULAR( $B, \Pi, \mathcal{C}, s_{init}, S_{goal}$ )

- $s_{i+1} = inc_{\Pi}[s_i][b_i], \forall i \in 1..m$

- COUNT( $b, B, >, 0$ )

- $status(p, s_{m+1}) = \Gamma_p$

**where**

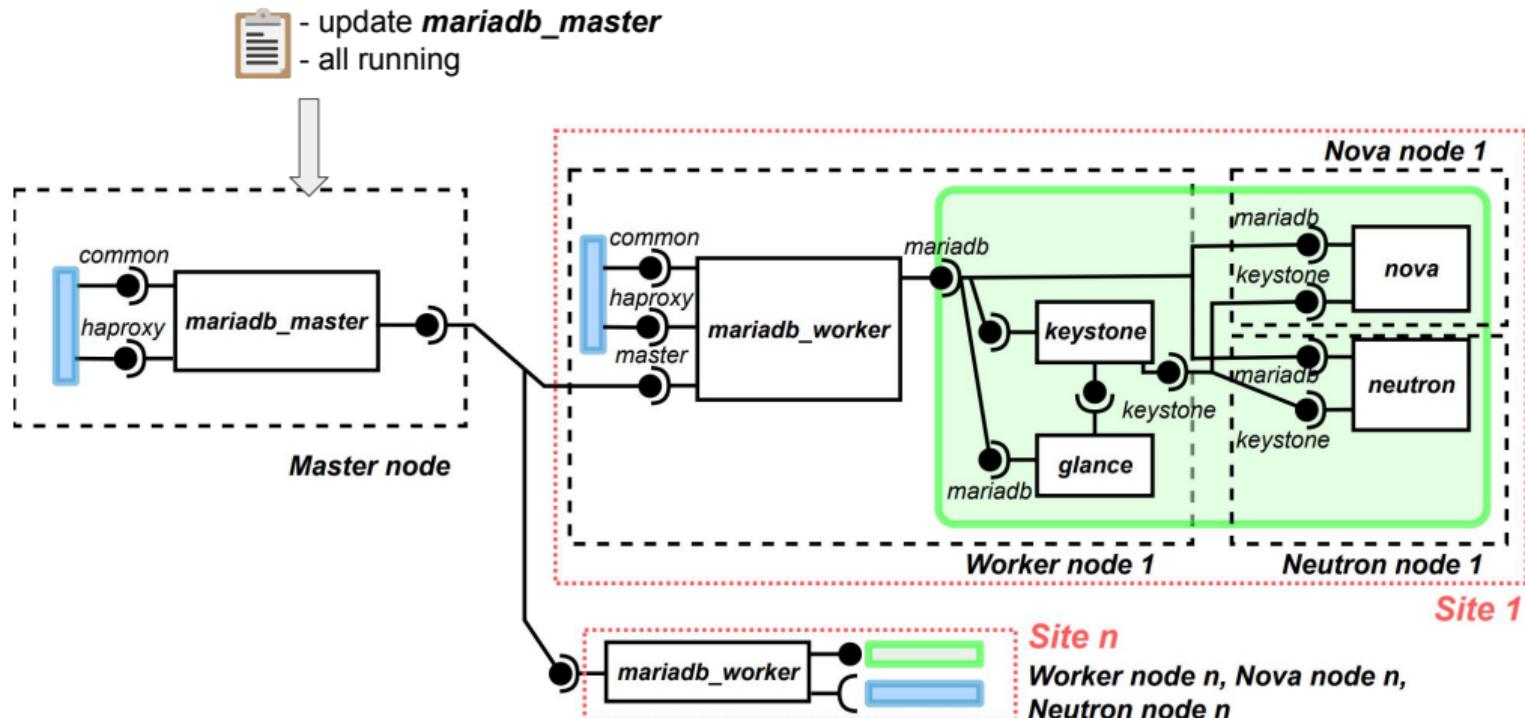
$\Pi$  an automaton with  $\mathcal{C}$  costs

$B$  a sequence of  $m$  behaviors

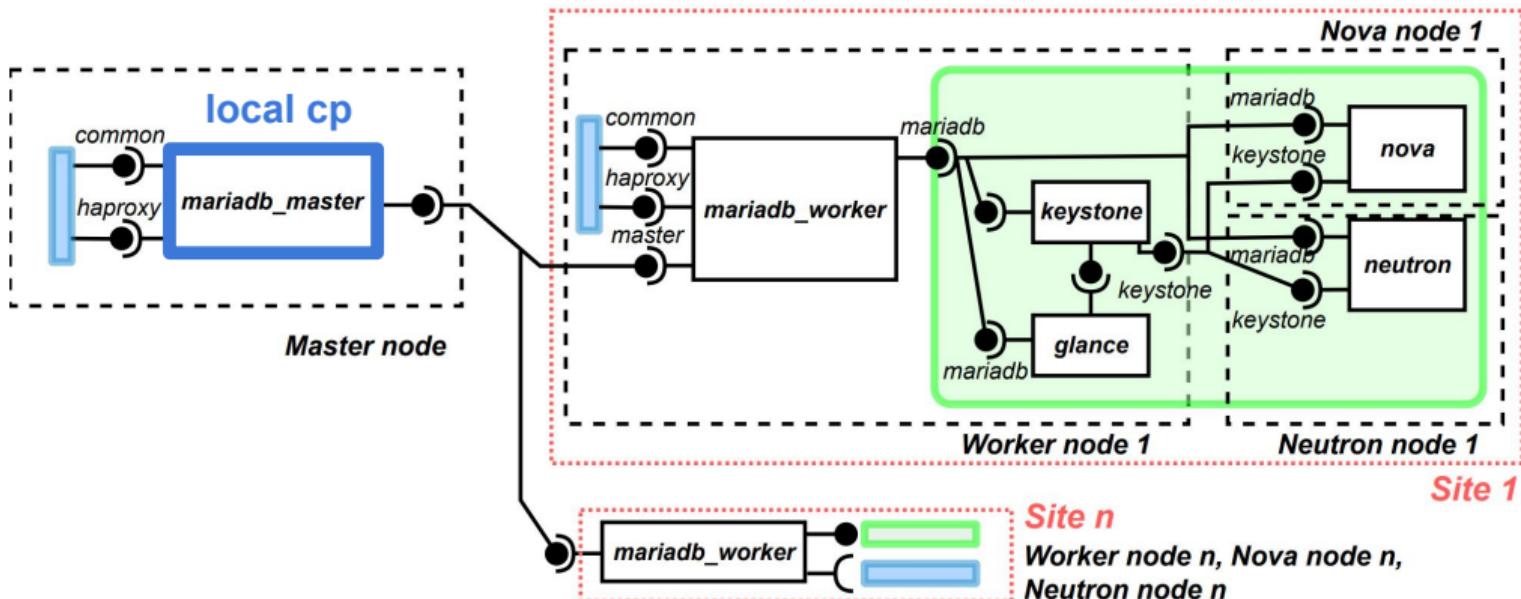
$\Gamma_p \in \{\text{active, inactive}\}$ i.e.  $\{ \checkmark, \times \}$

$b \in \{ \text{interrupt, deploy, pause, update, uninstall} \}$

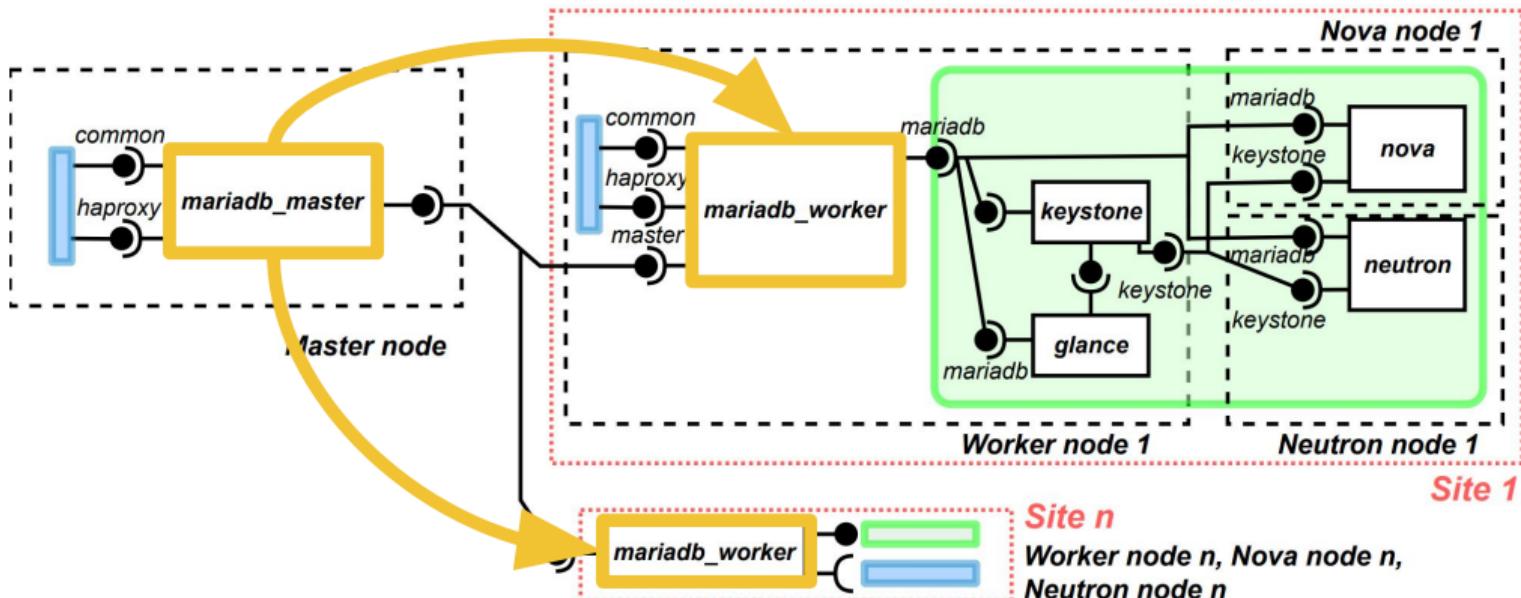
# Communication protocol



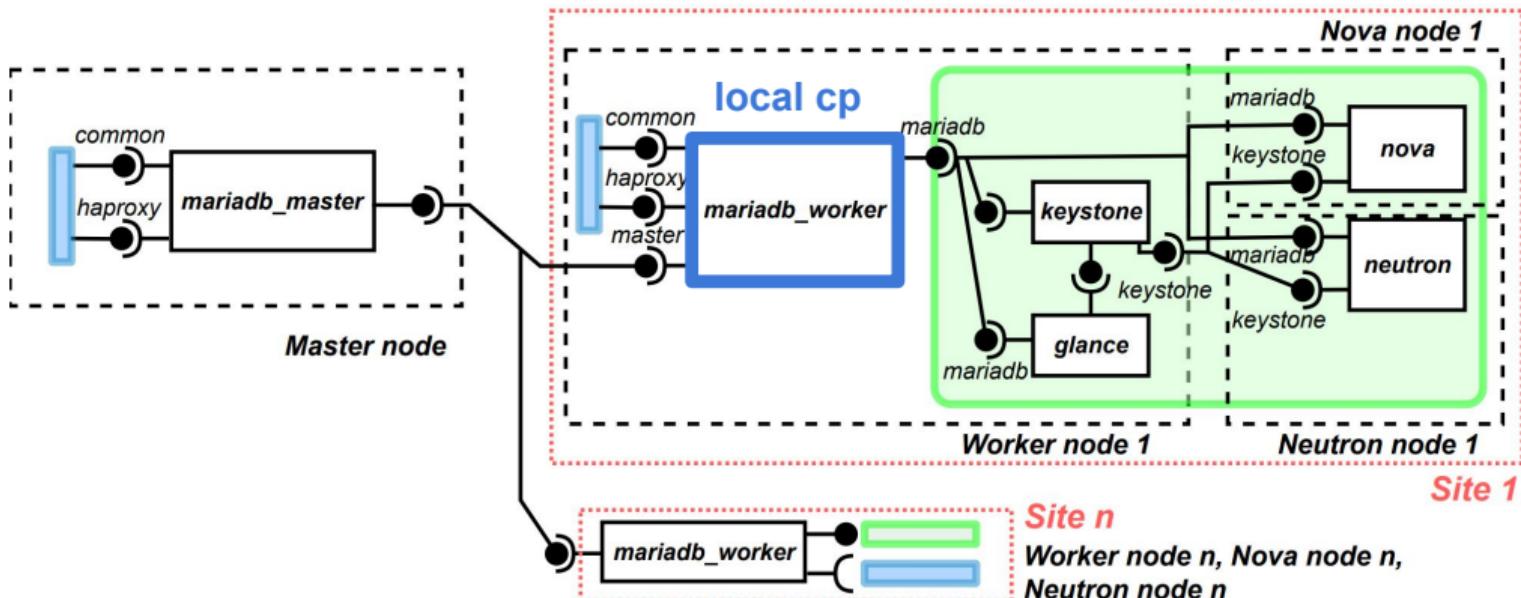
# Communication protocol



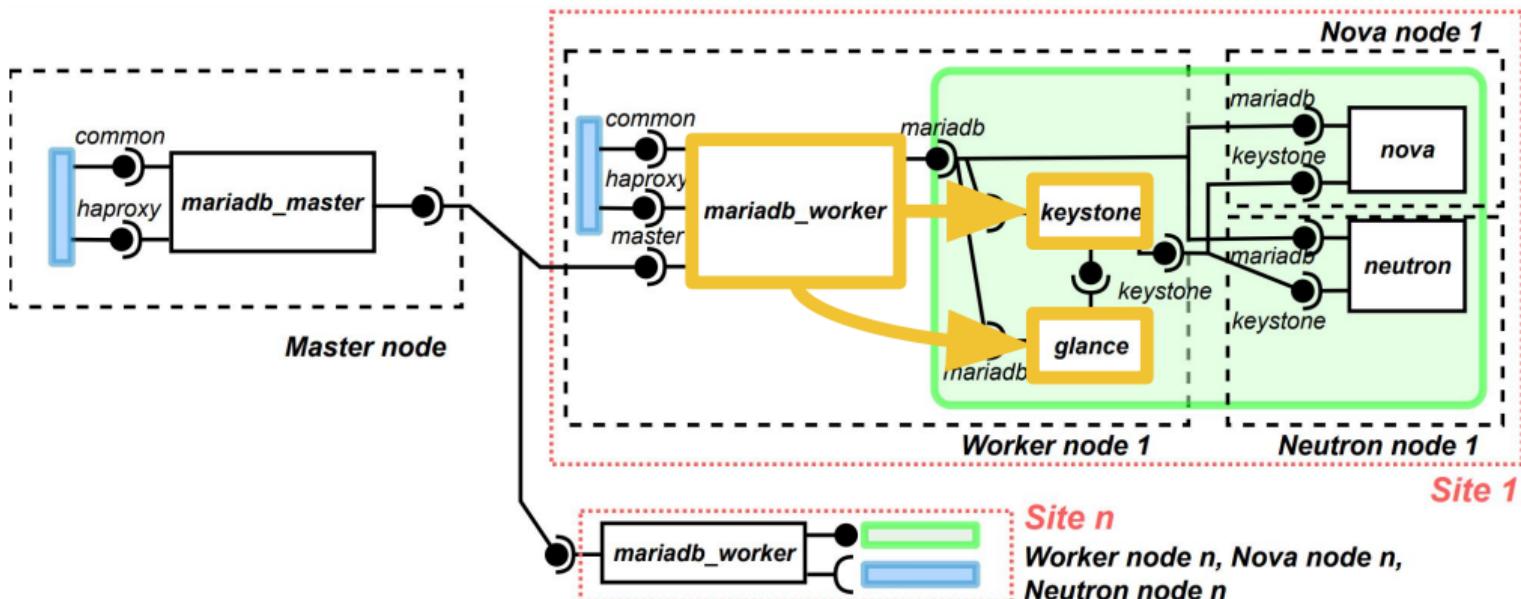
# Communication protocol



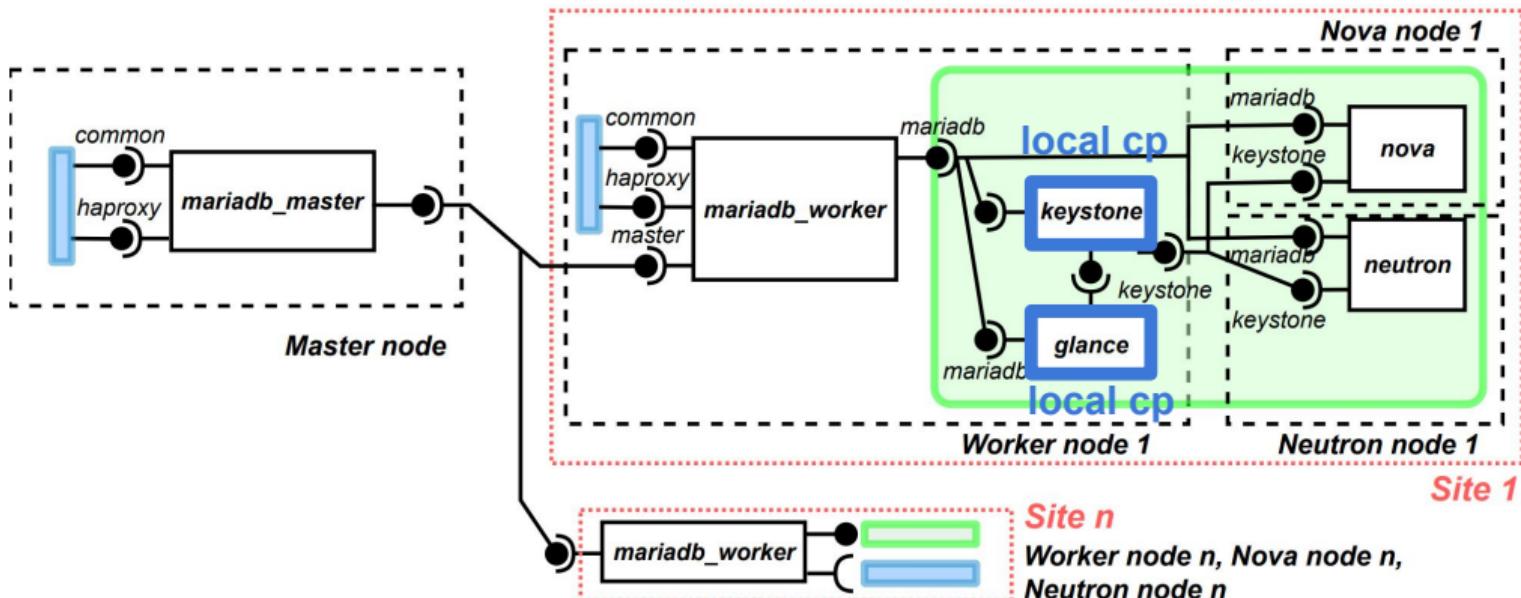
# Communication protocol



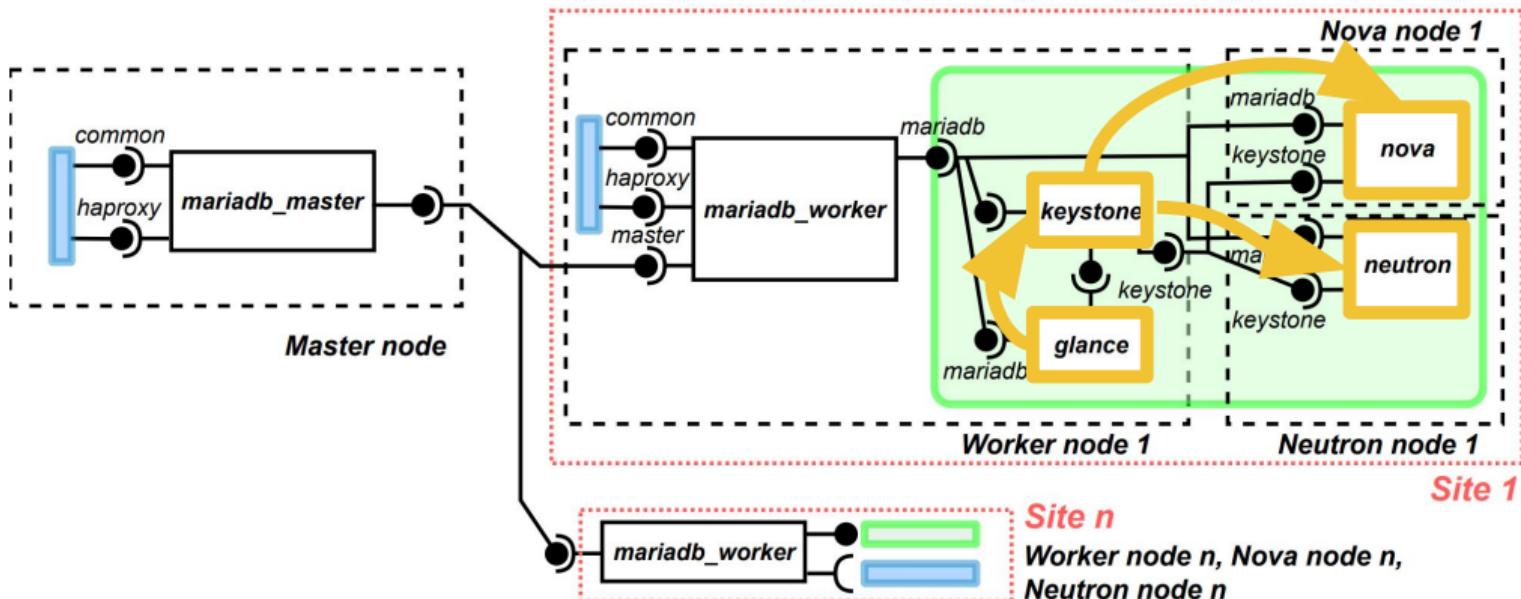
# Communication protocol



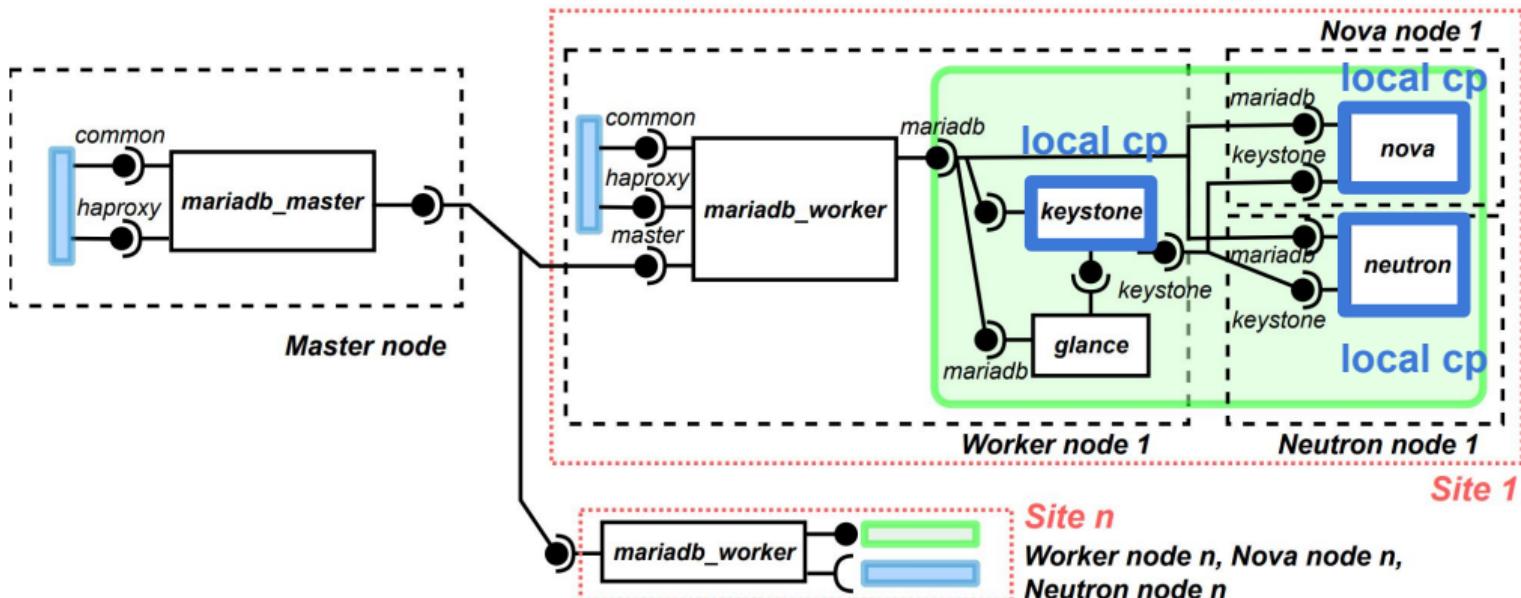
# Communication protocol



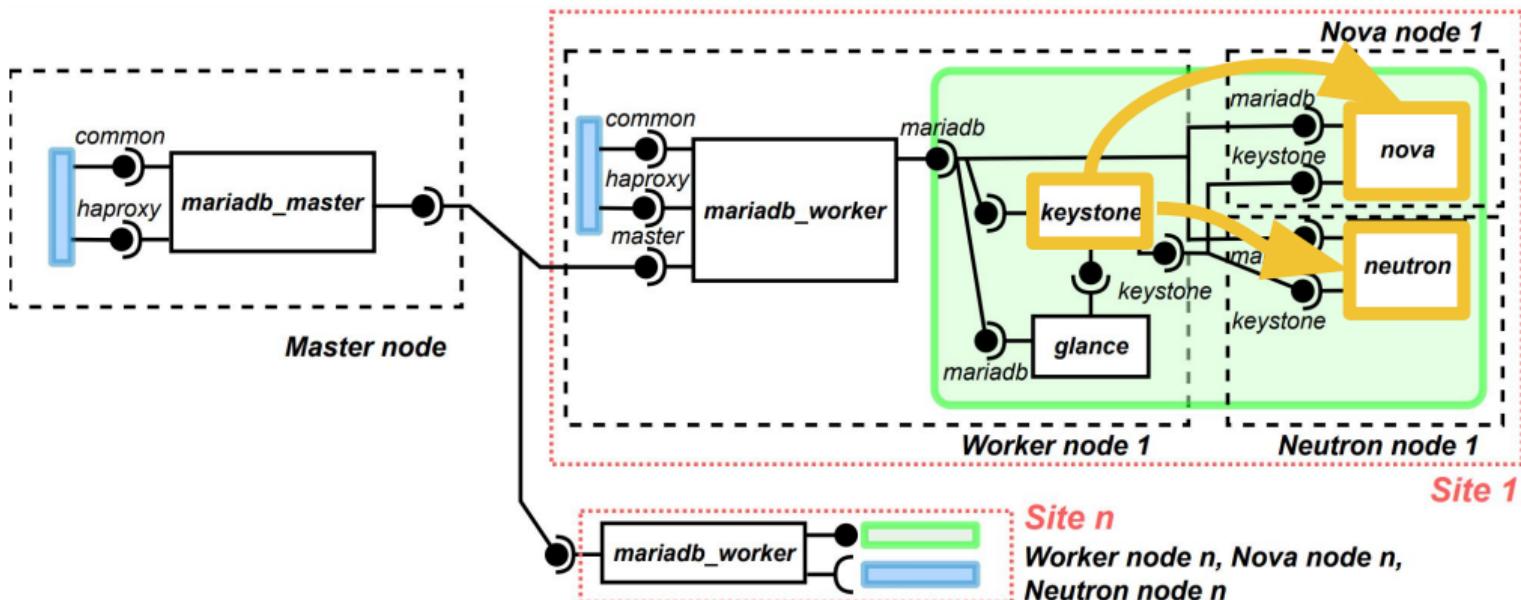
# Communication protocol



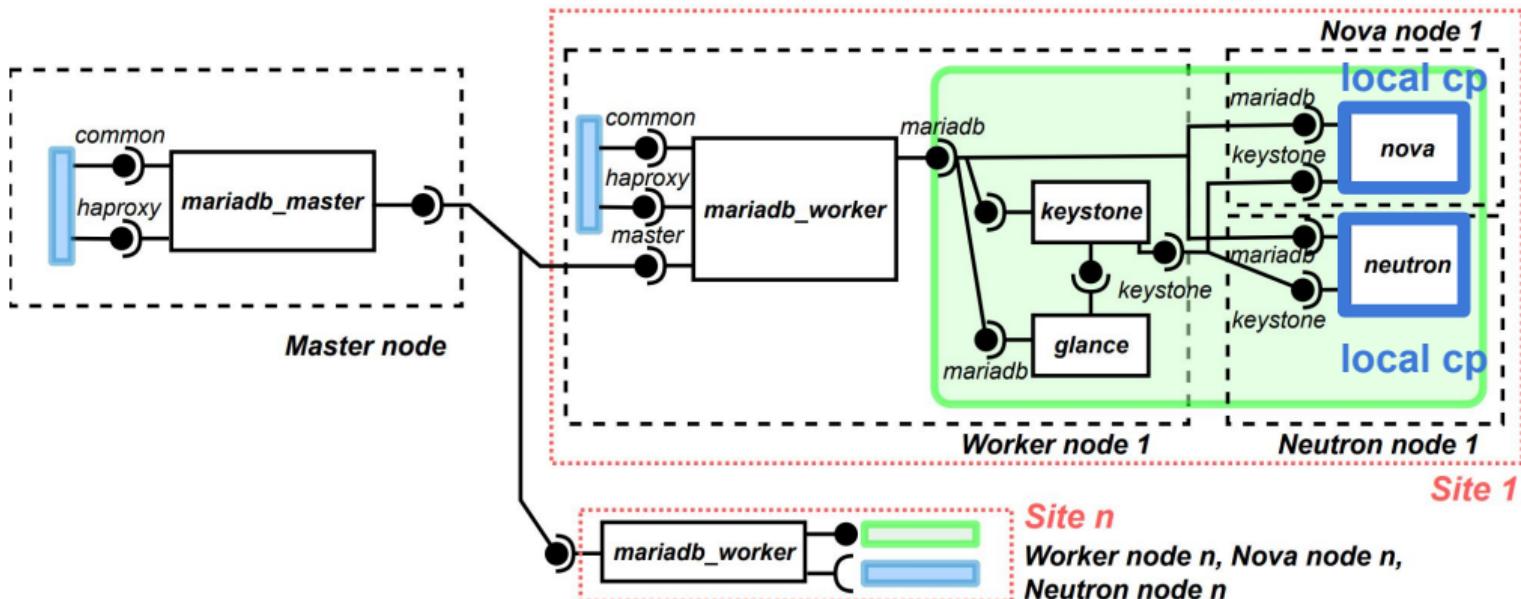
# Communication protocol



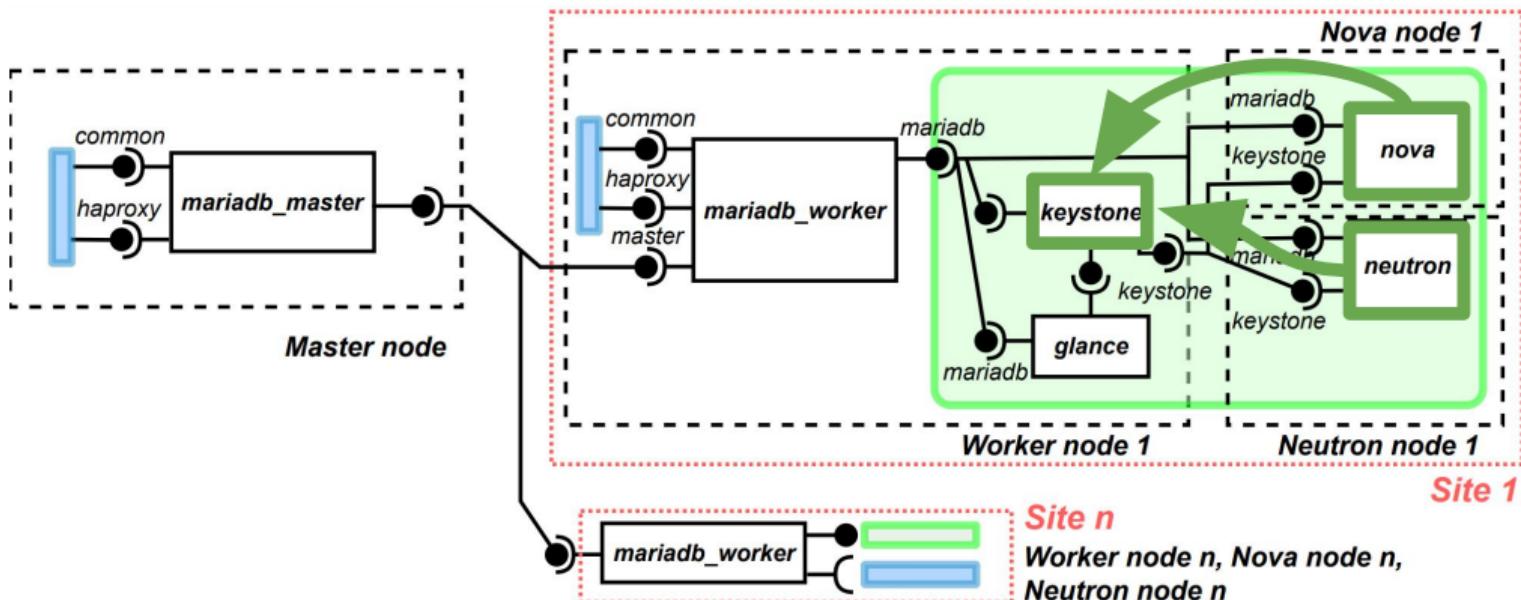
# Communication protocol



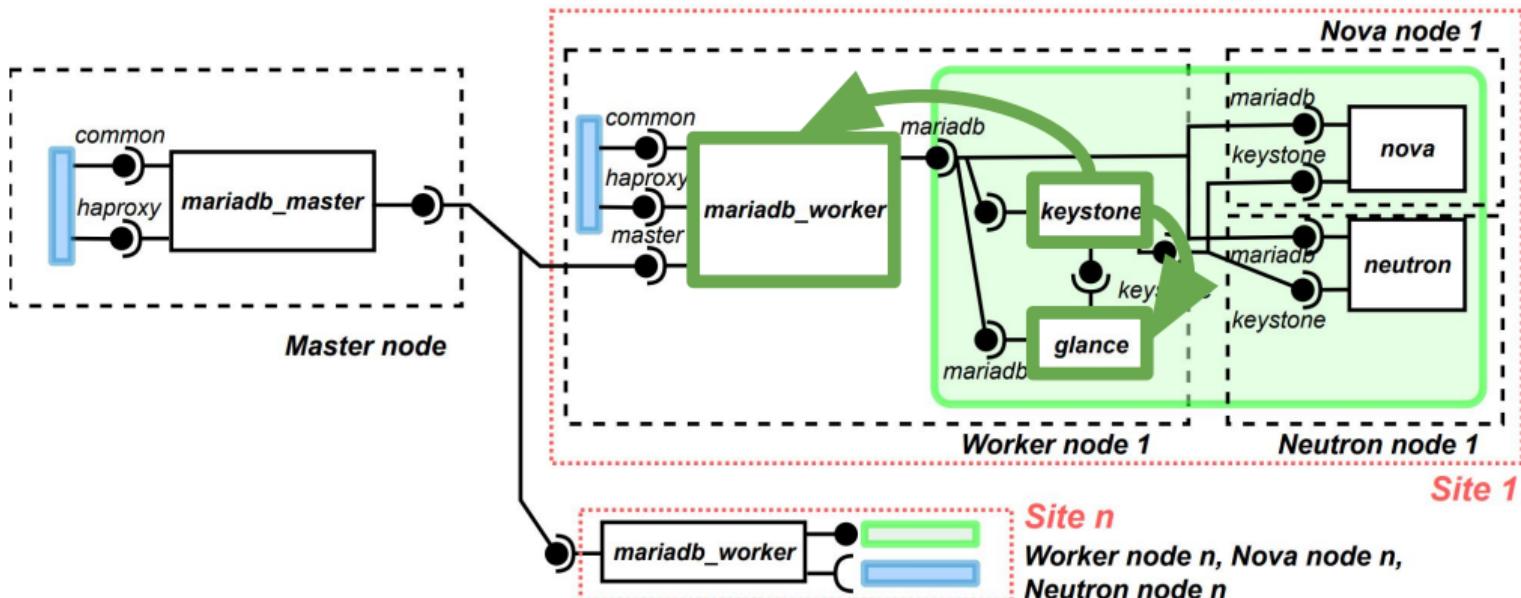
# Communication protocol



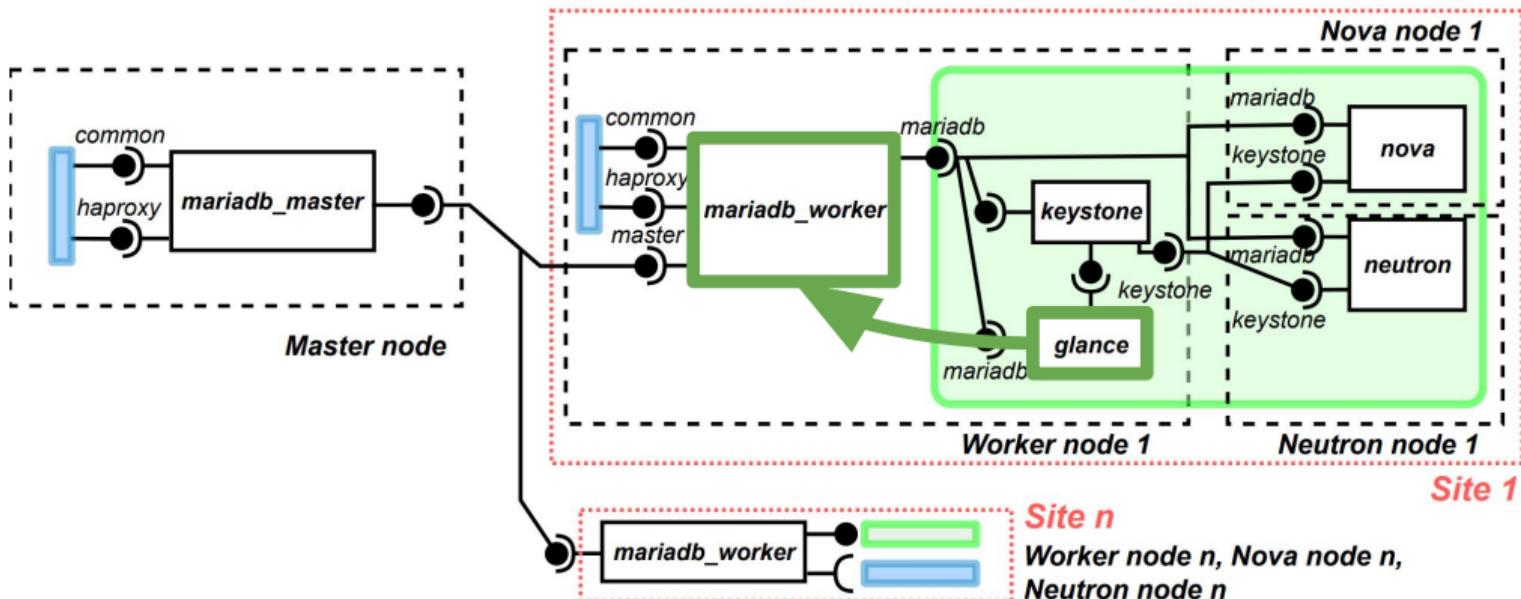
# Communication protocol



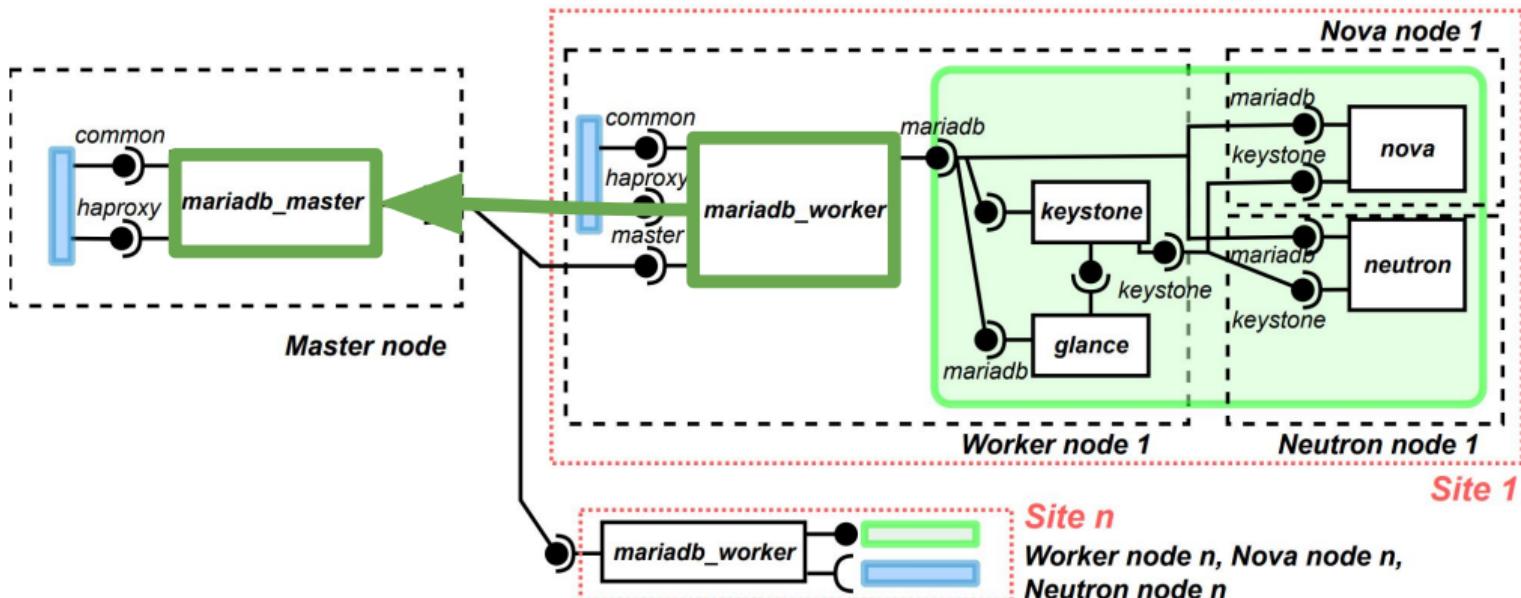
# Communication protocol



# Communication protocol



# Communication protocol



# Communication protocol

