

# **Correction in distributed systems, from usage to reconfiguration**

Progress and insights from my research journey

---

Jolan Philippe

January 24, 2025 - Séminaire STR

DiverSE team, IRISA (Univ. Rennes)

# Hello world



Jolan Philippe, Postdoc  
Université de Rennes, DiverSE Team  
 $(\lambda x.\lambda y.x@y)$  Jolan.Philippe.inria.fr

Topic of interest:

- Correctness in
- Distributed computing
  - Model driven engineering
  - Reconfiguration

More details on: <https://jolanphilippe.github.io/>

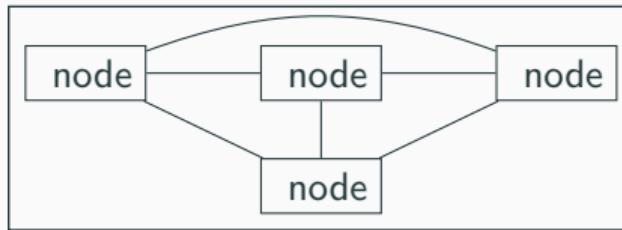
# Disclaimer



All the work presented in this talk is from different  
projects, contexts, and people.

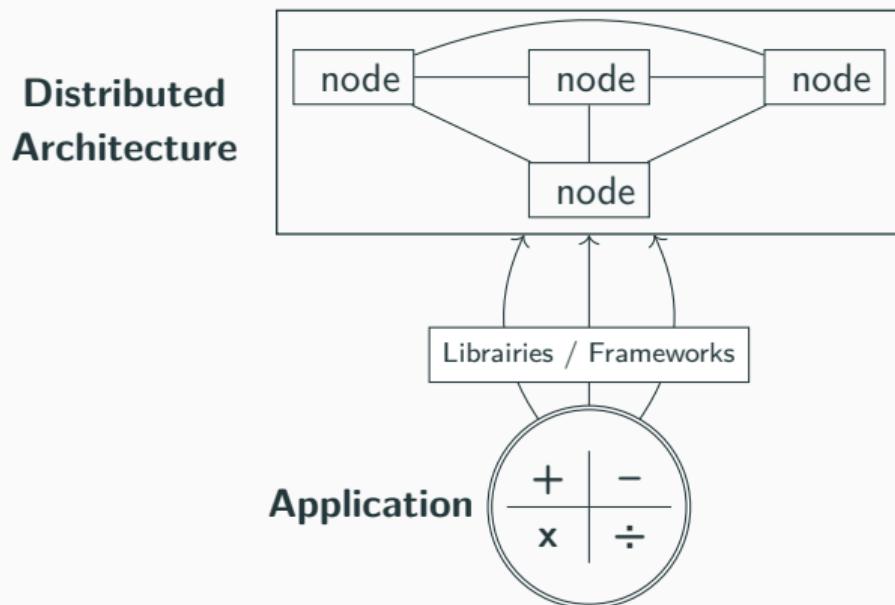
# Context: Using distributed infrastructure

## Distributed Architecture



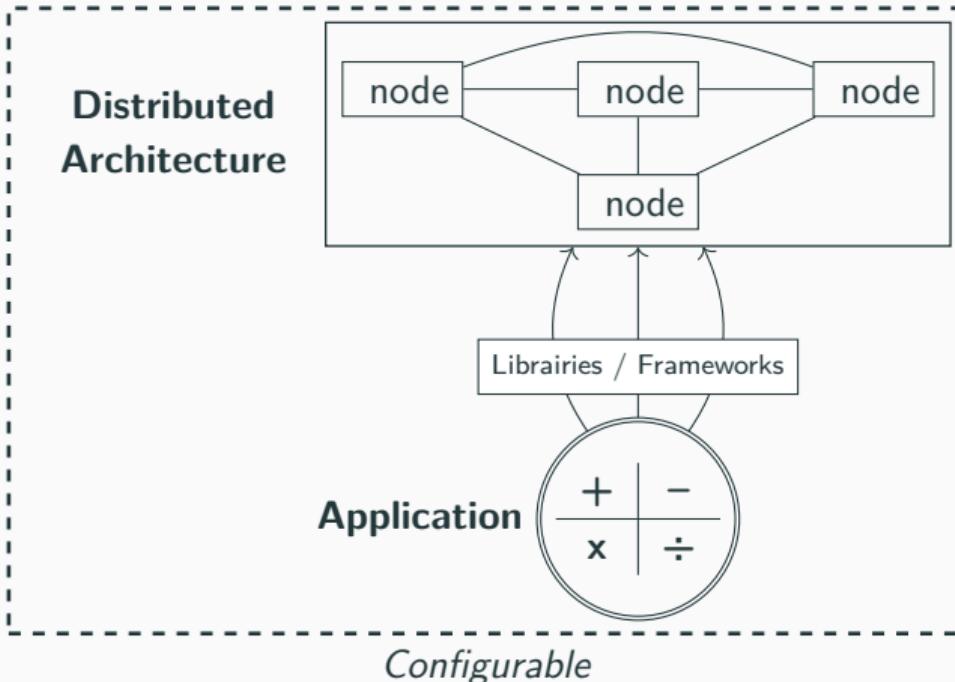
- Different topologies
  - Ring
  - Star
  - Bus
  - etc.

# Context: Using distributed infrastructure



- Different topologies
  - Ring
  - Star
  - Bus
  - etc.
- Different architecture
  - Single Instr. Single Data
  - Single Instr. Multiple Data
  - Multiple Instr. Single Data
  - Multiple Instr. Multiple Data

# Context: Using distributed infrastructure



- Configuring architecture
    - Parametrized resources
    - Services
  - Configuring application
    - Allocated resources
    - Features
- ⇒ DevOps perspective
  - Continuous Integration
  - Continuous Deployment

## Developing and managing application on distributed architecture is error-prone

### Ensuring correctness

**Goal:** Concrete application meets expectations, at different level:

- The application itself
- The used libraries / frameworks
- The reconfiguration

**Hint :** Using formal approaches

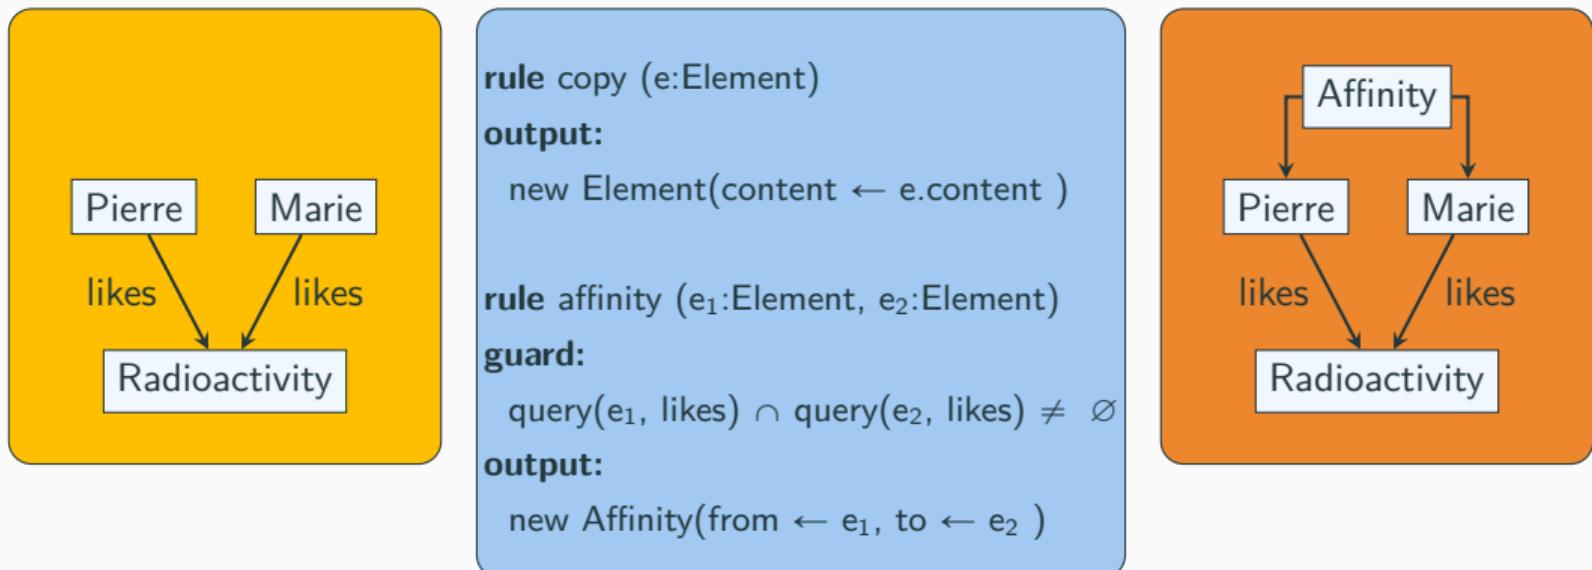
## Running example - A distributed model transformation engine



## Running example - A distributed model transformation engine



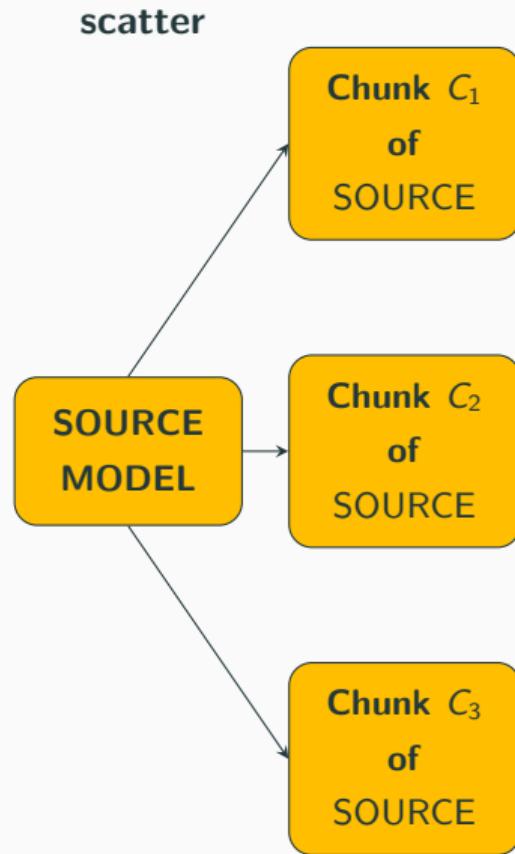
# Running example - A distributed model transformation engine



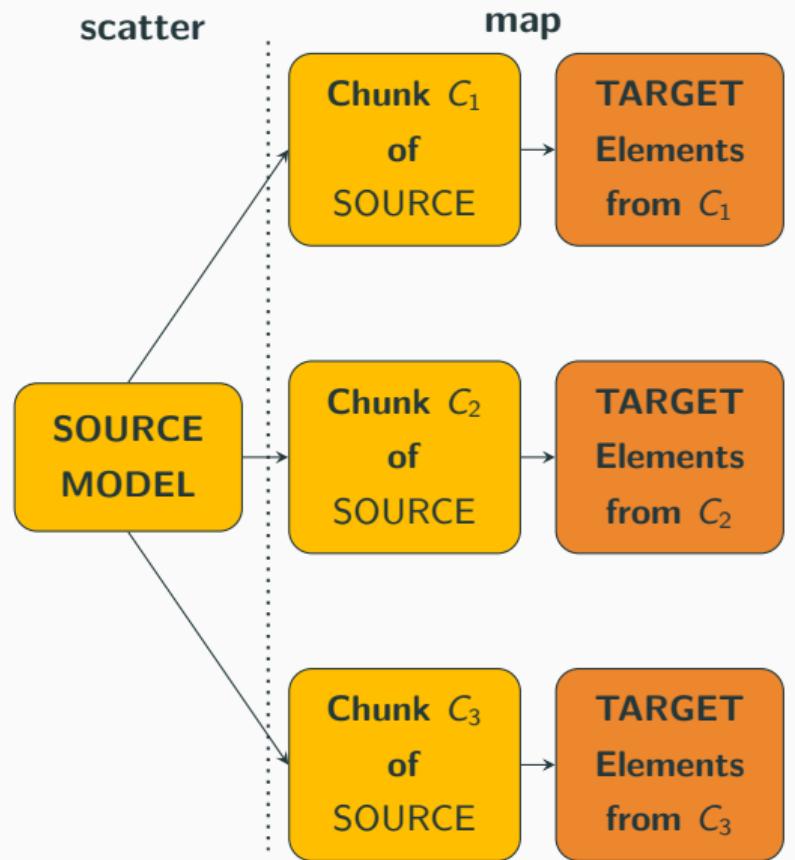
## Running example - A distributed model transformation engine

SOURCE  
MODEL

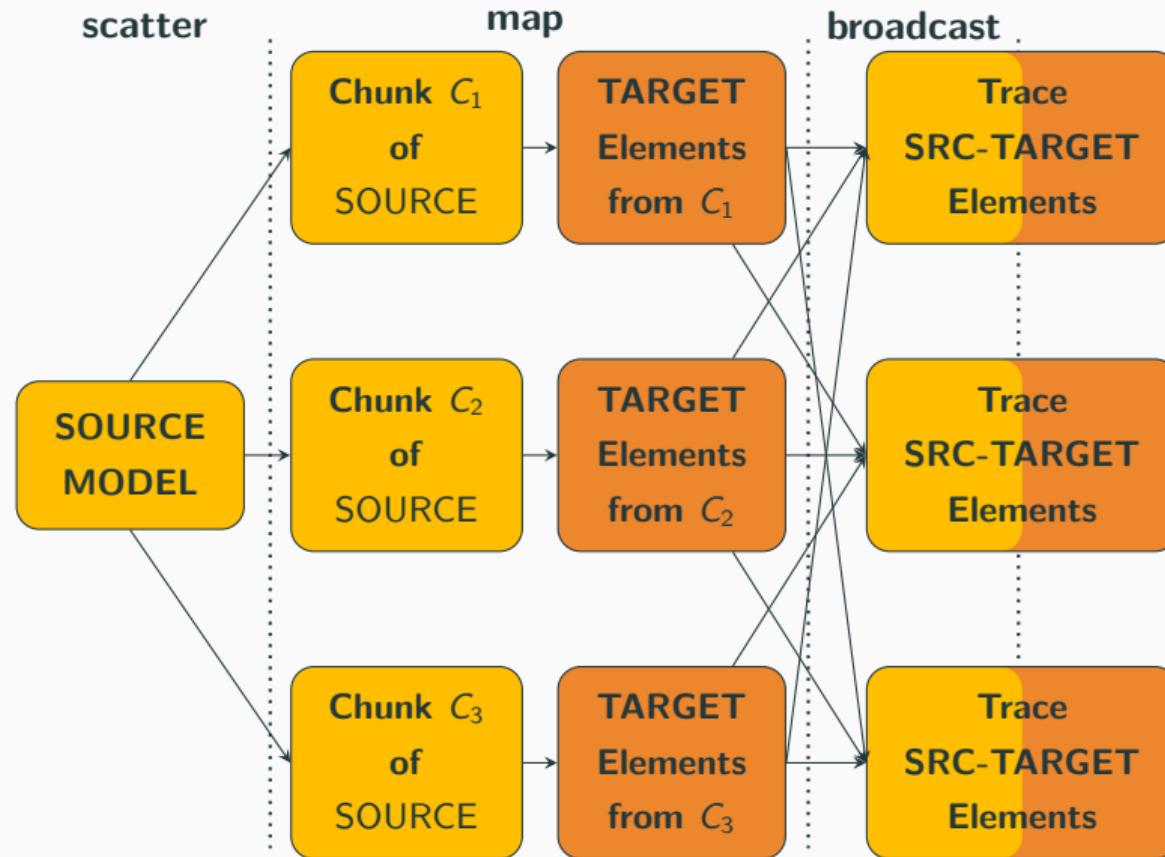
## Running example - A distributed model transformation engine



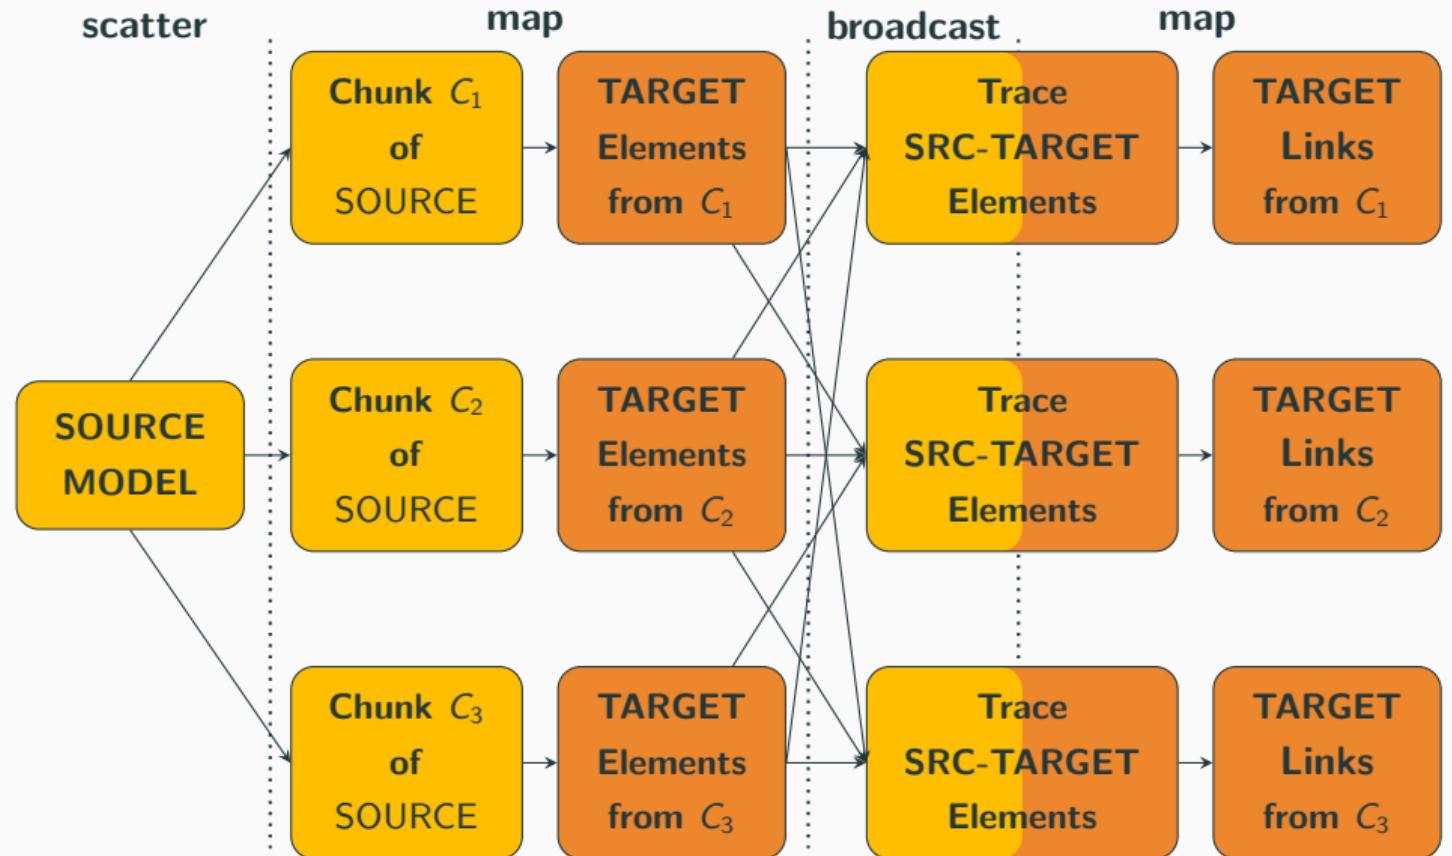
## Running example - A distributed model transformation engine



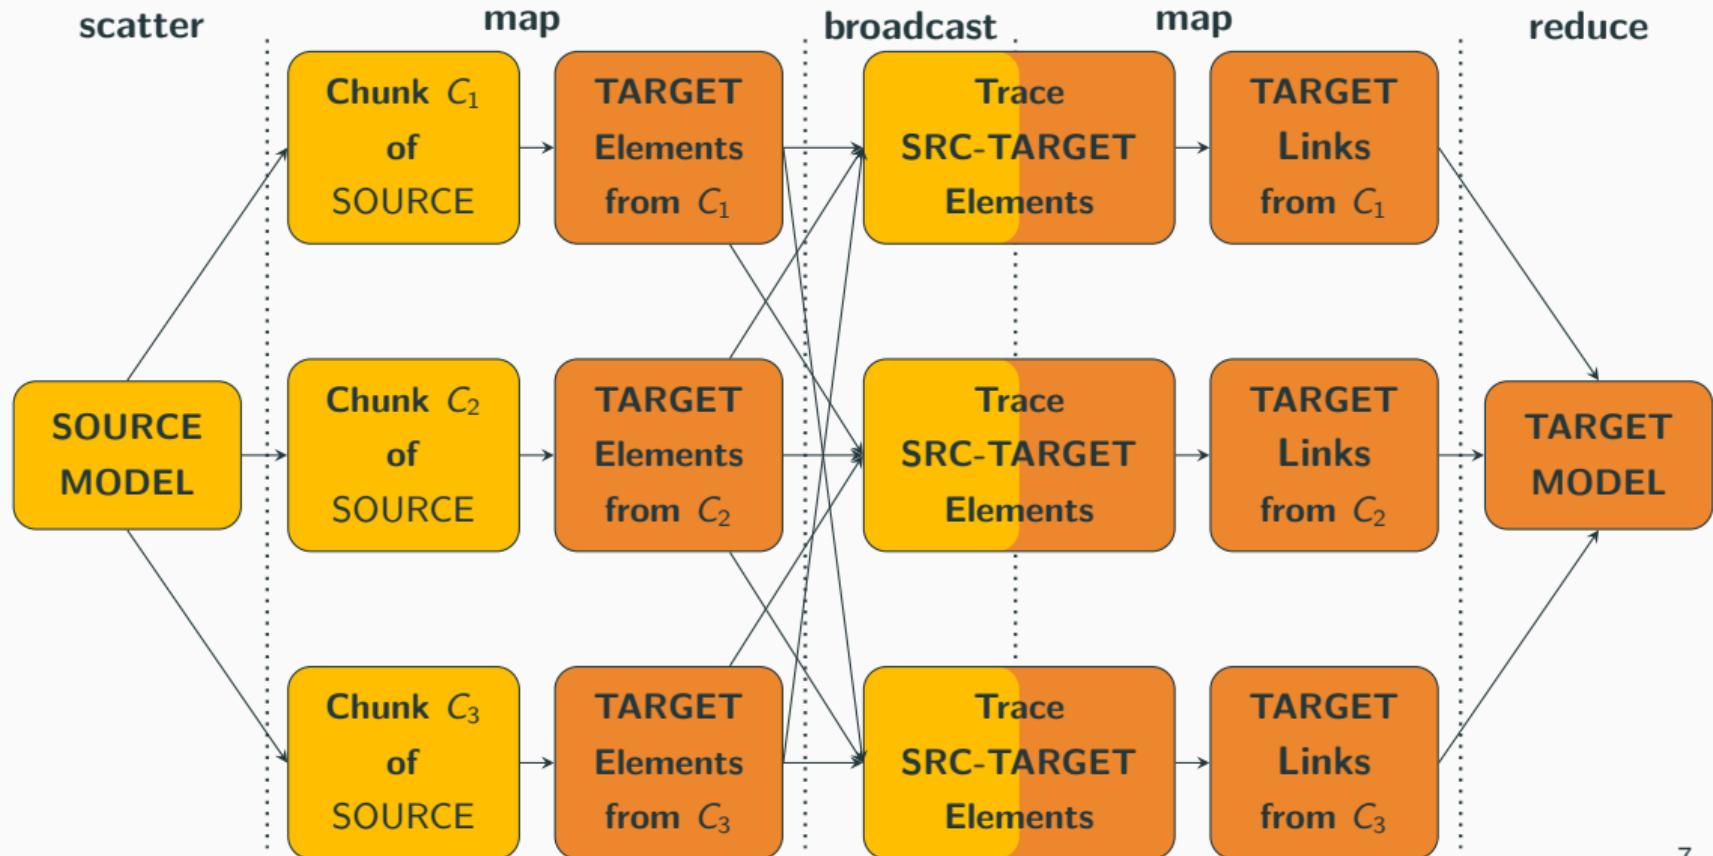
## Running example - A distributed model transformation engine



## Running example - A distributed model transformation engine



## Running example - A distributed model transformation engine



# Outline

1. **SparkTE, a correct-by-construction model transformation engine**
  - A configurable engine
  - Running correct-by-construction transformations
2. **Verifying frameworks for distributed calculation**
  - Skeletons and correctness
  - A Coq library: SyDPaCC
  - SyDPaCC for Spark
3. **Coordinated reconfiguration**
  - Complex architecture for SparkTE
  - Decentralized reconfiguration
  - Ballet for reconfiguring
  - Model-checking on Ballet

# SparkTE - A configurable engine

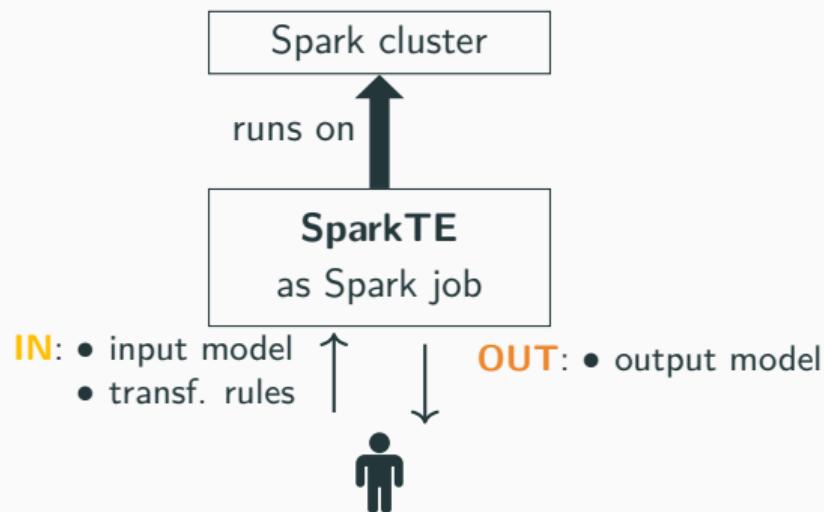
## Distributed model transformation engine

- Offers scalability
- Based on Apache Spark
  - Popular framework for large-scale data processing
  - Support for many paradigms
  - Open-source

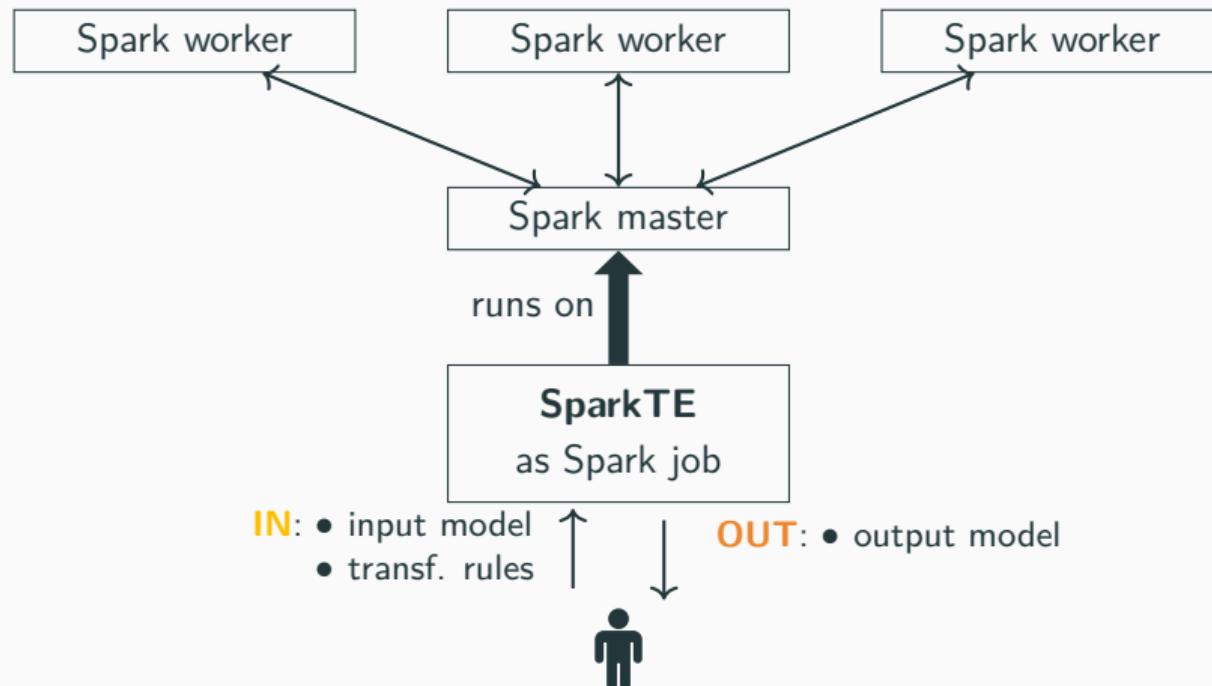
## Highly configurable transformation

- Several possible execution semantics
- Multi-paradigm approach for querying input model
- Engineering design choices configuration

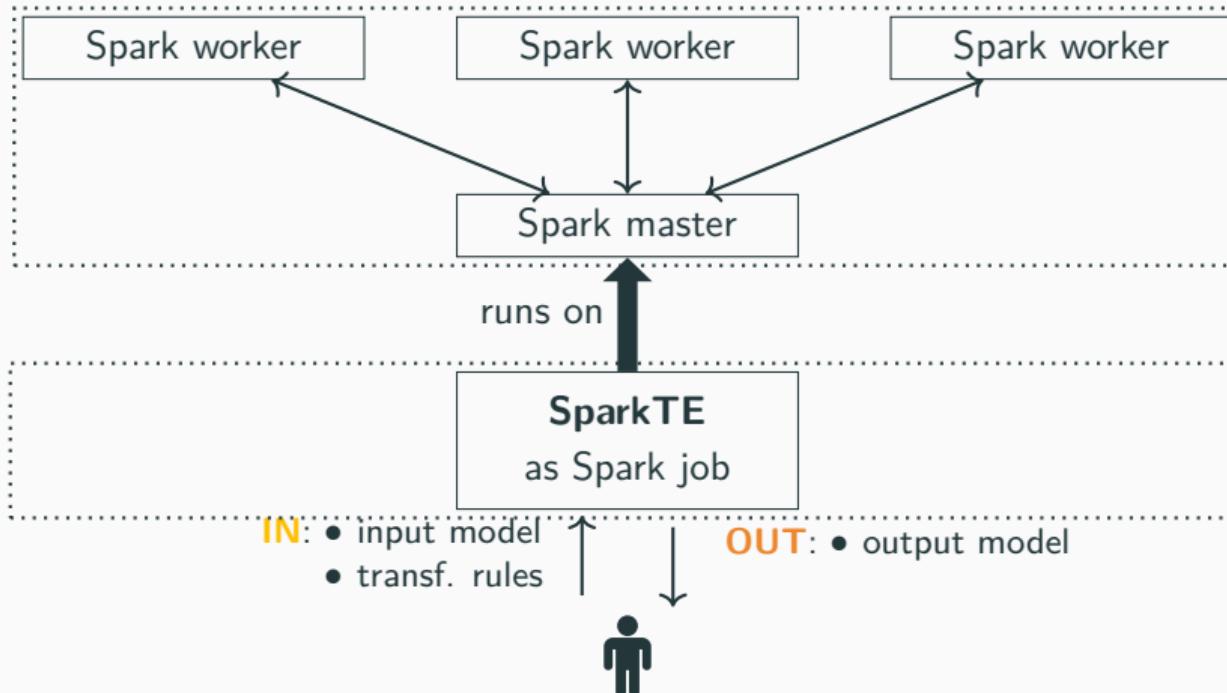
# SparkTE - A configurable engine



# SparkTE - A configurable engine



# SparkTE - A configurable engine



## Configure

- memory
- num workers
- num threads
- etc.

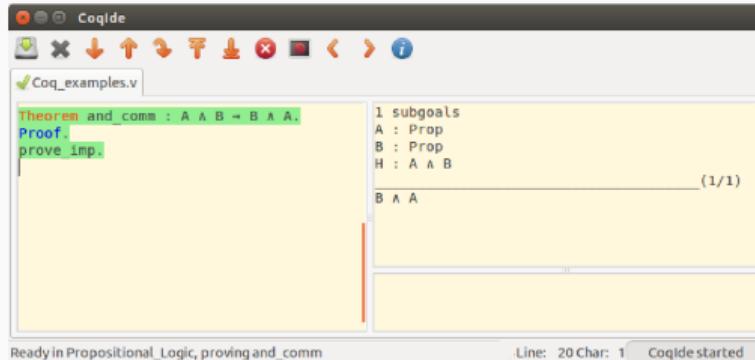
## Configure

- queries on input model
- semantics execution
- engineering choices

## General approach for reasoning on model transformation

- Formalize transformation in the proof assistant Coq
- Refine the formalization for performances
- Extract a running engine
  - ⇒ Extract spec. into Scala code
  - ⇒ Run Scala code on Spark Cluster

# Use of CoqTL for reasoning



## The Coq proof assistant

- Designed for specifying semantics
- A proof assistant based **calculus of constructions** and **Hoare logic**
- **Extraction** mechanism (to ML langs)

## CoqTL

- **DSL** for rule-based model transformation
- Made for **reasoning on model transformations**
- Can **reason on the semantic** of the transformation

A screenshot of a GitHub repository page for 'coqtl' by 'atlanmod'. The page includes a search bar, a pull requests button (12), actions, projects, and a wiki link. Below the header are three small icons. The main content area describes CoqTL as a tool for writing model transformations and proving their correctness in Coq. It includes a 'View license' link and social sharing icons. At the bottom, it shows statistics: 14 stars, 12 forks, 7 watching, 1 branch, 18 tags, and activity links.

CoqTL allows users to write model transformations and prove engine/transformation correctness in Coq

View license

14 stars 12 forks 7 watching 1 Branch 18 Tags Activity

# Correct-by-construction: Parallelizable CoqTL

## Increase parallelization

1. Two distinct phases : **instantiate** and **apply**
  - Defined as map-reduce phases
2. Iterate on rules instead of source patterns
  - Avoid unnecessary computations
3. Iterate on trace for apply instead of source patterns
  - Reuse intermediate results while everything is redefined in CoqTL

	spec (loc)	cert (loc)	effort (man-days)
1.	69	484	10
2.	42	487	7
3.	69	520	4



# Correct-by-construction: Build SparkTE

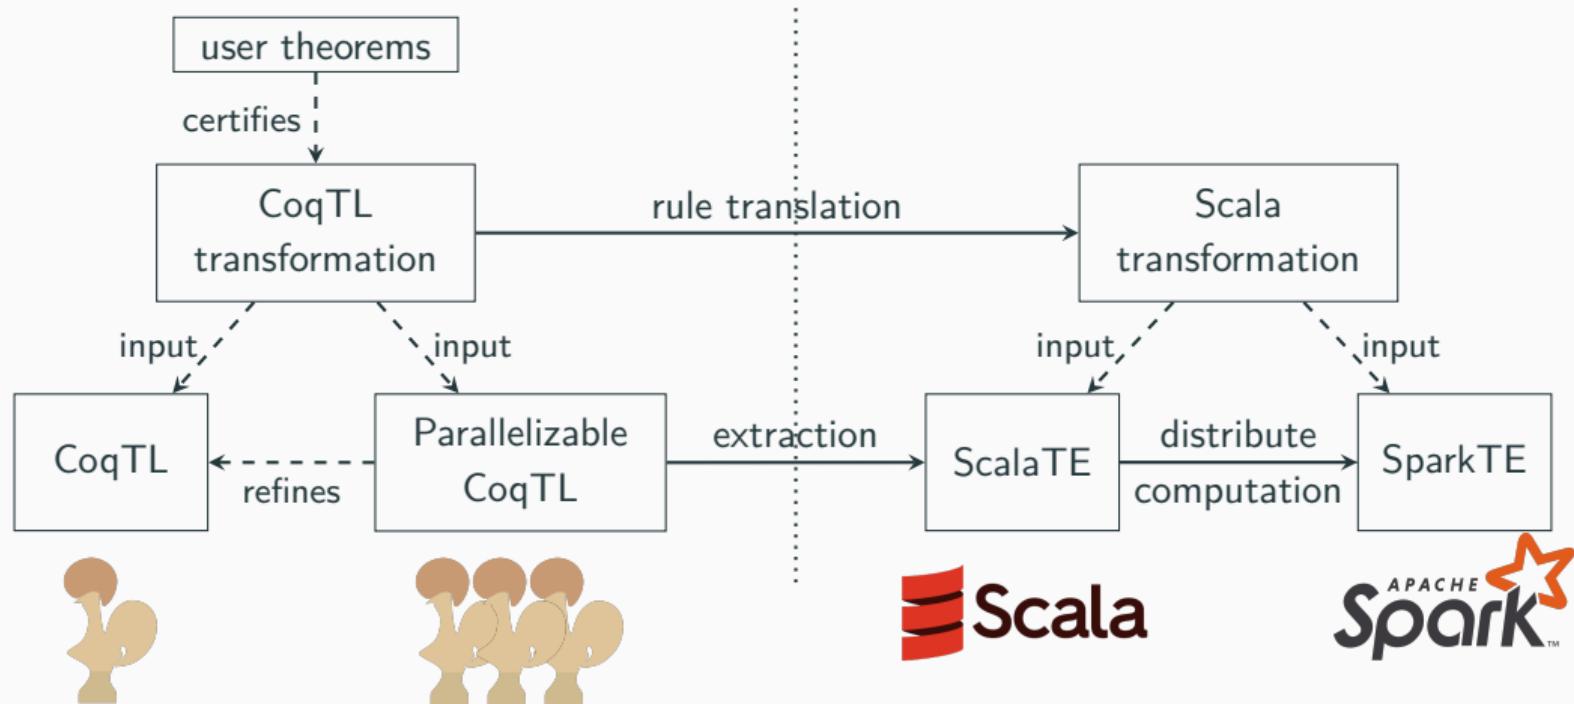


## CoqTL to SparkTE

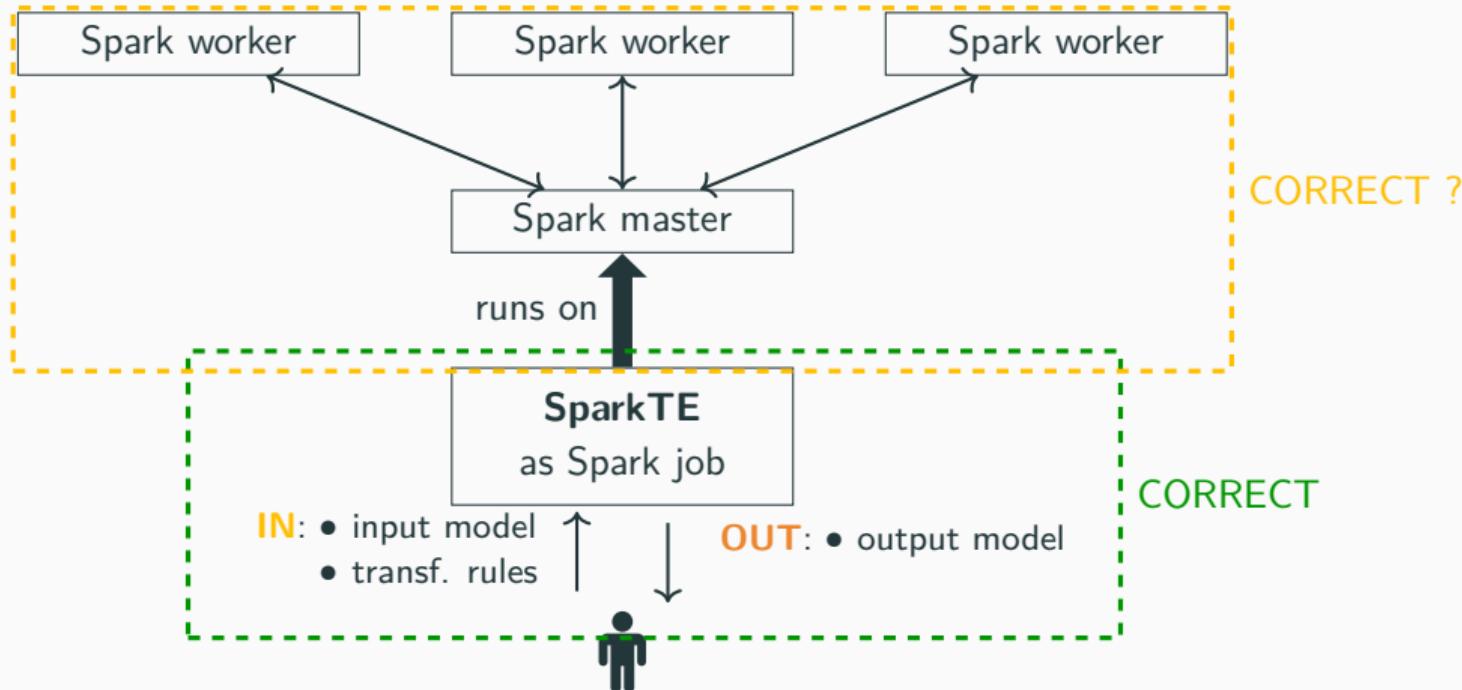
1. Produce executable and maintainable code
  - By hand: Object-oriented approach, with pure Scala functions
  - With Scallina: Not maintainable, but certified
2. Distribute the computation
  - Distribute data-structure
  - Explicit communication operations (scatter, broadcast and reduce)



# Correct-by-construction



# Correctness... really ?



## Certify Spark

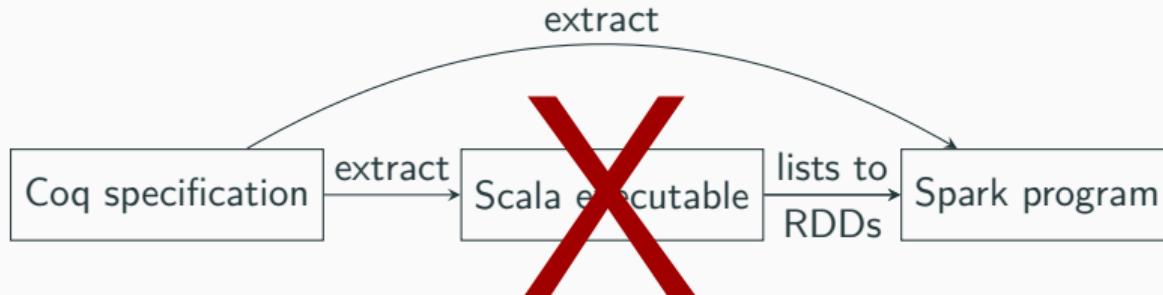
Spark defines program as

- Usage of high-order functions (e.g., *map*, *reduce*)
- Using a distributed implementation (i.e., skeletons)
- Considering the sequential implementation equivalent to distributed one

```
val instantiatedElements =  
    transformation.rules.map { rule => instantiate(model.elements, rule) }
```

**No guarantee that the parallel implementation behaves the same as the sequential implementation.**

# Coq to Spark



## Extract Coq code into Spark program

- Formalize Spark's distributed structure (i.e., RDD) in Coq
- Formalize computation on RDDs
- Prove the equivalence between function on lists and on RDDs

## SyDPaCC

- Coq library for writing data-parallel program specification
- Code can be extracted into BSML (BSP for Ocaml)
- Ensure the correctness of the extracted parallel program
- Based on type equivalences (with composition)

$$\begin{array}{ccc} A_p & \xrightarrow{f_p} & B_p \\ \downarrow join_A & & \downarrow join_B \\ A & \xrightarrow{f} & B \end{array}$$

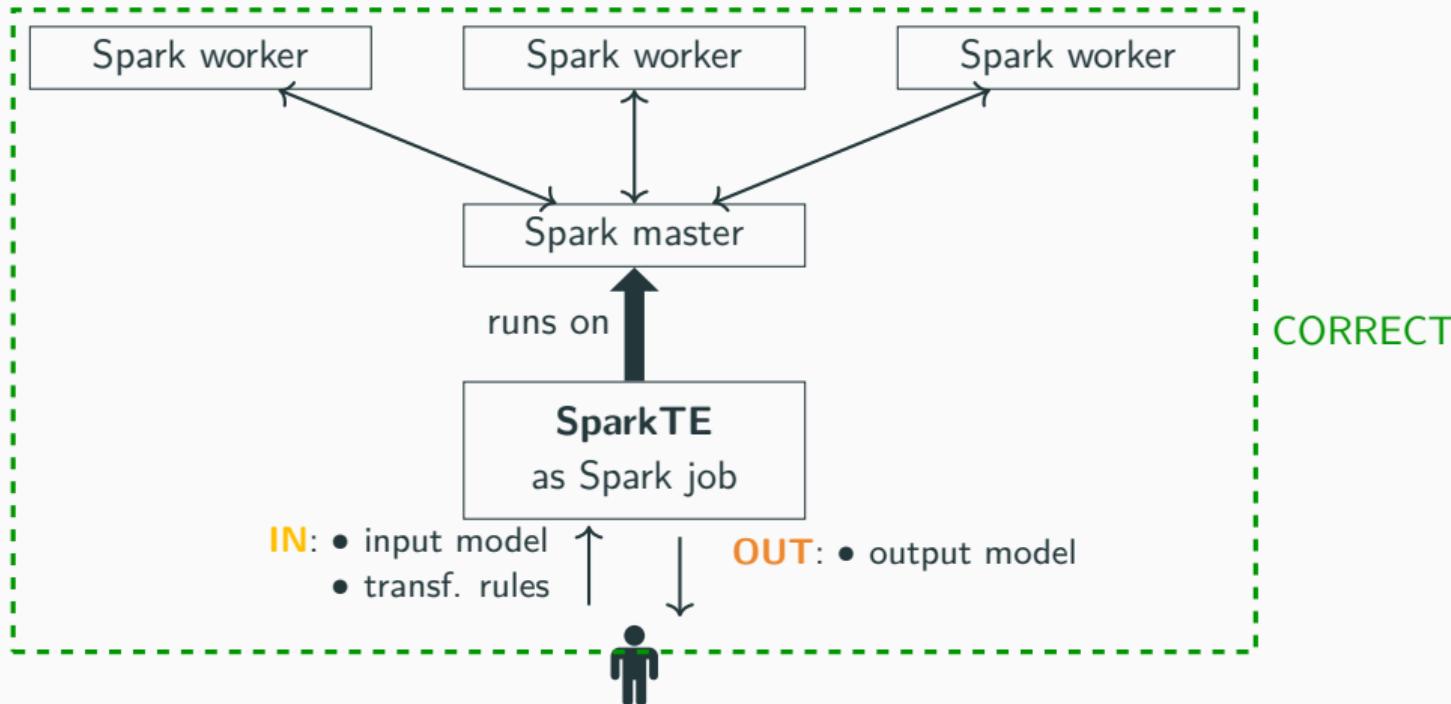
## Example - Equivalence of map



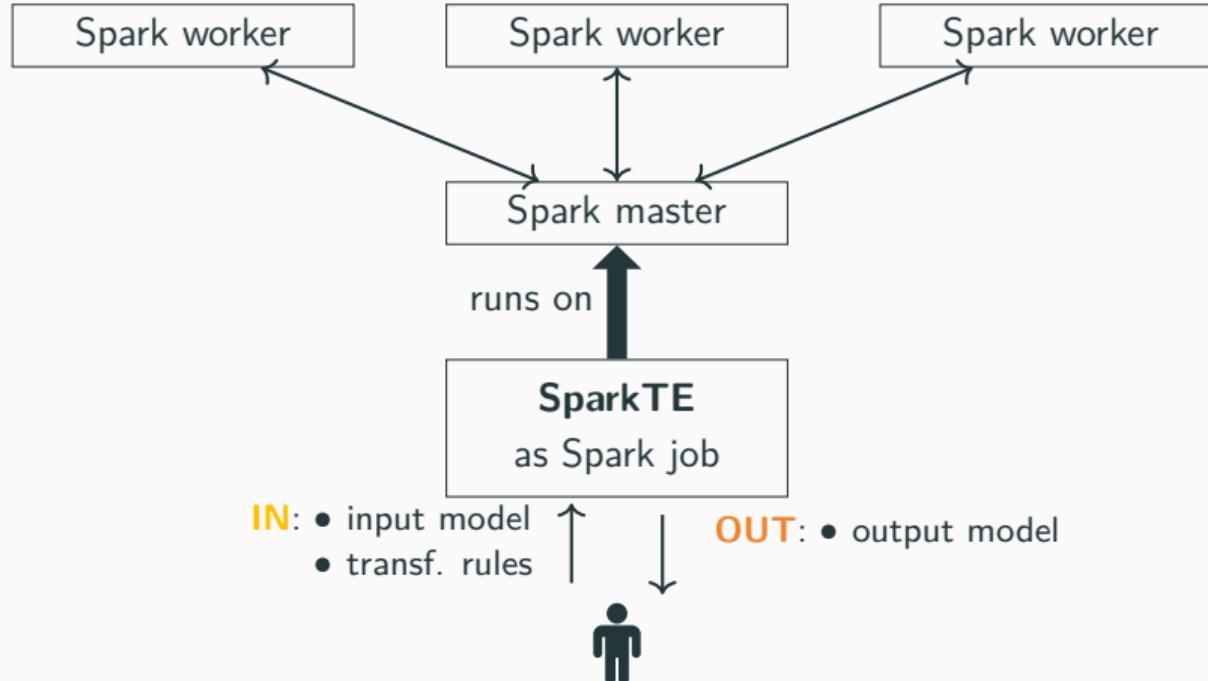
### Extending SyDPaCC

- Formalize RDDs
- Additional proofs
  - $RDD.map \circ collect = collect \circ List.map$
  - Surjectivity of  $collect$

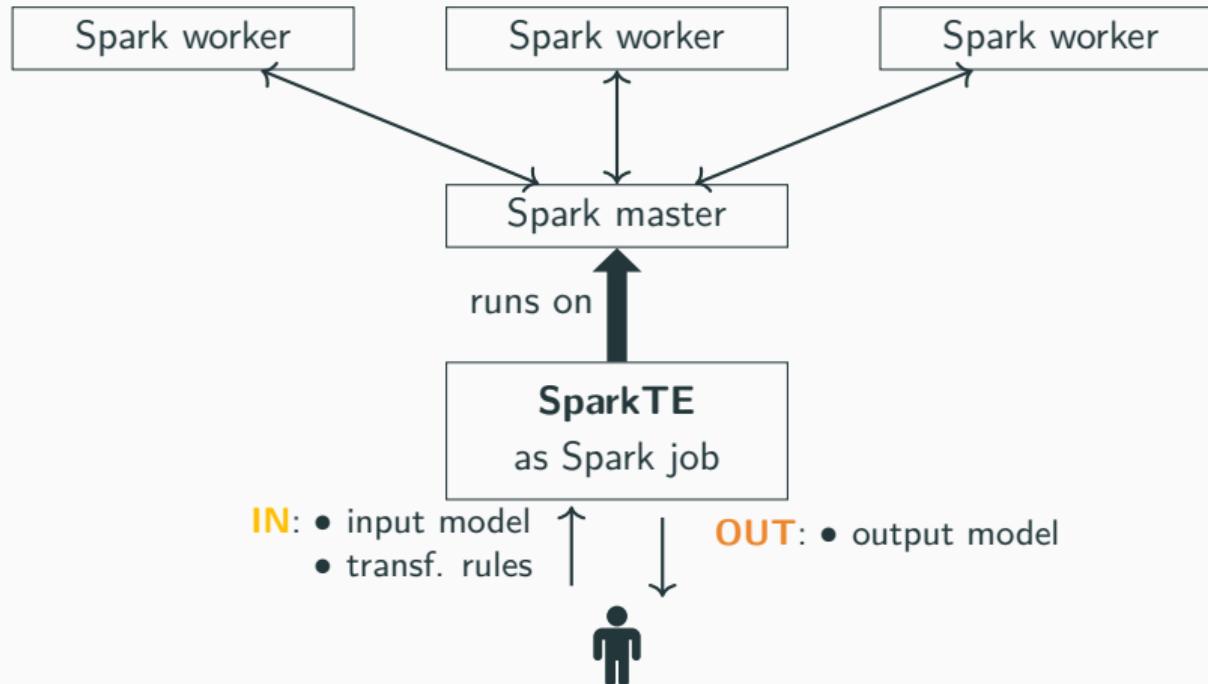
# Towards verified parallel computing with Coq and Spark



# SparkTE, as a “prototype”, architecture



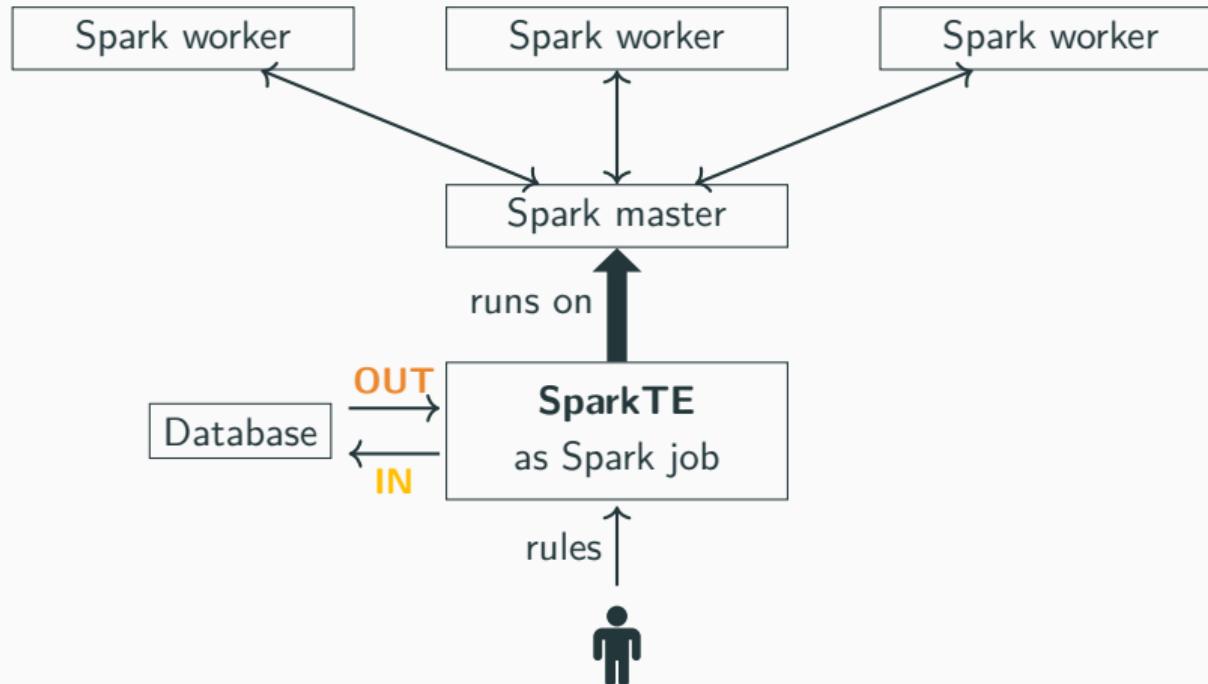
# SparkTE, as a “production tool”, architecture



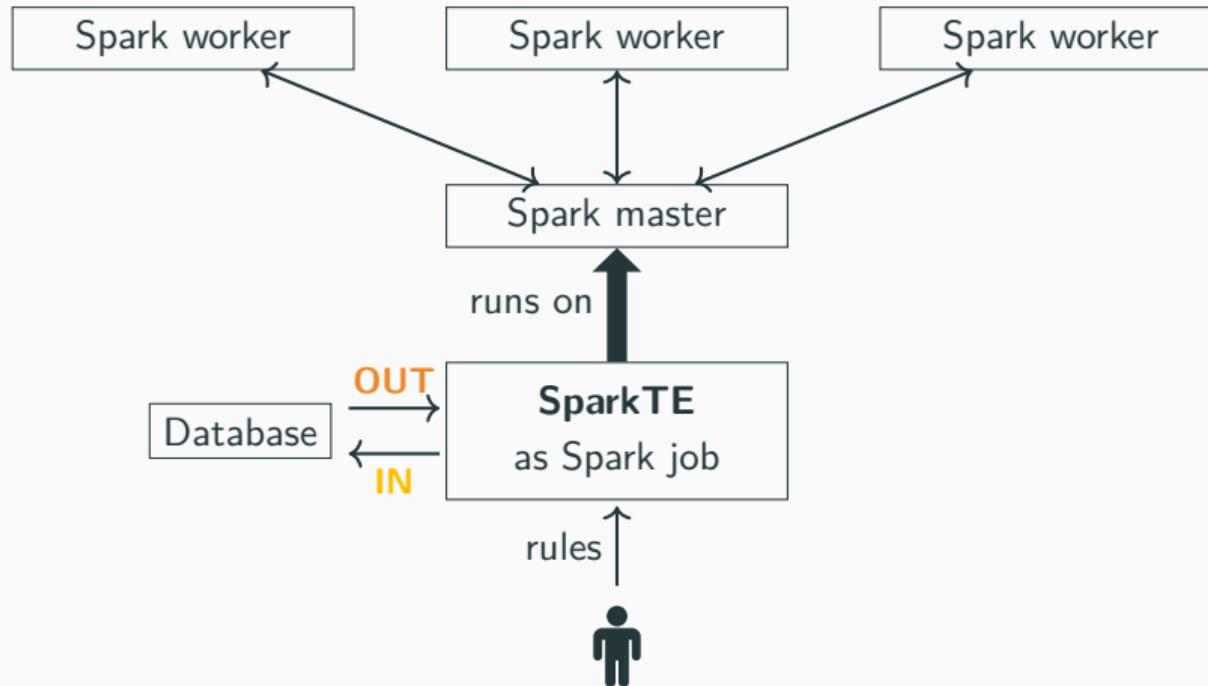
We want to store the model within a database

NB: We could use HDFS files alongside Apache Hive

# SparkTE, as a “production tool”, architecture

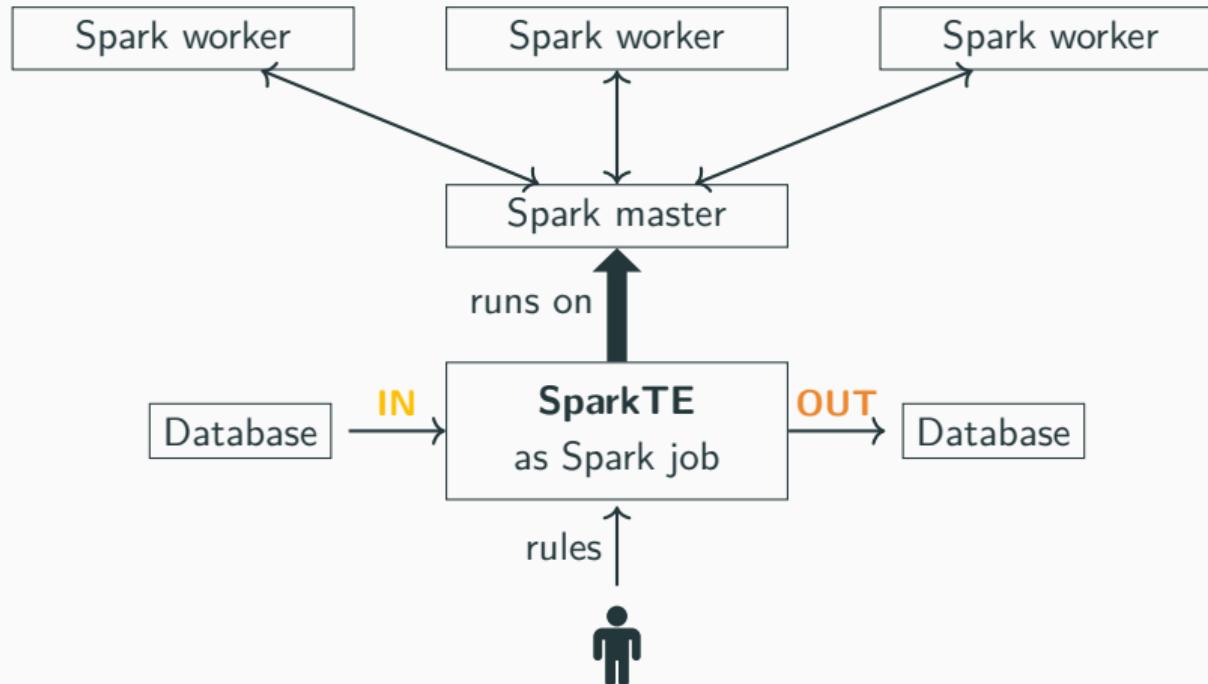


# SparkTE, as a “production tool”, architecture

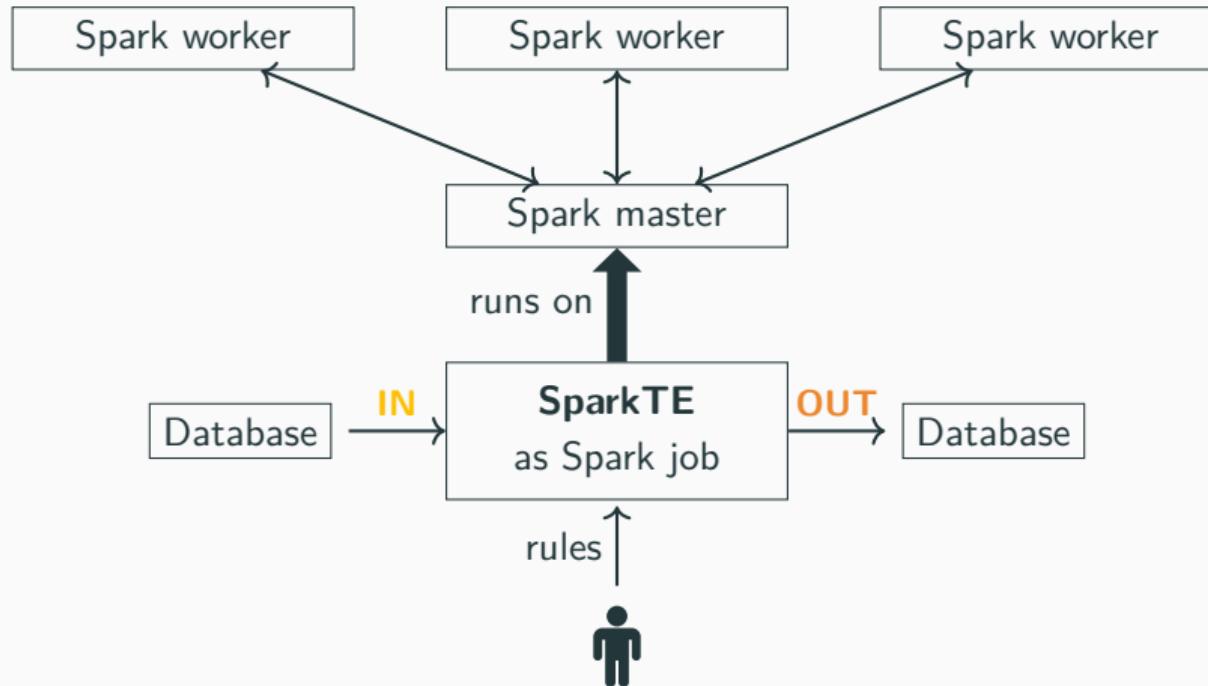


Since the transformation is not necessarily “in-place”  
we might want to store the input model and output model in different databases

# SparkTE, as a “production tool”, architecture

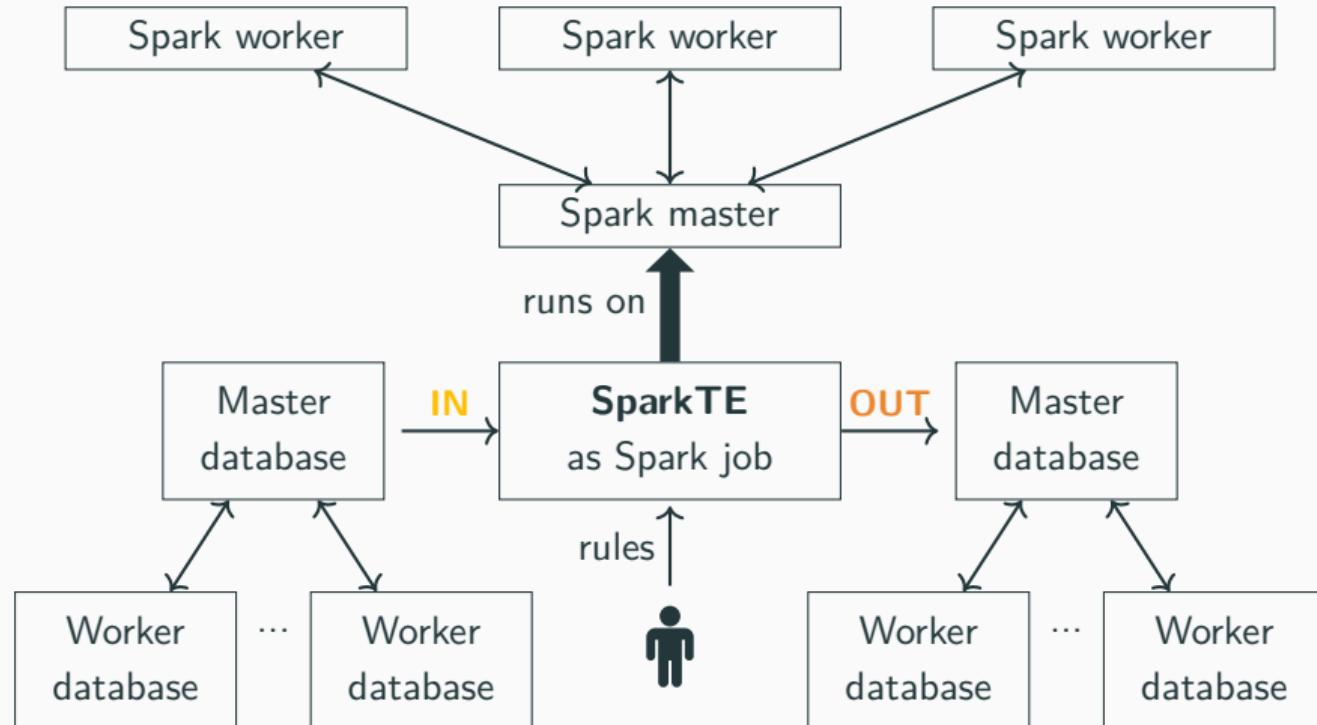


# SparkTE, as a “production tool”, architecture

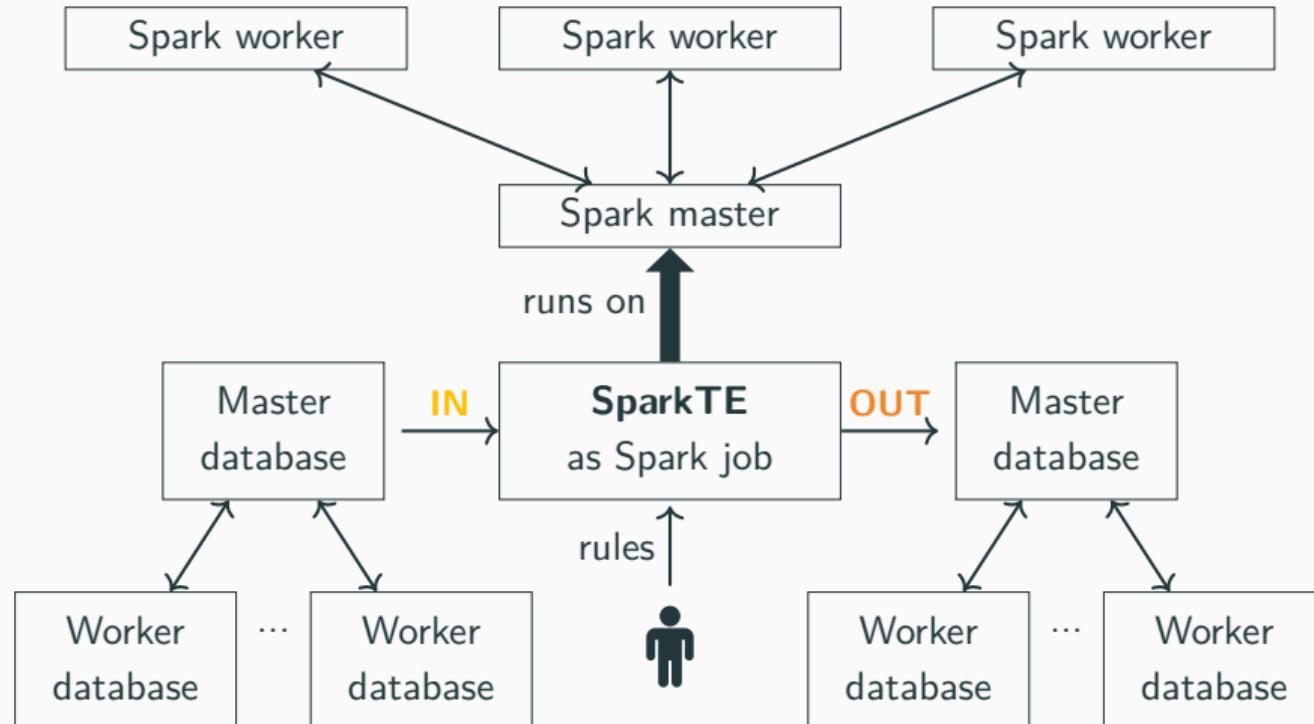


Because we handle very large model, we want distributed databases

# SparkTE, as a “production tool”, architecture

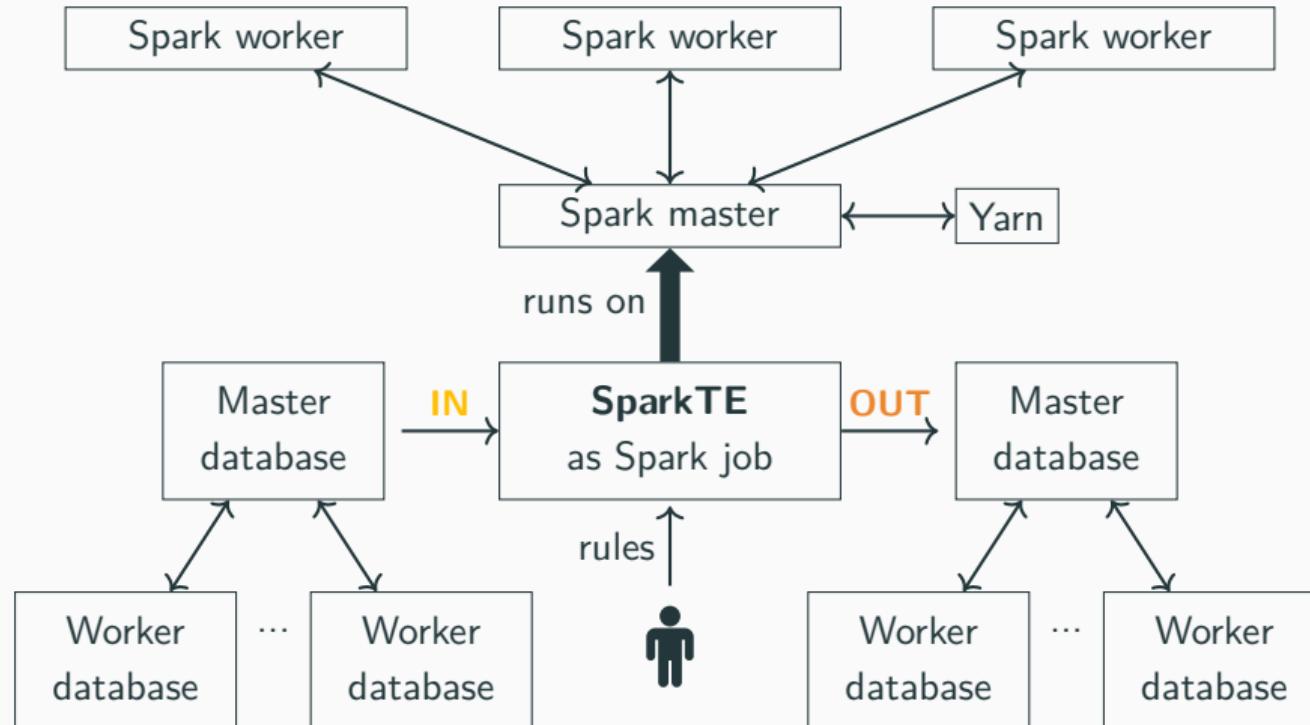


# SparkTE, as a “production tool”, architecture



We want more fine-grained management of Spark resources using Yarn

# SparkTE, as a “production tool”, architecture

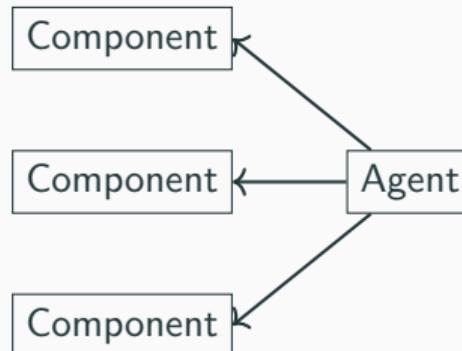


# Deploy and reconfigure SparkTE

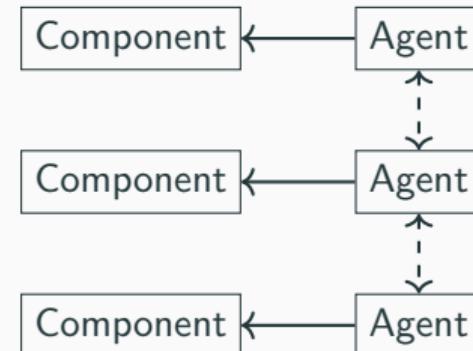
## Two approaches for reconfiguring distributed systems

**Reconfiguration** : Change of the state of entities by applying operations (e.g., deploy, update, destroy)

- Centralized: single agent manages the reconfiguration with control components
- Decentralized: several agents manage the reconfiguration with control components



Centralized approach



(Fully) Decentralized approach

# Decentralized approach

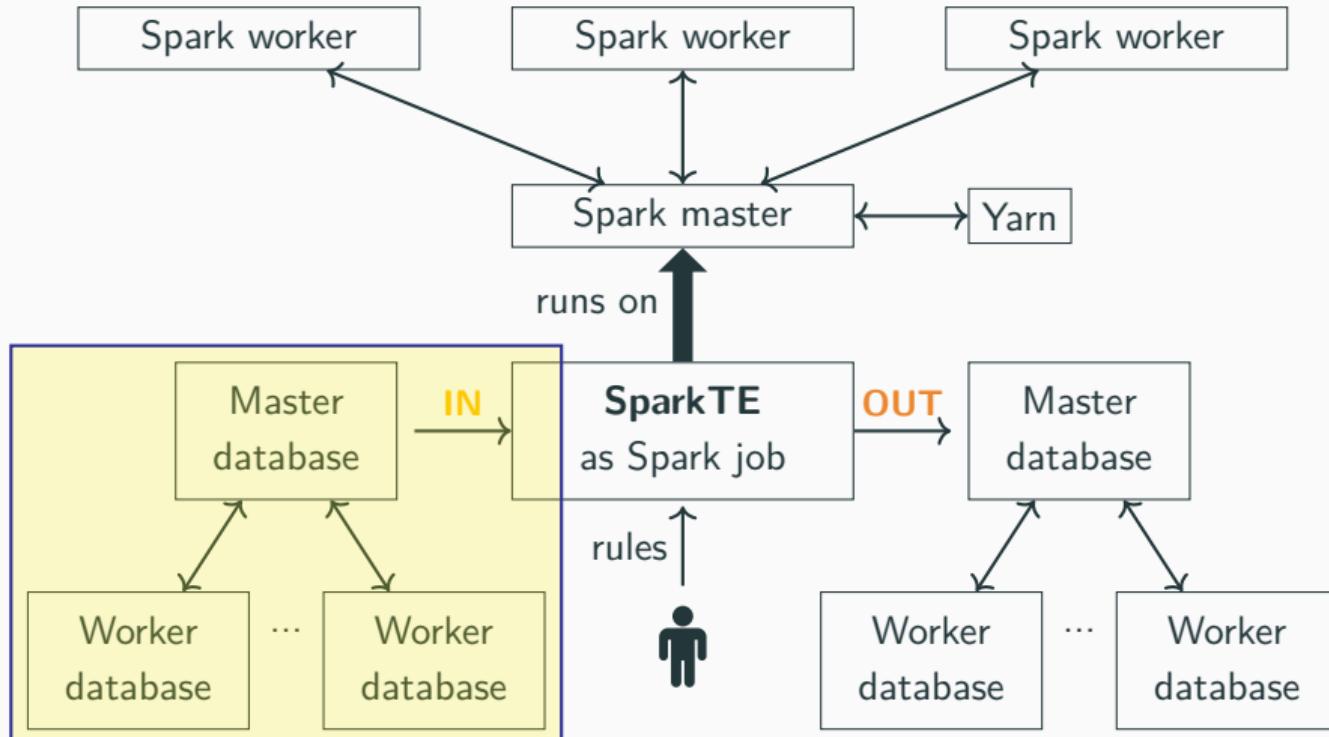
## Strength

- Not a single point of failure
- Separation of information
- Scalability
- Allow geo-distribution

## Challenges

- All agents must coordinate
- Operate communications

# Reconfiguring distributed databases



How to reconfigure ?

Let's have a look

## Example: Deploy distributed databases



### Database Master (DM) plan

1. Configure the service
2. Bootstrap the database
3. Start the service
4. Expose API

### Database Worker<sub>i</sub> (DW<sub>i</sub>) plan

1. Configure the service
2. Register to master
3. Bootstrap the database
4. Start the service
5. Expose API

- **Component granularity:** DM << DW<sub>i</sub>
- **Lifecycle granularity:** DM(4) << DW<sub>i</sub>(2) (partial order)

## Example: Update distributed databases



### Database Master (DM) plan

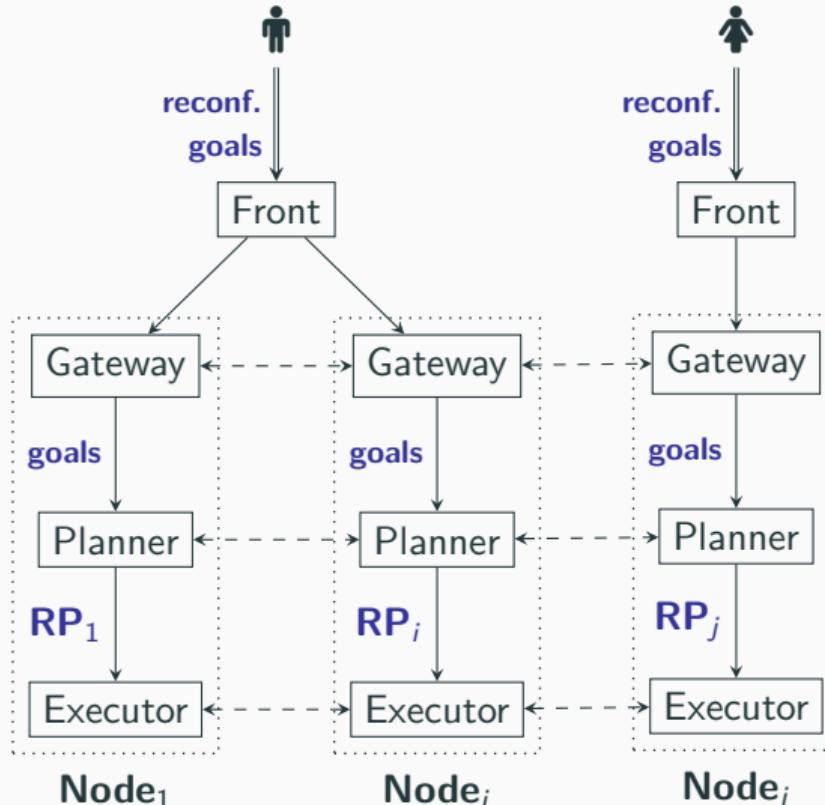
1. Interrupt the service
2. Make the update
3. Start the service
4. Expose API

### Database Worker<sub>i</sub> (DW<sub>i</sub>) plan

1. Interrupt the service
2. Make the update
3. Register to master
4. Start the service
5. Expose API

- **Component granularity?** Destroy DW<sub>i</sub> << Update DM << Deploy DW<sub>i</sub>
- **Lifecycle granularity:** DW<sub>i</sub>(1) << DM(1) & DM(4) << DW<sub>i</sub>(3)

# Ballet for decentralized reconfiguration



- Decentralized tool (one instance of Ballet on each node)
- Declarative input as goals
- Reconfiguration with automatic planning and efficient execution

## Gateway

Global knowledge building of reconfiguration goals

## Planner

Decentralized inference of reconfiguration plans (RPs)

## Executor

Coordinated execution of RP

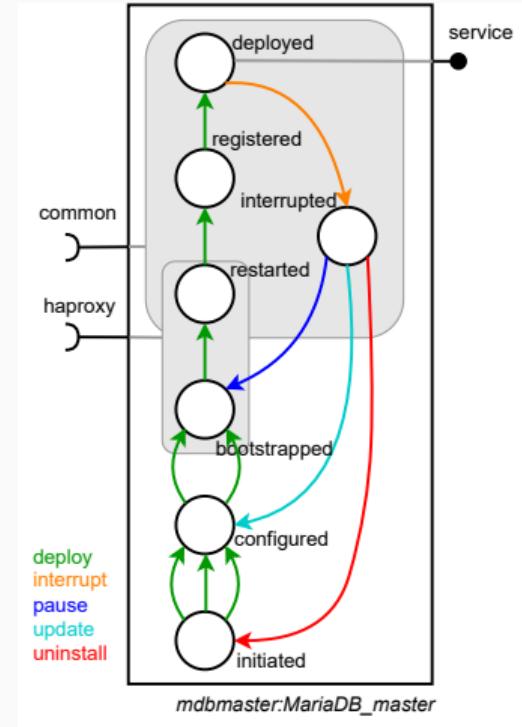
# Specify lifecycle and dependencies for control components

## Lifecycle and dependencies

- **Places:** milestones of the reconfiguration
- **Behaviors:** interface for executable actions
- **Transitions:** concrete actions between places, associated to behaviors
- **Ports:** Provide (resp. use) information to (resp. from) external components
  - Ports are bounded to places and transitions

## Exemple: Database for SparkTE

Lifecycle representation of a MariaDB database with 5 executable behaviors: **deploy**, **interrupt**, **pause**, **update**, and **uninstall**



# Usage of Ballet : goals

## Reconfiguration goals

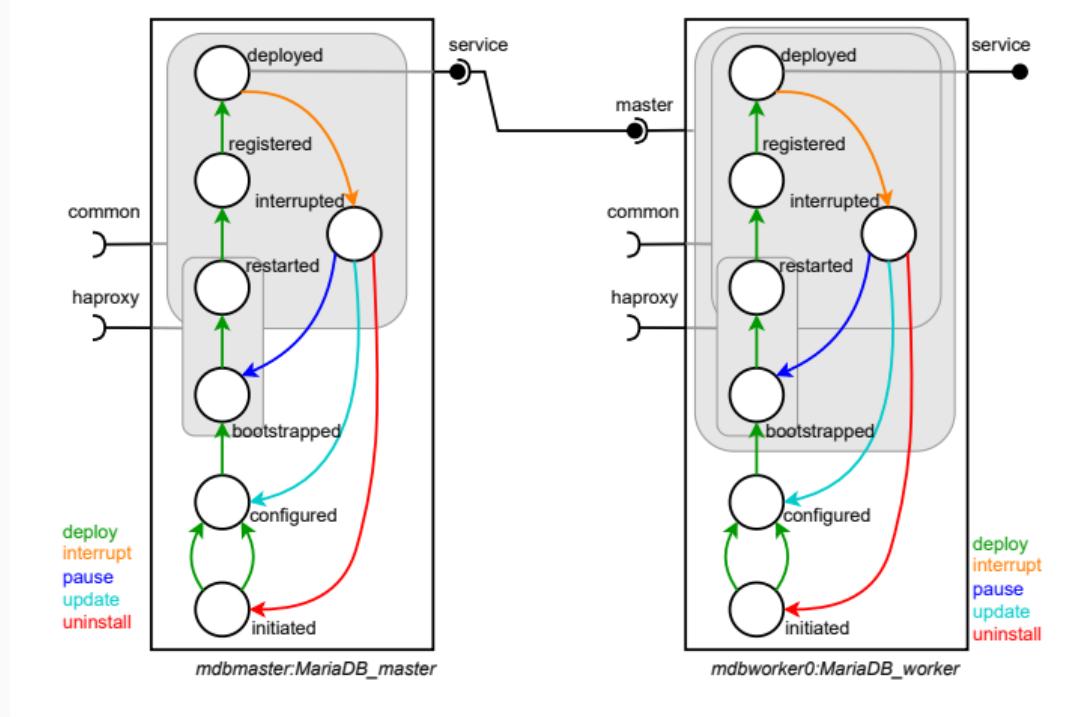
Declarative language for defining reconfiguration goals

- **Behavior goal:** Specify a behavior that must be executed
- **Port goal:** Specify a port status (active, inactive)
- **State goal:** Specify a component state (specific, running, initial)

**Listing 1:** Language to define reconfiguration goals for DevOps usage

```
<goals> ::= behaviors: <bhvr_list>
           ports: <port_list>
           components: <comp_list>
<bhvr_list> ::= ...
<bhvr_item> ::= - forall: <bhvr_name>
                  | - component: <comp_name>
                    behavior: <bhvr_name>
<port_list> ::= ...
<port_item> ::= - forall: <port_status>
                  | - component: <comp_name>
                    port: <port_name>
                    status: <port_status>
<comp_list> ::= ...
<comp_item> ::= - forall: <comp_status>
                  | - component: <comp_name>
                    status: <comp_status>
```

# Assembly of components



Assembly of MariaDB master and worker components  
Similar to synchronous Petri nets

## A simple language to interact with components - i.e., write a reconfiguration plan

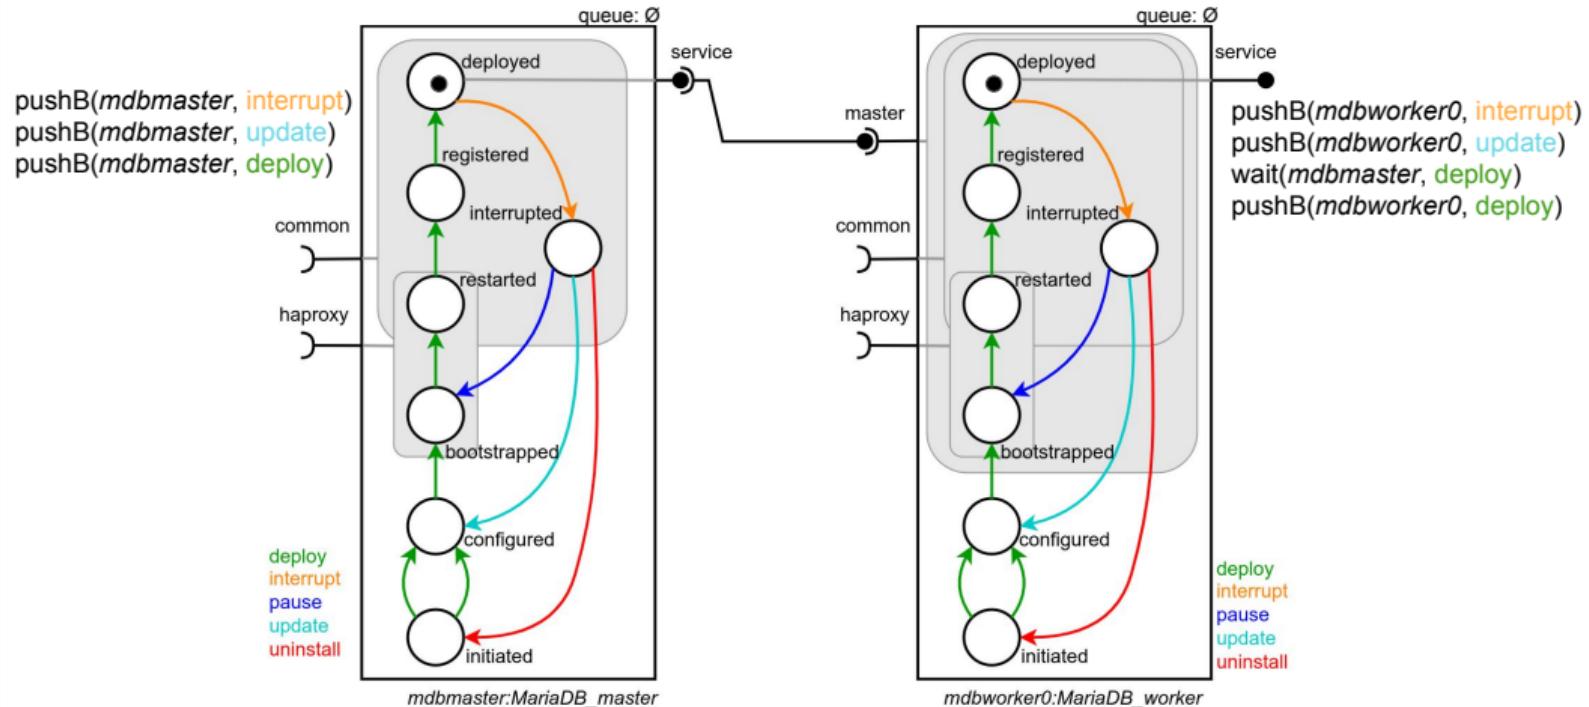
- **Add/remove** a component instance to the current assembly
- **Connect/disconnect** two component instances with compatible ports
- **Push behavior** to the behavior queue on a component instance

$\text{pushB}(id_C, bhv)$

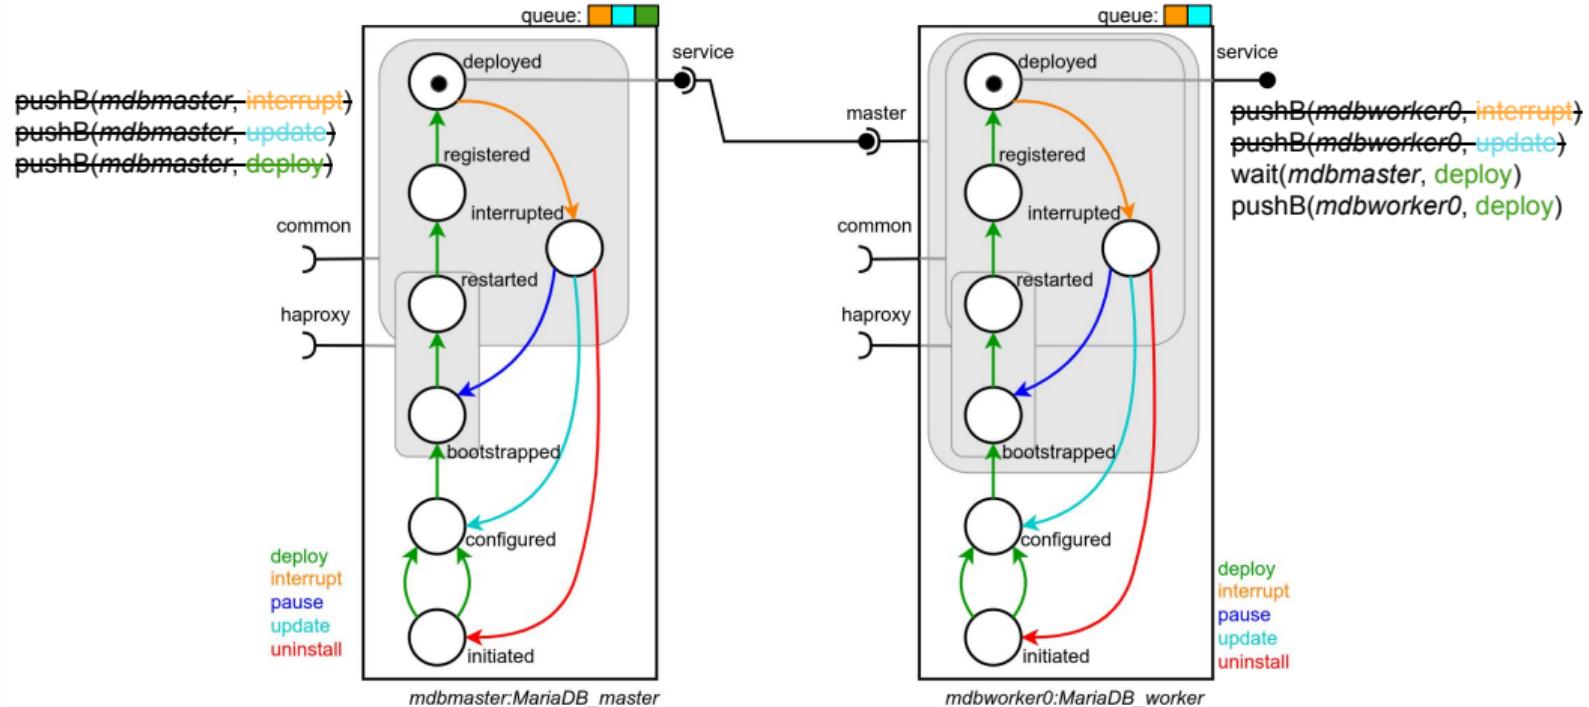
- **Wait** for a given component instance to execute a behavior

$\text{wait}(id_C, bhv)$

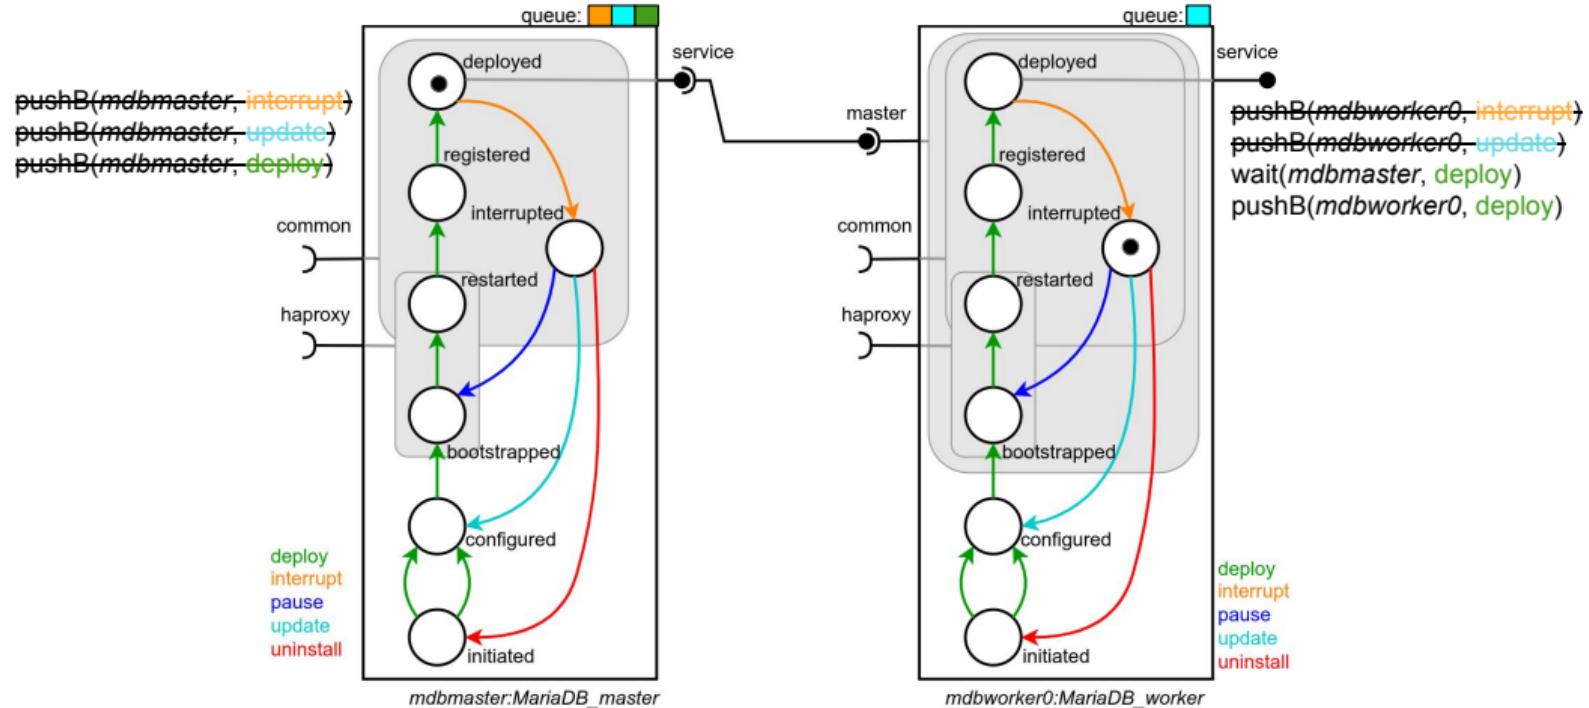
# Execution example



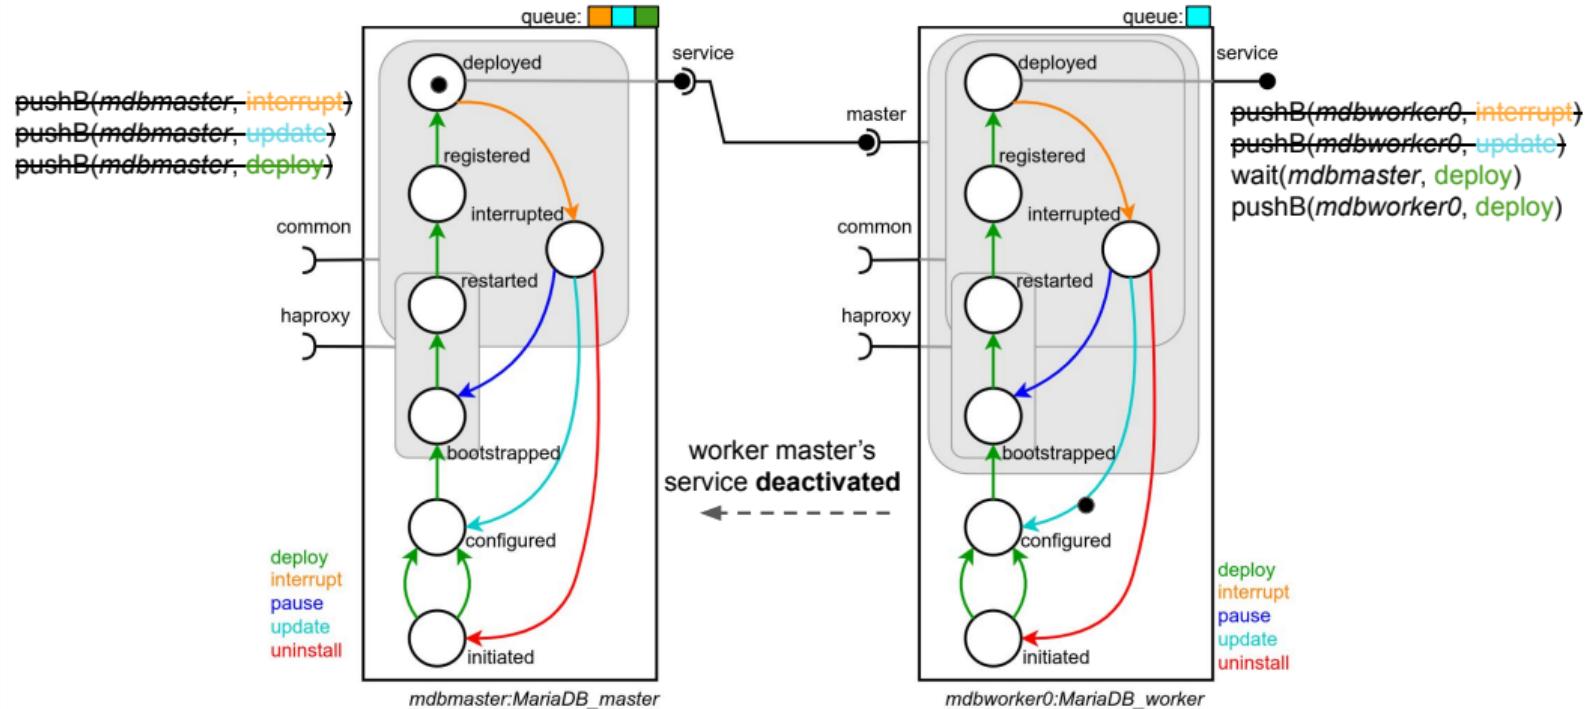
# Execution example



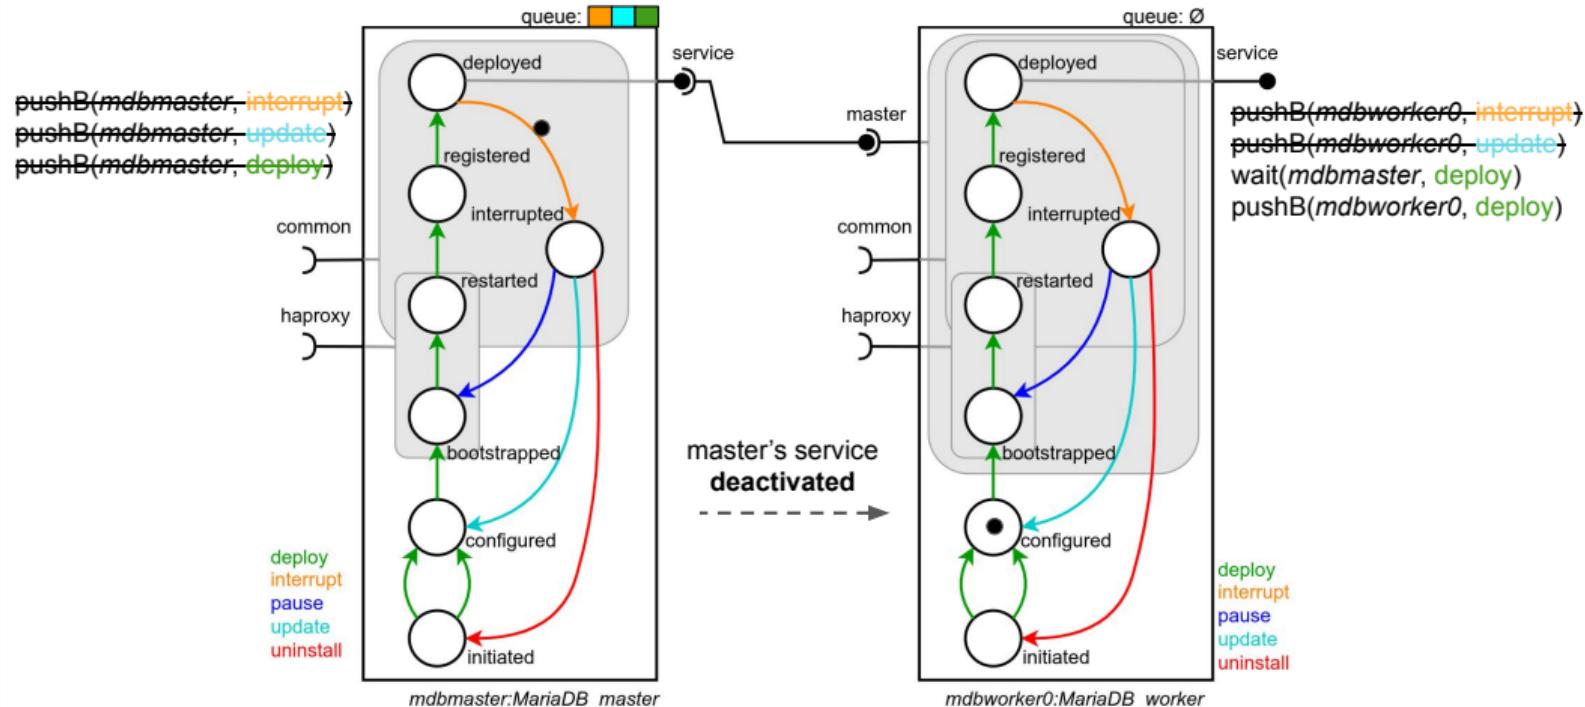
# Execution example



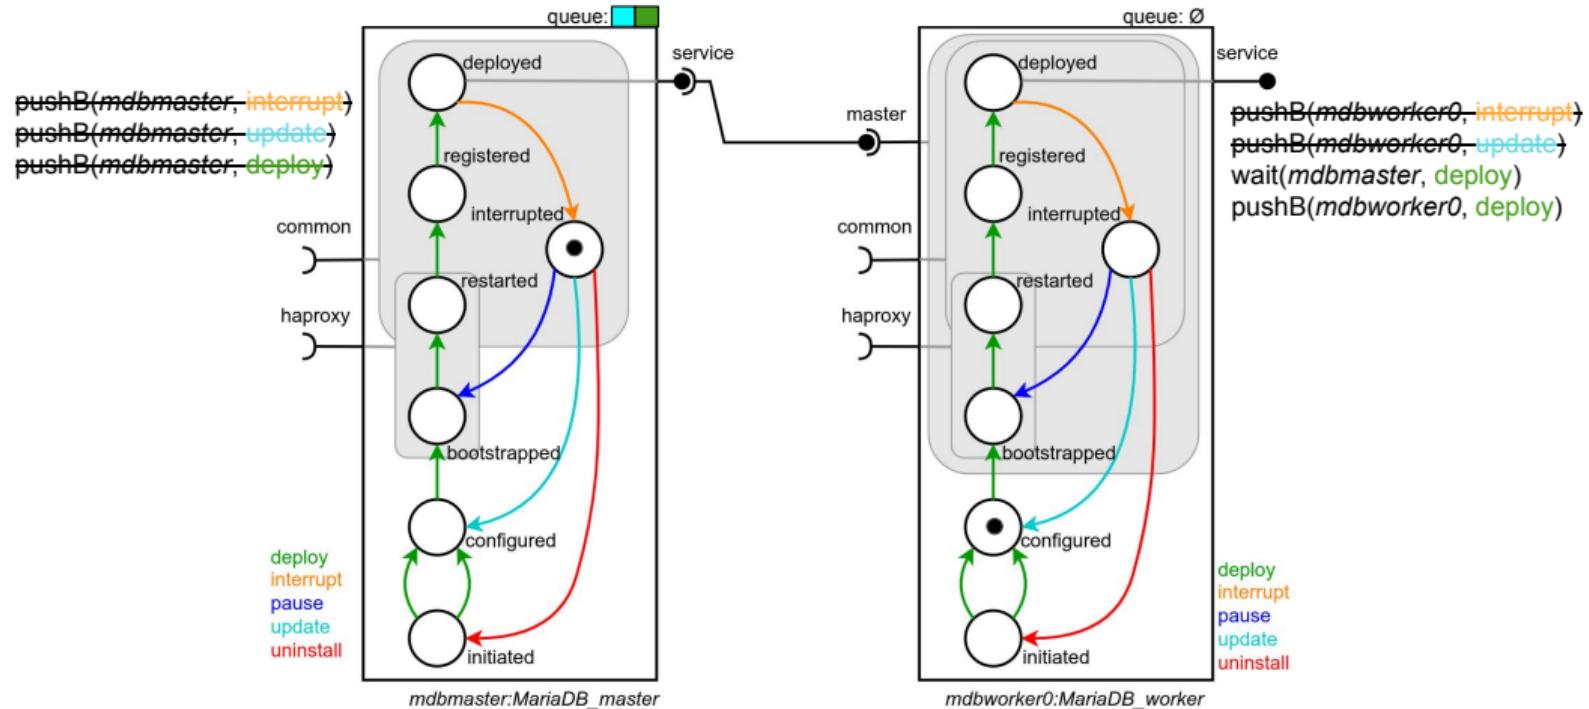
# Execution example



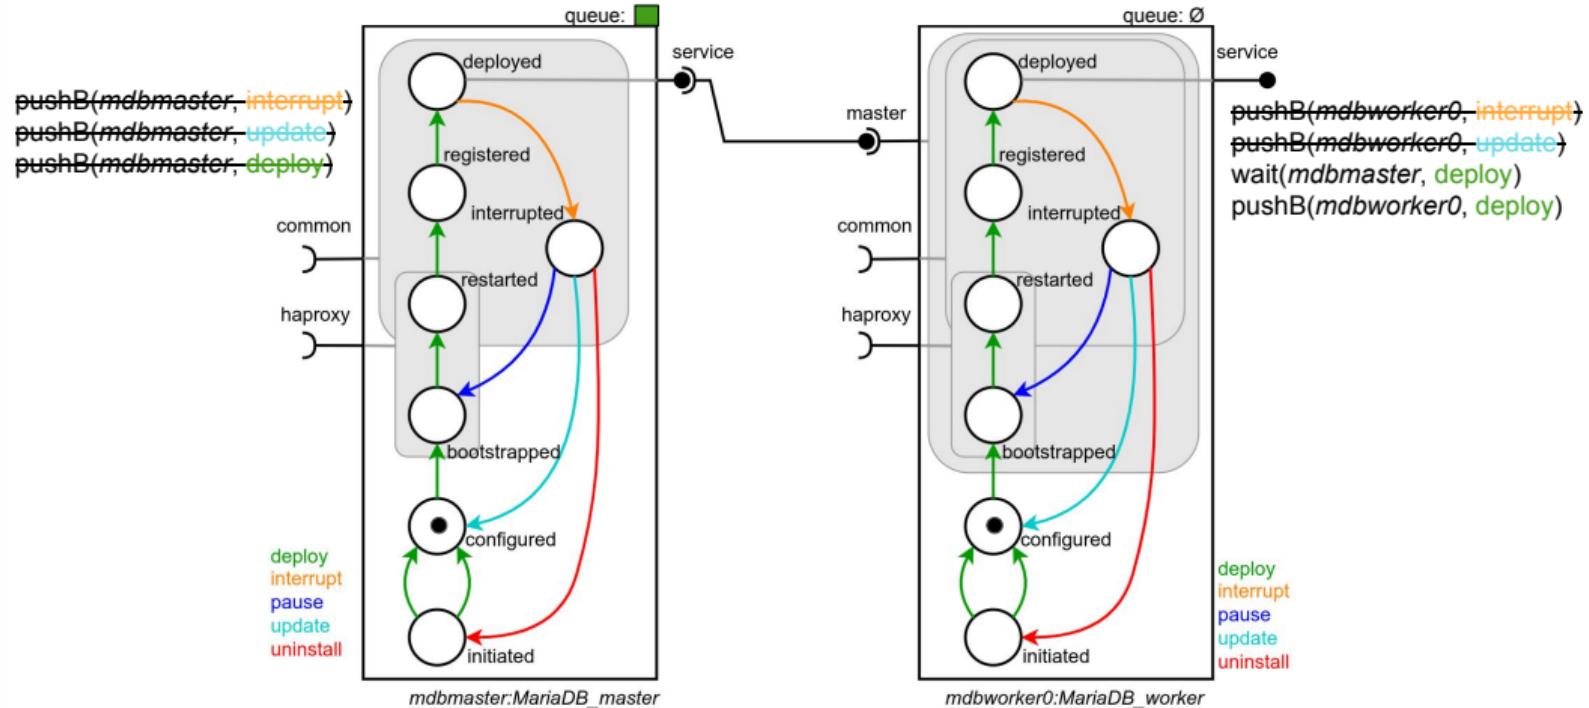
# Execution example



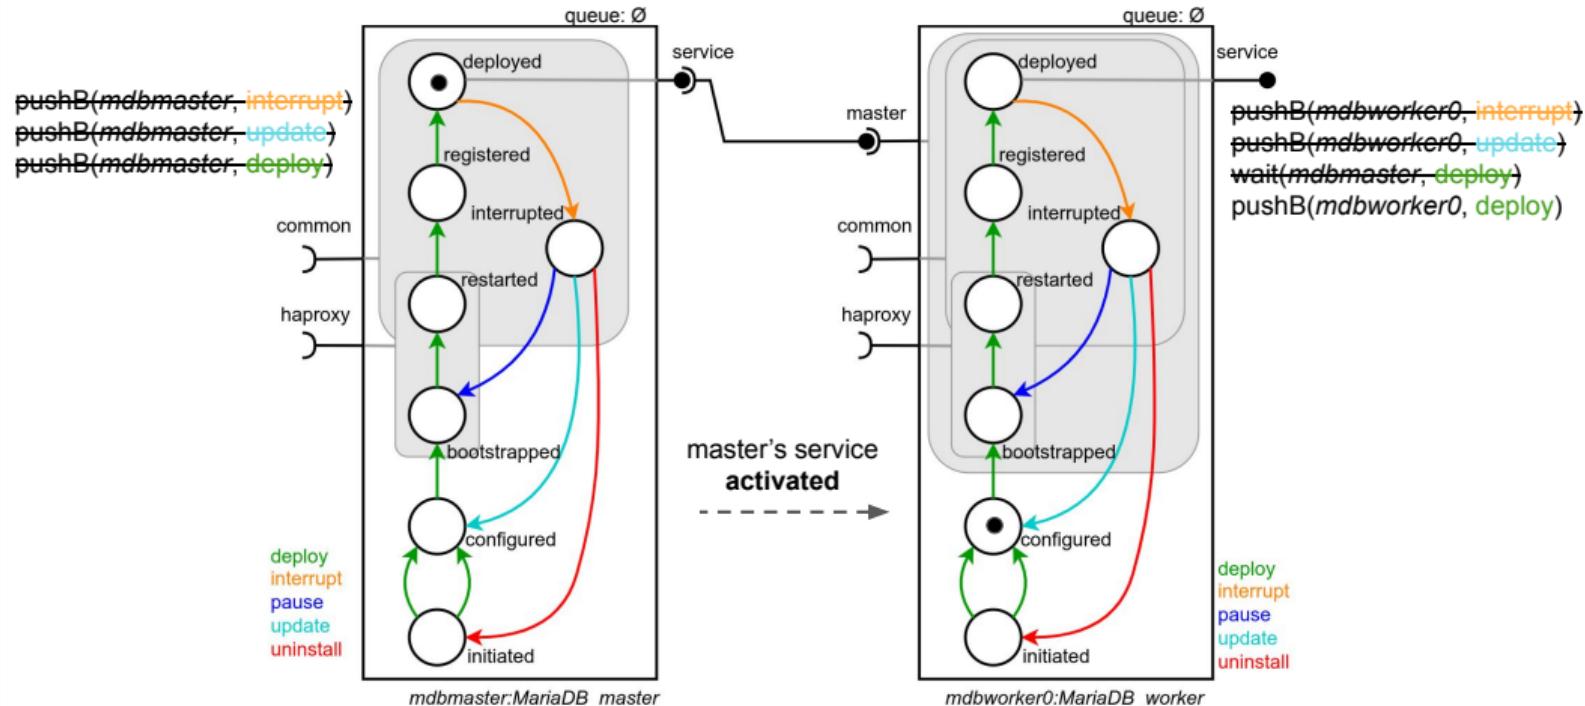
# Execution example



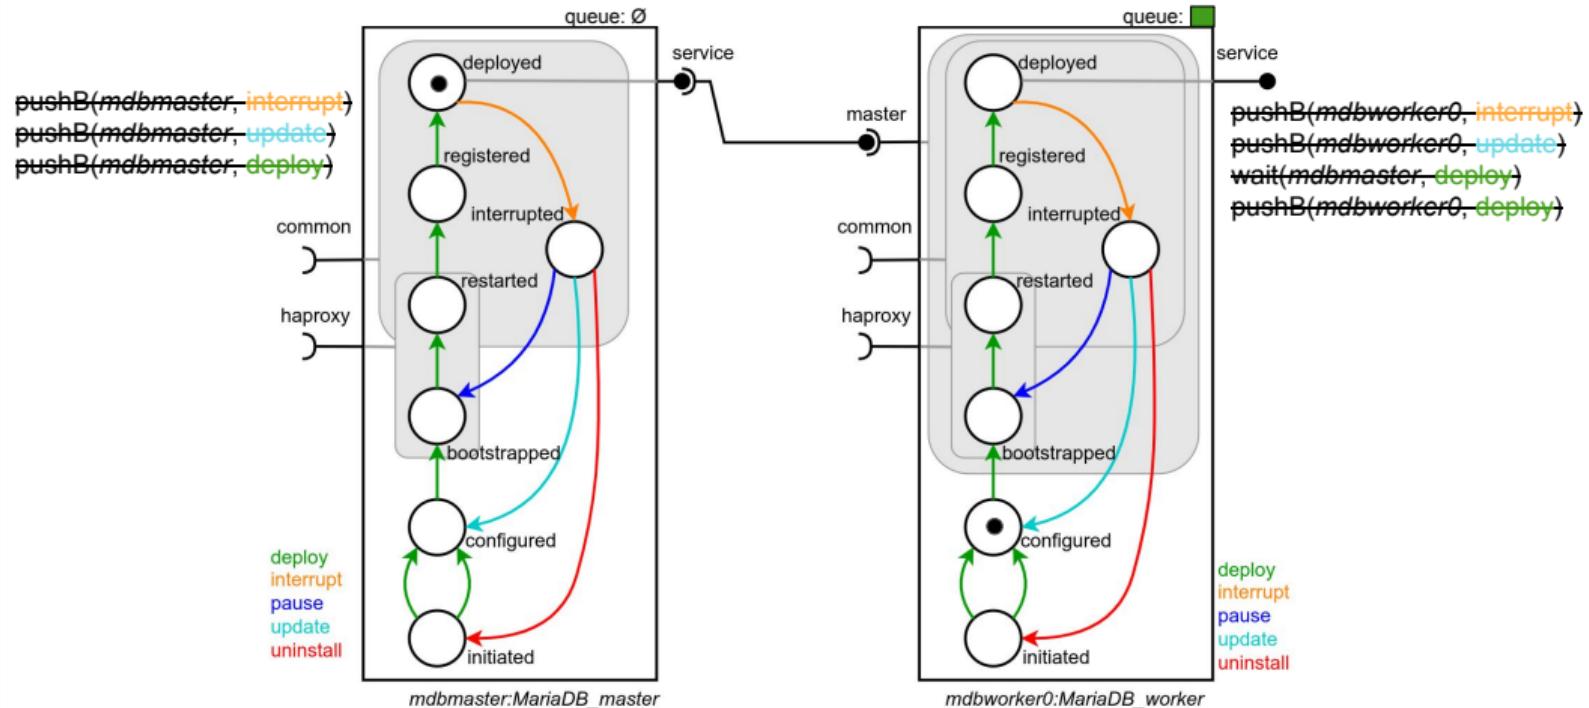
# Execution example



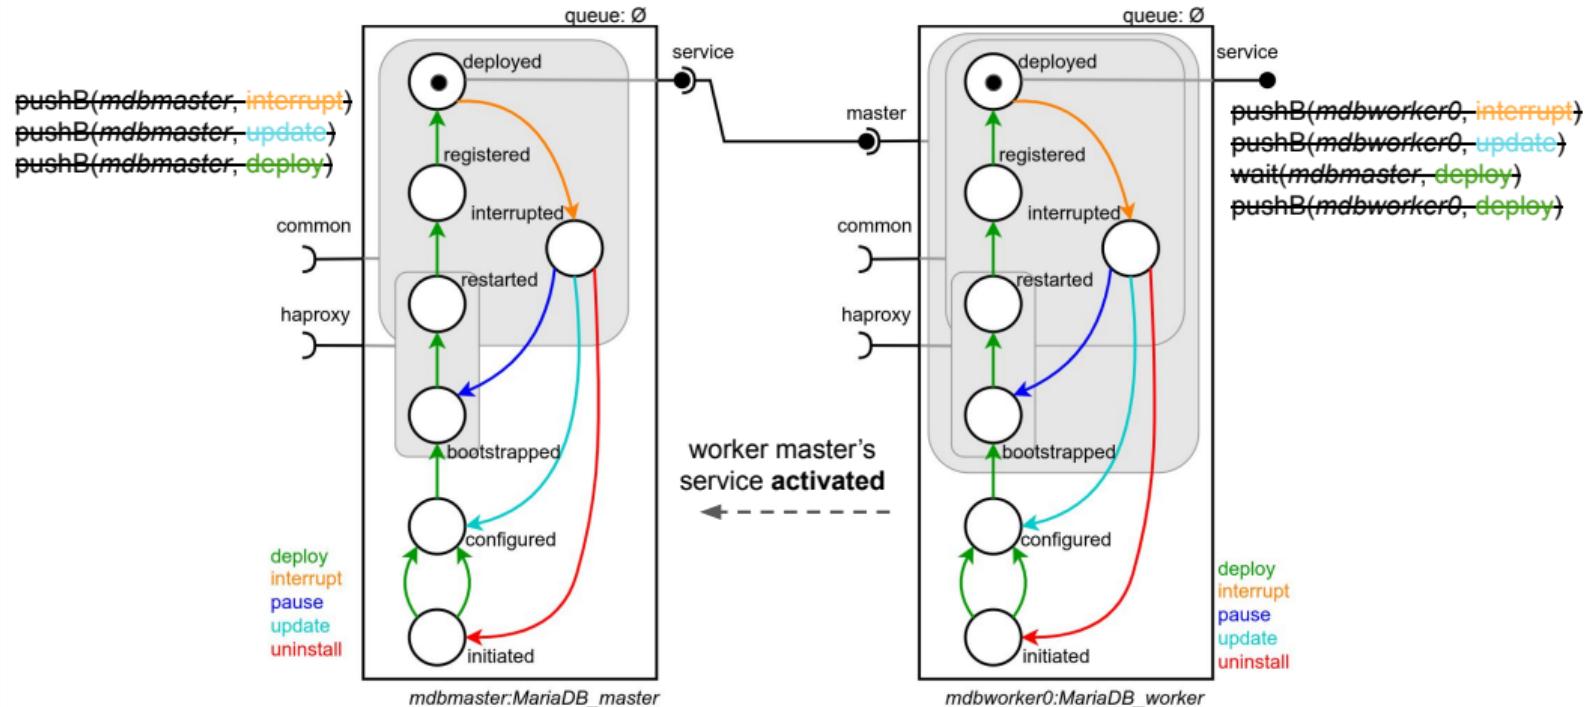
# Execution example



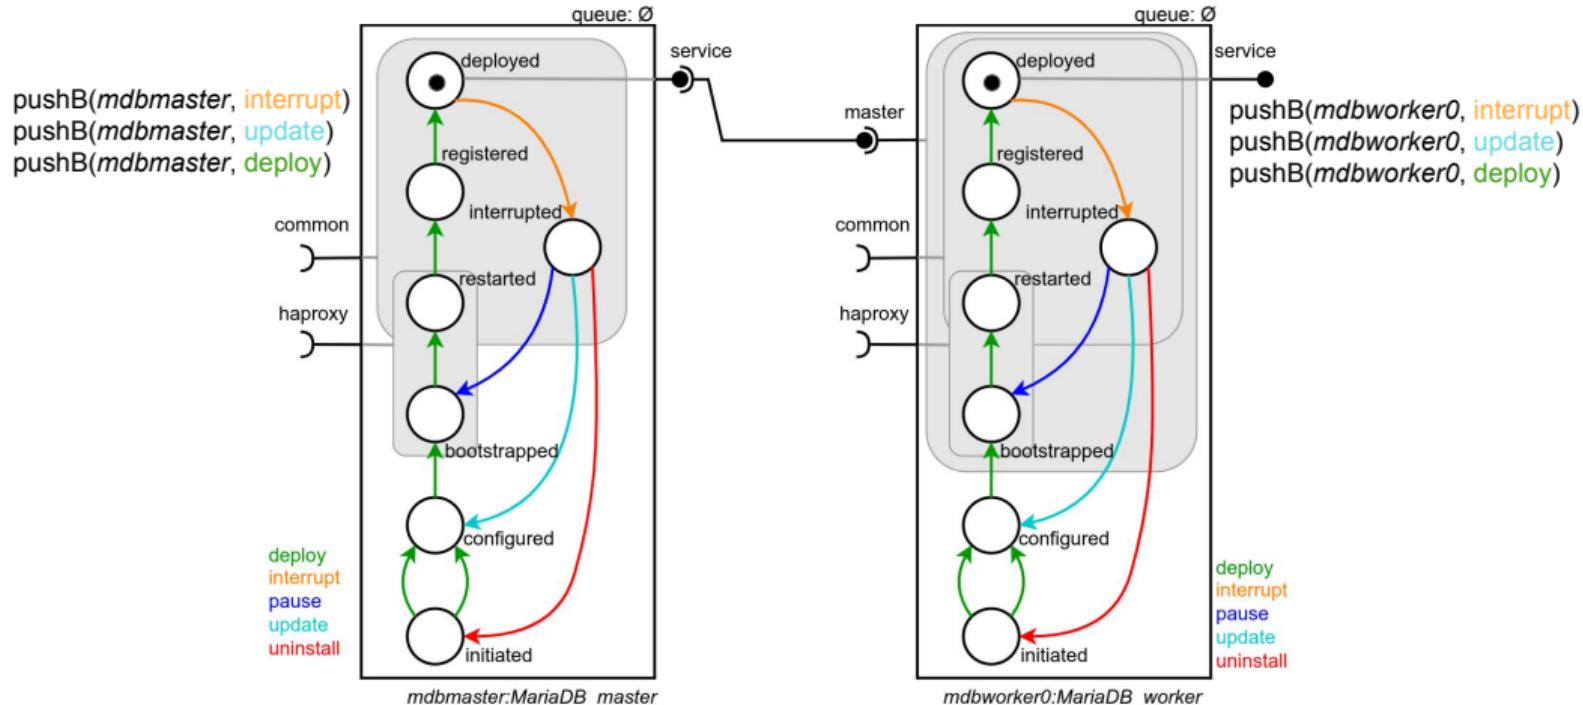
# Execution example



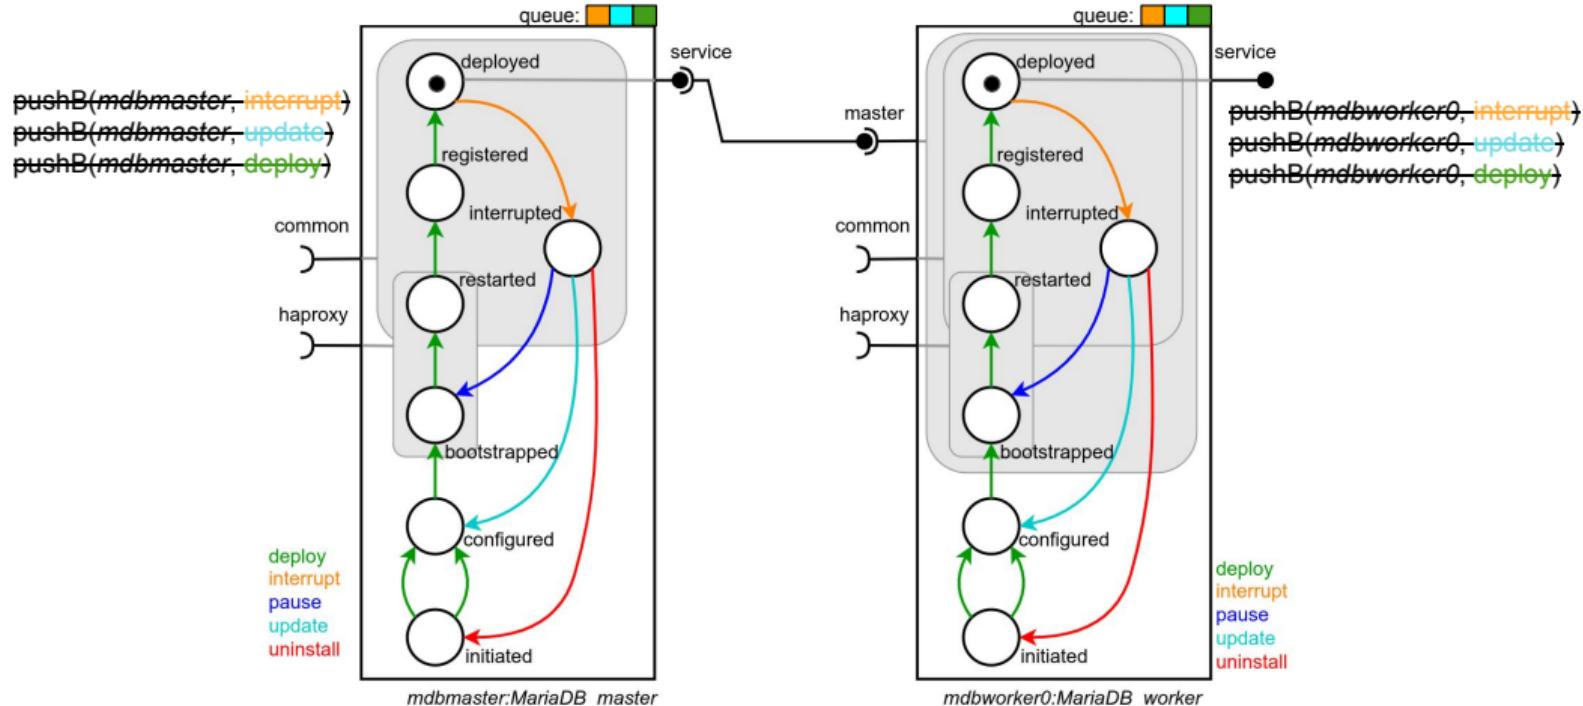
# Execution example



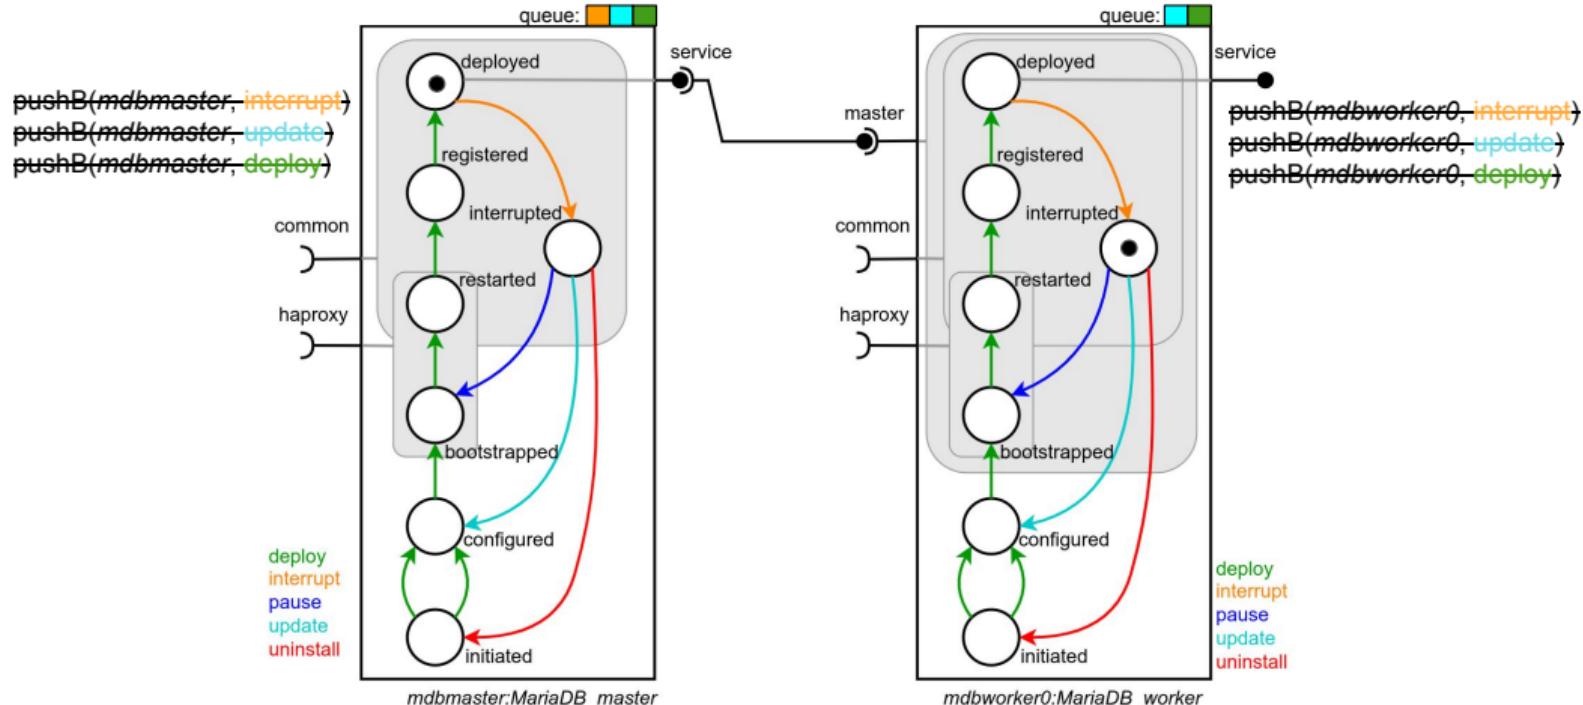
# Without the wait ? Failing execution



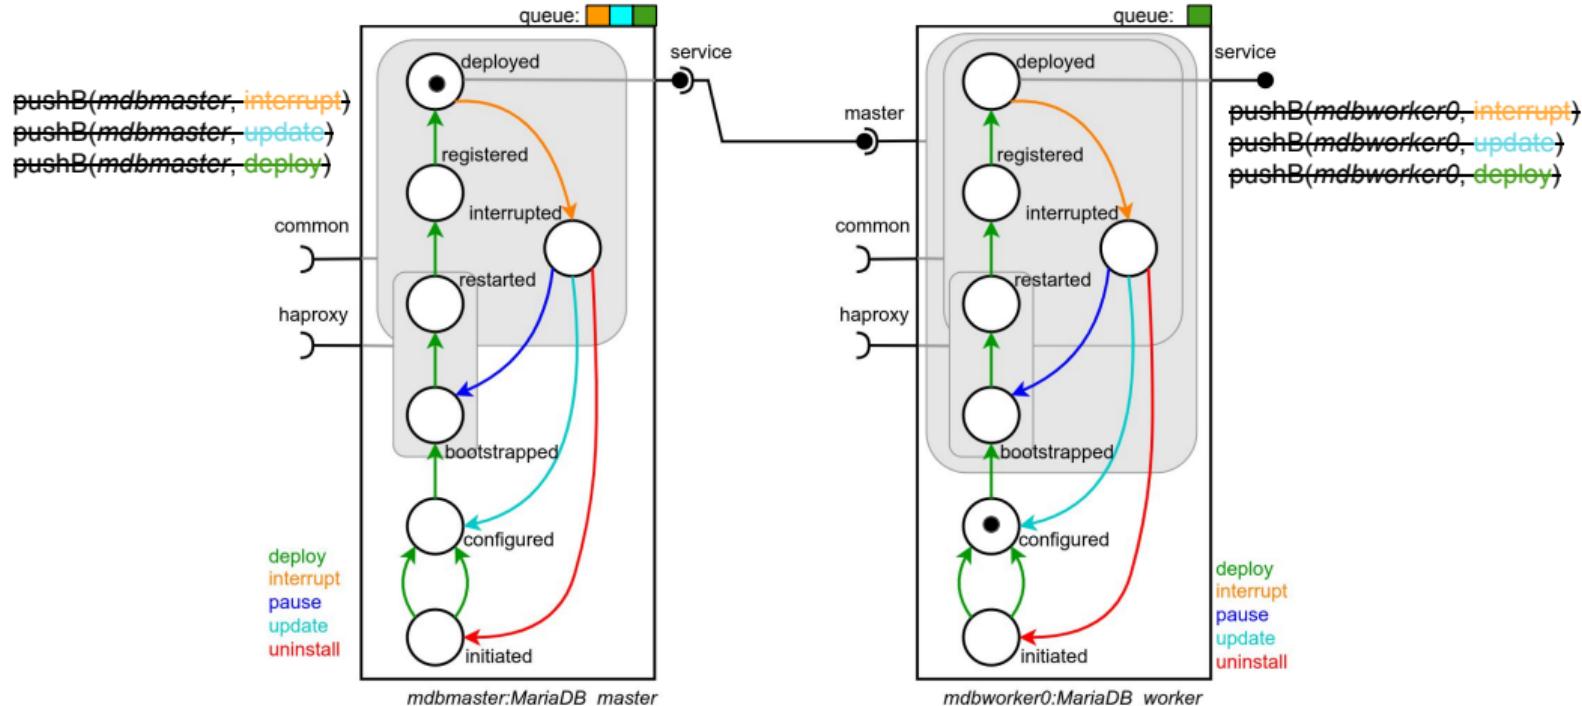
# Without the wait ? Failing execution



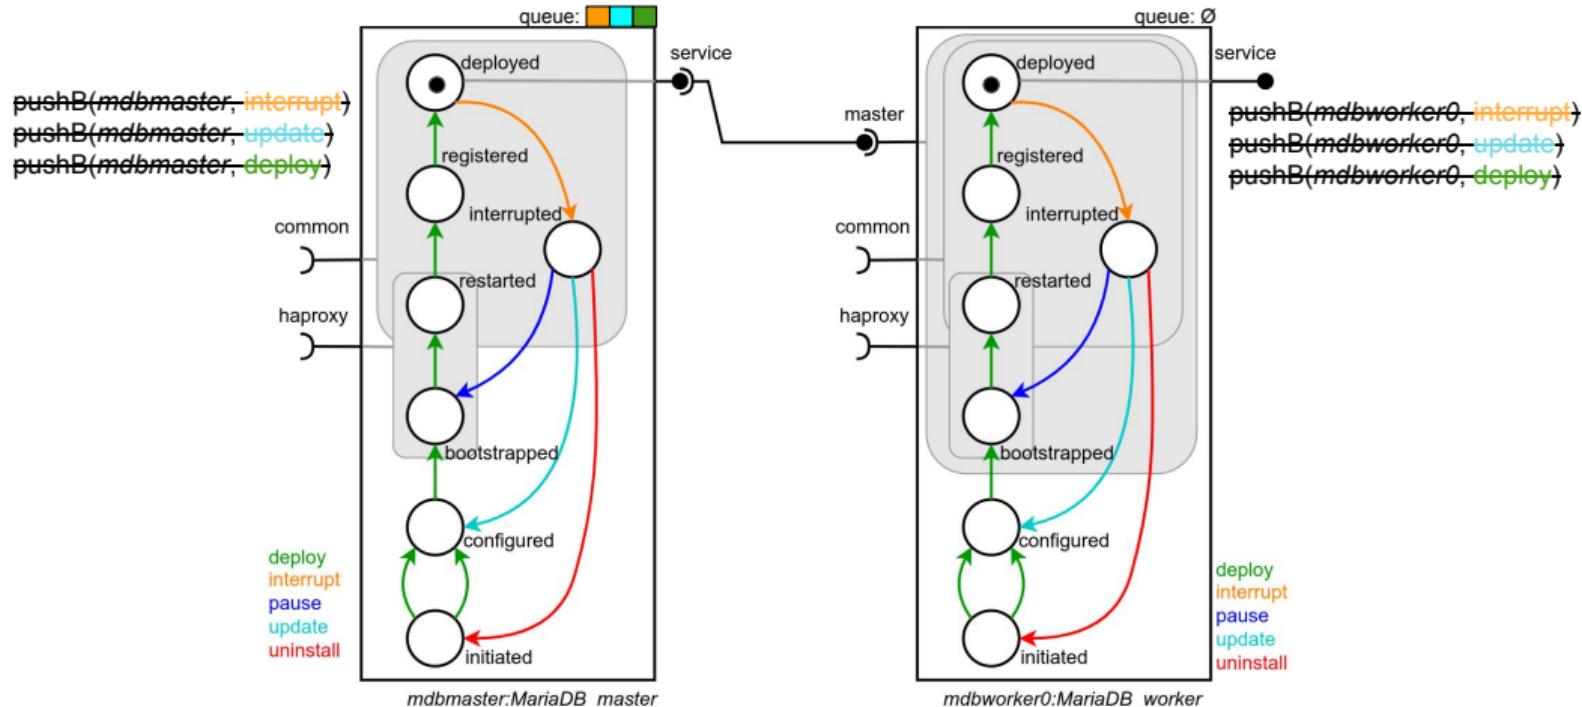
# Without the wait ? Failing execution



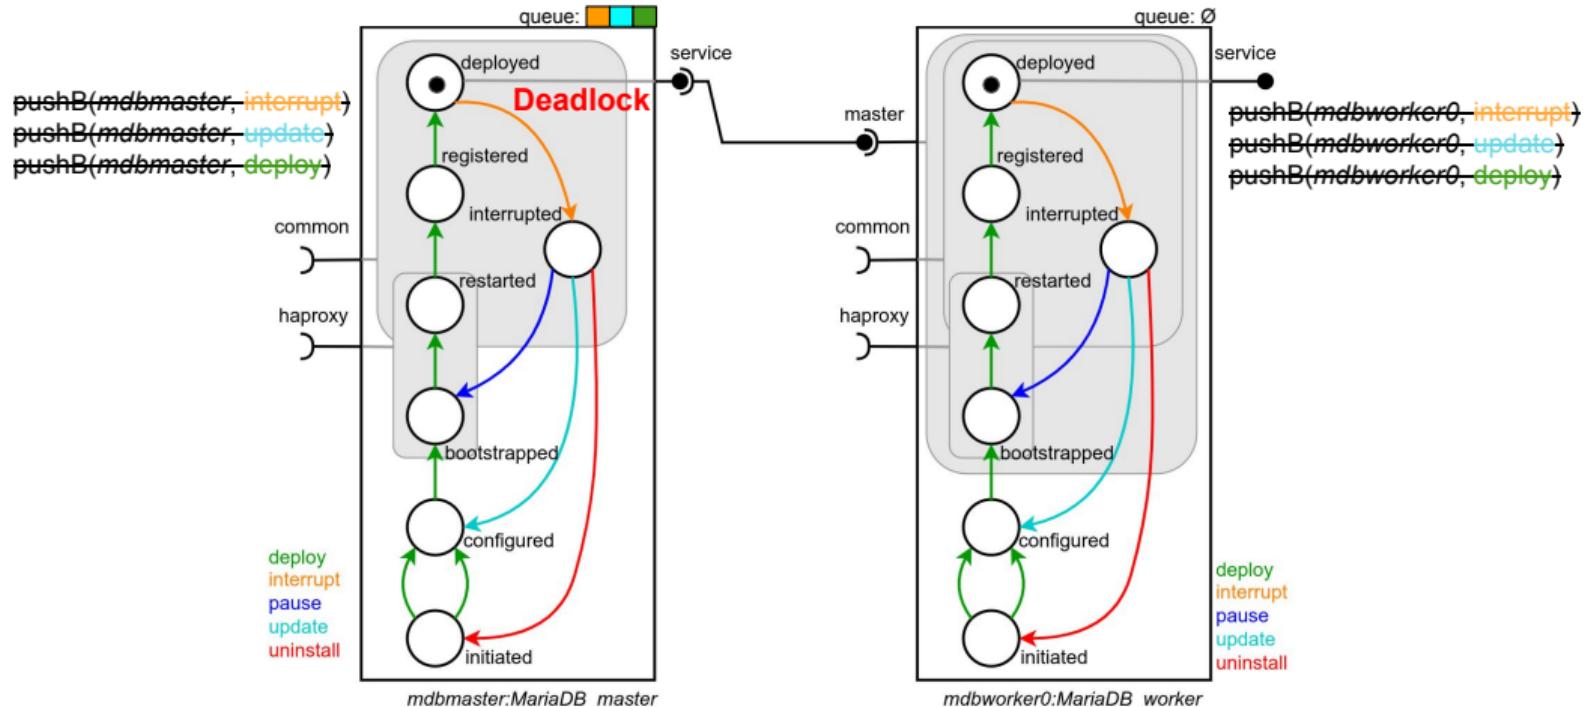
# Without the wait ? Failing execution



# Without the wait ? Failing execution



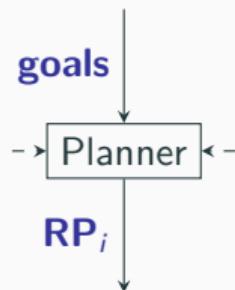
# Without the wait ? Failing execution



# Planning Concerto-D programs

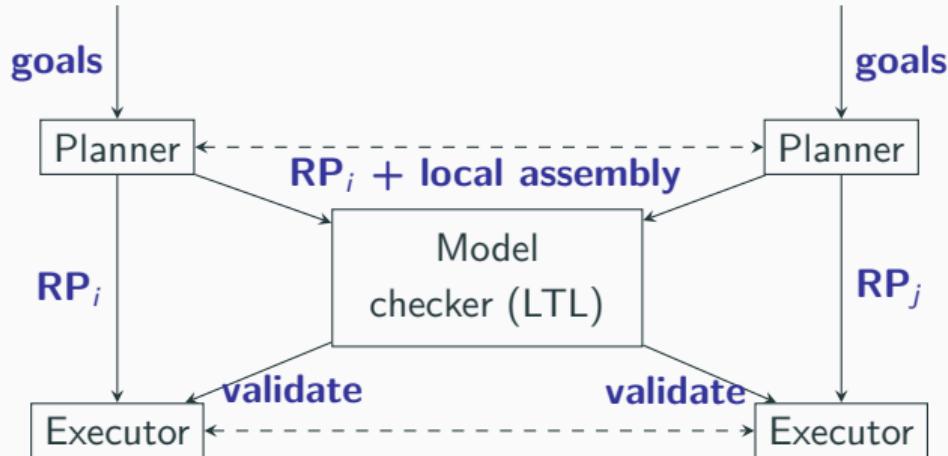
## Decentralized planner

- **Input:** goals and lifecycle
- **Output:** a reconfiguration plan
- On each node, iterative resolution :
  - Using SAT solver for intermediate plans
  - Diffusing port constraints, to enrich neighborhood constraint models



- Sat solver ensure validity of the Reconfiguration Plan (RP)
- If the model is unsat, we find the MUS (Minimum Unsat Satisfiability) for explainability, and return error to user

# Verify Ballet's execution ?



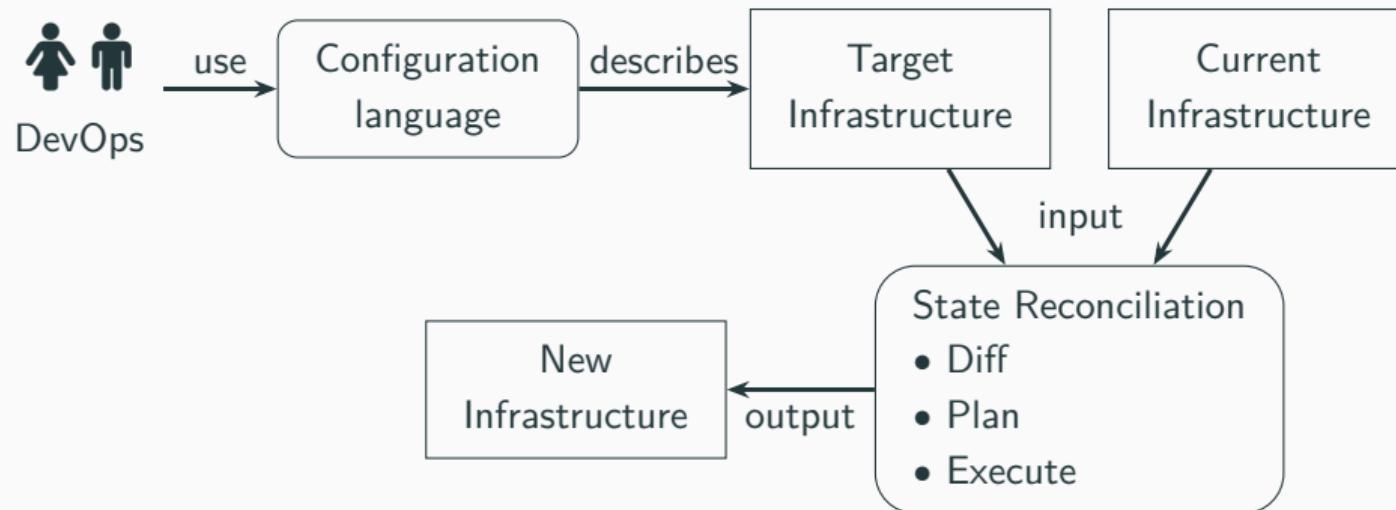
## Model checker

- Formalized Ballet's executor within Maude
- Model checking with linear temporal logic (LTL)
- **Pros:** A first step for verifying properties
- **Cons:** Works with all plans and full assembly
- **Cons:** Current formalization does not scale for realistic applications

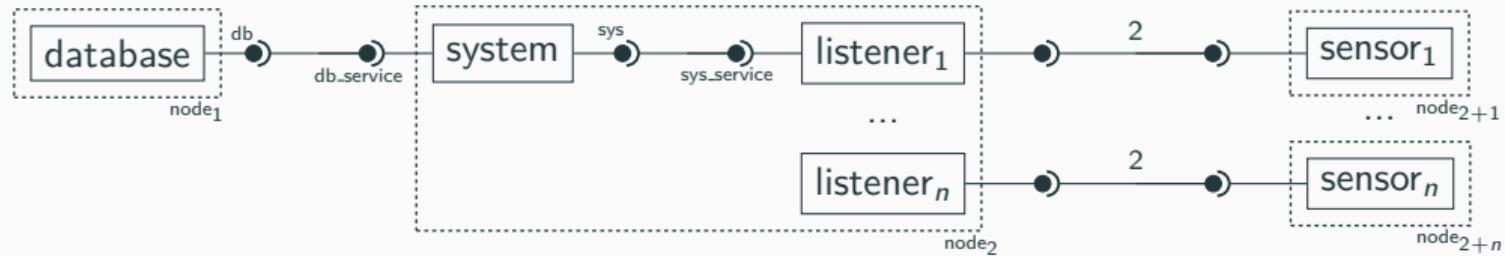
# Integrate Ballet - Infrastructure as code

## Infrastructure as code (IaC)

Infrastructure as Code (IaC) is the practice of defining and managing infrastructure using code (written in configuration languages). Tools then take this code and automatically deploy the infrastructure as specified.



# Integrate Ballet - Control fleet of CPS



# Conclusion

## SparkTE

- A configurable model transformation engine
- A correct by construction engine on top Spark

## Verifying parallel implementation with SyDPaCC

- Extended a Coq library for verifying skeletons
- Formalized a subpart of Spark

## Reconfiguration with Ballet

- Declarative tool for decentralized reconfiguration
- Decentralized planning with SAT solver
- Decentralized execution of plans
- Premises of model checking

# Perspectives

## Verifying Spark

- Formalize distribution process of Spark RDDs calculation
- More support for Spark functions
- Implements additional skeletons in Spark

## Model checking for Ballet

- Define additional properties (e.g., safety)
- Use partial order reduction techniques for reducing state-space exploration
- Decentralized checking ? by composition + distribution

## Distributed systems

- Energetic optimization (e.g., placement problem)
- Energetic model for configuration space

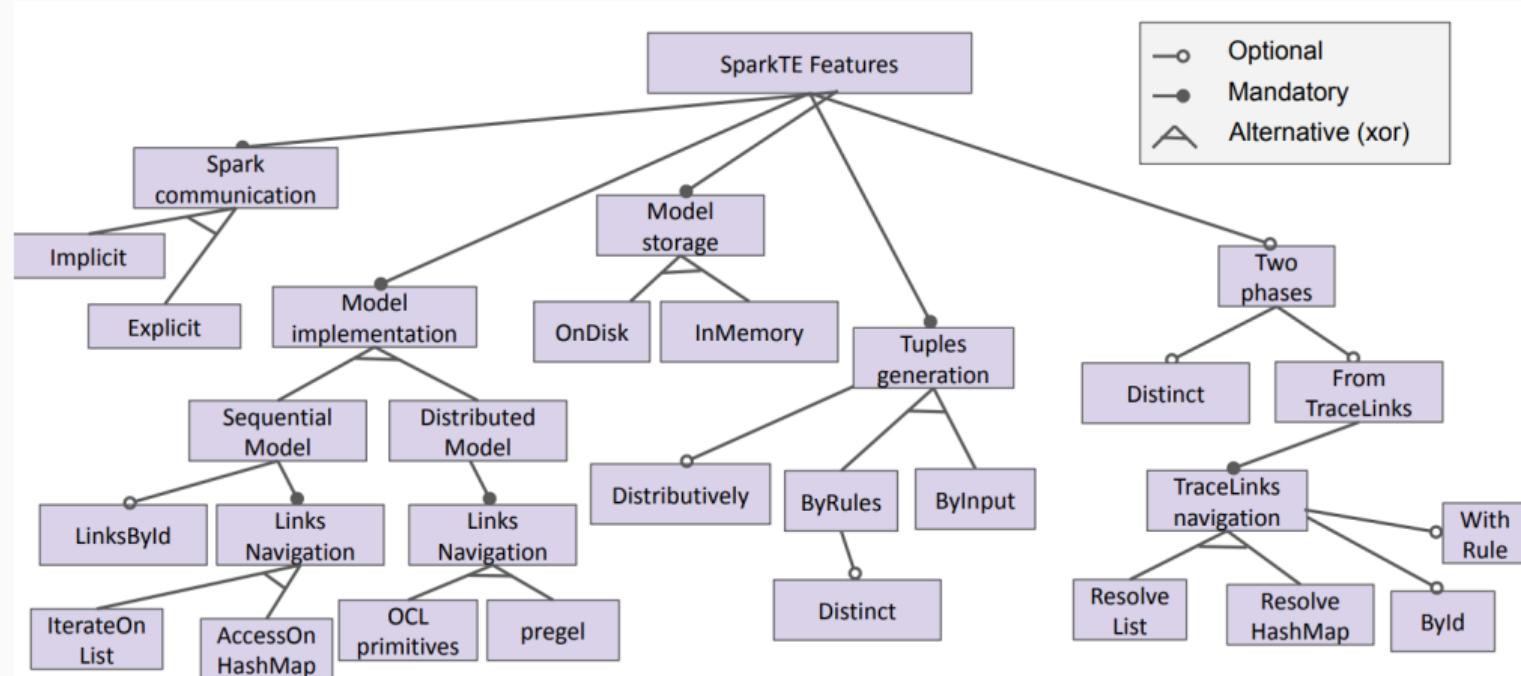
## References

- [1 ] Jolan Philippe. Contribution to the Analysis of the Design-Space of a Distributed Transformation Engine. Ph.D thesis. IMT Atlantique, 2022.
- [2 ] Jolan Philippe, Massimo Tisi, Hélène Coullon, and Gerson Sunyé. Executing Certified Model Transformations on Apache Spark. In 14th ACM SIGPLAN International Conference on Software Language Engineering (SLE), pages 36–48. ACM, 2021.
- [3 ] Frédéric Loulergue, and Jolan Philippe. Towards Verified Scalable Parallel Computing with Coq and Spark. In 25th ACM International Workshop on Formal Techniques for Java-like Programs (FTfJP), pages 11–17. ACM, 2023.
- [4 ] Jolan Philippe, Antoine Omond, Hélène Coullon, Charles Prud'Homme, and Issam Raïs. Fast Choreography of Cross-DevOps Reconfiguration with Ballet: A Multi-Site OpenStack Case Study. In IEEE International Conference on Software Analysis, Evolution and Reengineering (SANER). IEEE, 2024.
- [5 ] Farid Arfi, Hélène Coullon, Frédéric Loulergue, Jolan Philippe, and Simon Robillard. A Maude Formalization of the Distributed Reconfiguration Language Concerto-D. In 17th Interaction and Concurrency Experience (ICE). EPTCS 2024.

## Backup

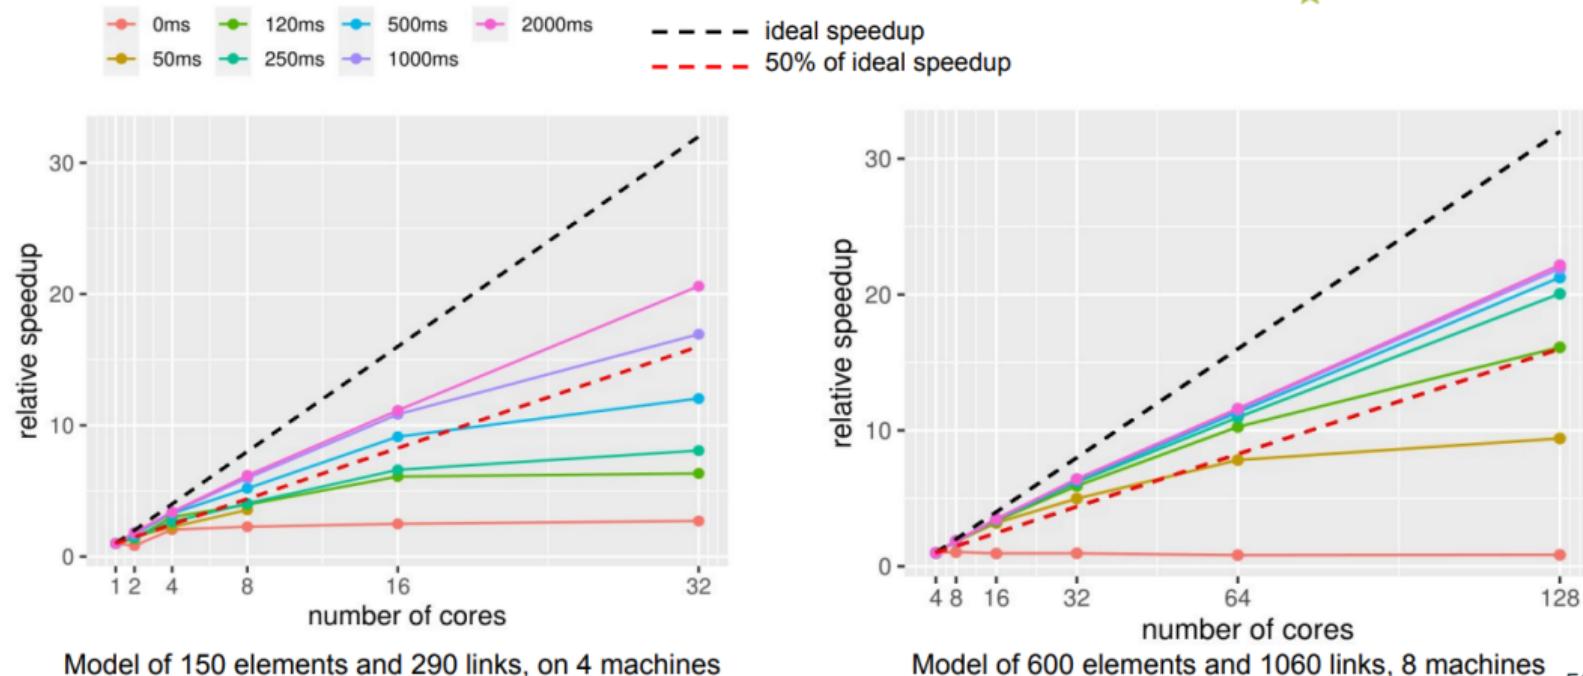
---

# SparkTE - Configuration space overview

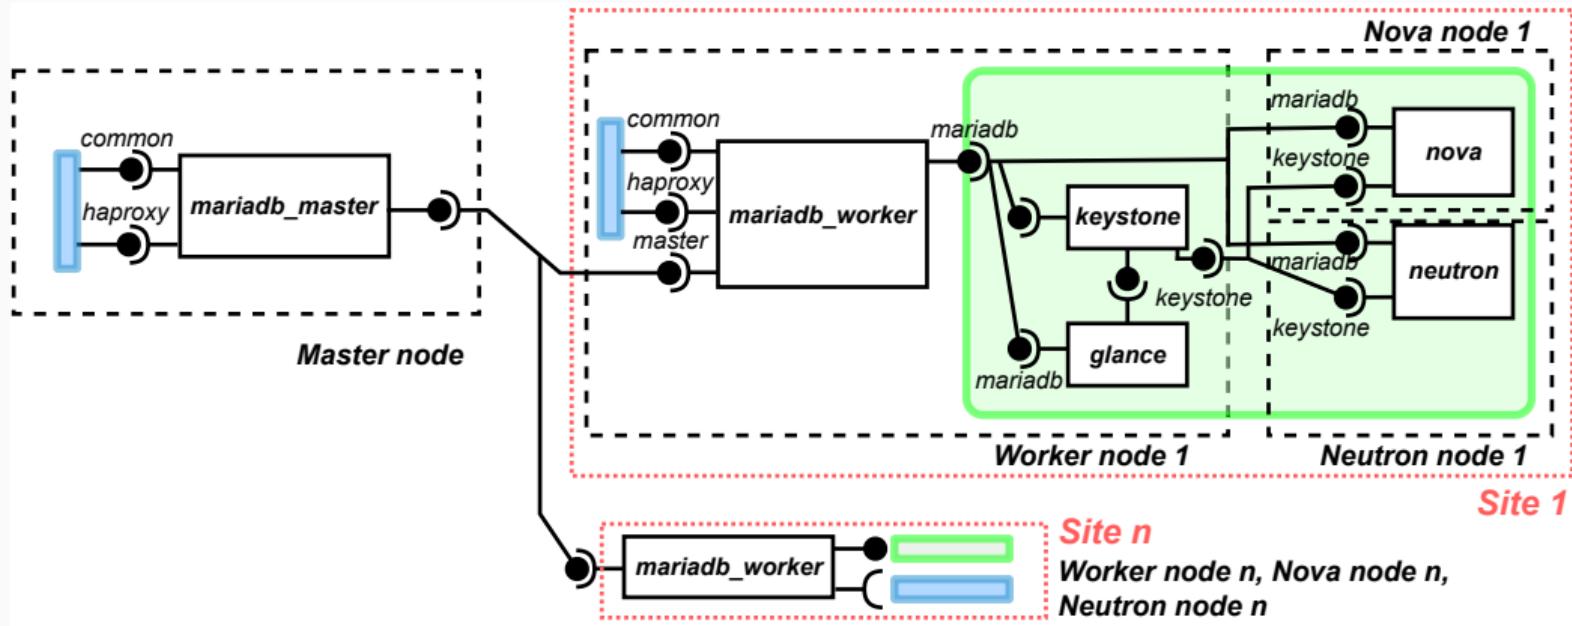


# SparkTE Performances

- Simulate a uniform amount of computation on nodes
  - fixed time for each task



# Ballet performances on real use-case



## Ballet performances on real use-case

Sc.	# Sites	Planning	Ballet Execution	Total	Muse	Gain
Deploy	1	1.69s	306.02s	307.71s	536.57s	42.7%
	2	1.78s	306.09s	307.86s	536.69s	42.6%
	5	1.77s	306.19s	307.97s	537.09s	42.7%
	10	2.02s	306.14s	308.19s	538.13s	42.7%
Update	1	3.36s	416.84s	420.20s	555.56s	24.4%
	2	4.39s	416.92s	421.31s	555.70s	24.2%
	5	6.05s	417.17s	423.22s	556.08s	24.0%
	10	5.97s	417.46s	423.43s	556.77s	24.0%

**Table 1:** Comparison of time for planning and executing a deployment and an update of the MariaDB\_master instance with Ballet and Muse.

- $(B, \Pi, \mathcal{C}, s_{init}, S_{goal})$
- $s_{i+1} = inc_{\Pi}[s_i][b_i], \forall i \in 1..m$
- $(b, B, >, 0)$
- $status(p, s_{m+1}) = \Gamma_p$

**where**

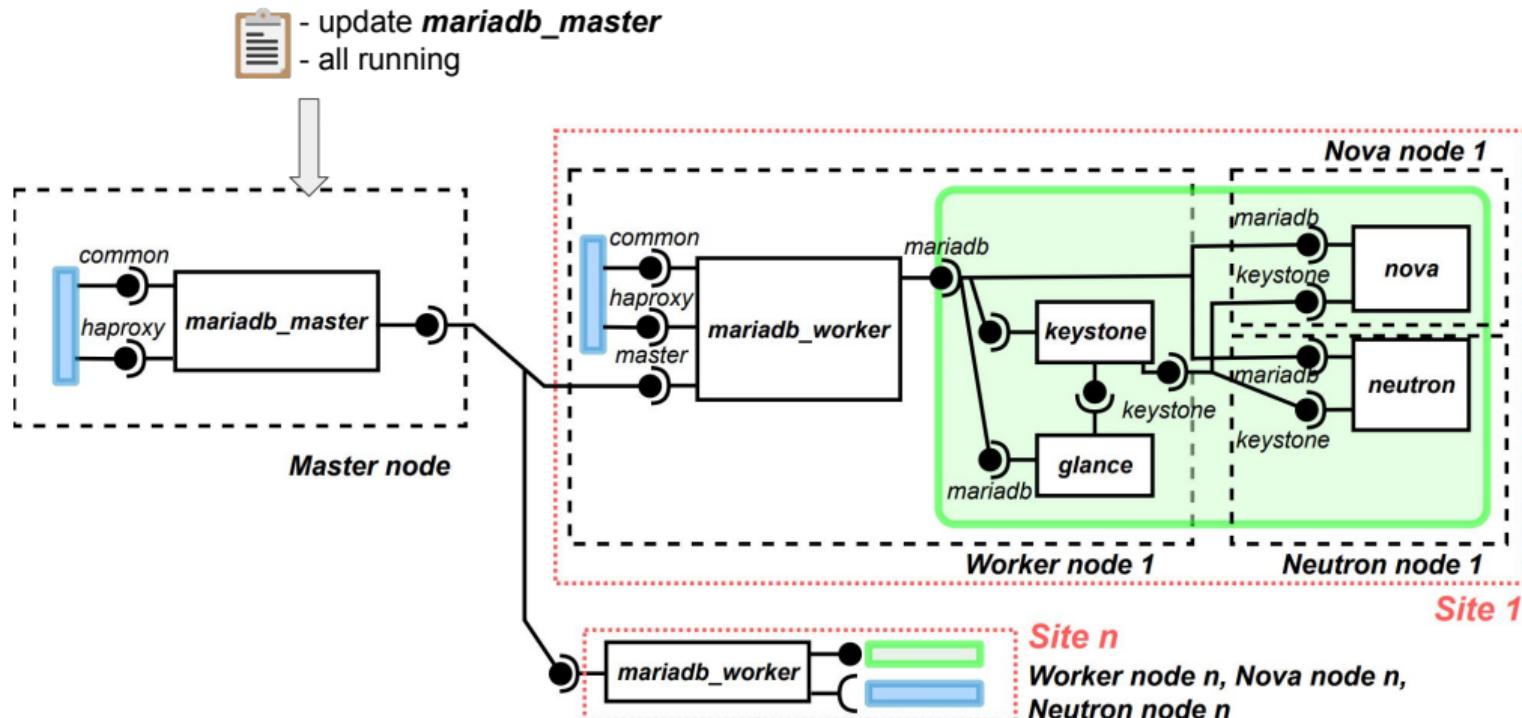
$\Pi$  an automaton with  $\mathcal{C}$  costs

$B$  a sequence of  $m$  behaviors

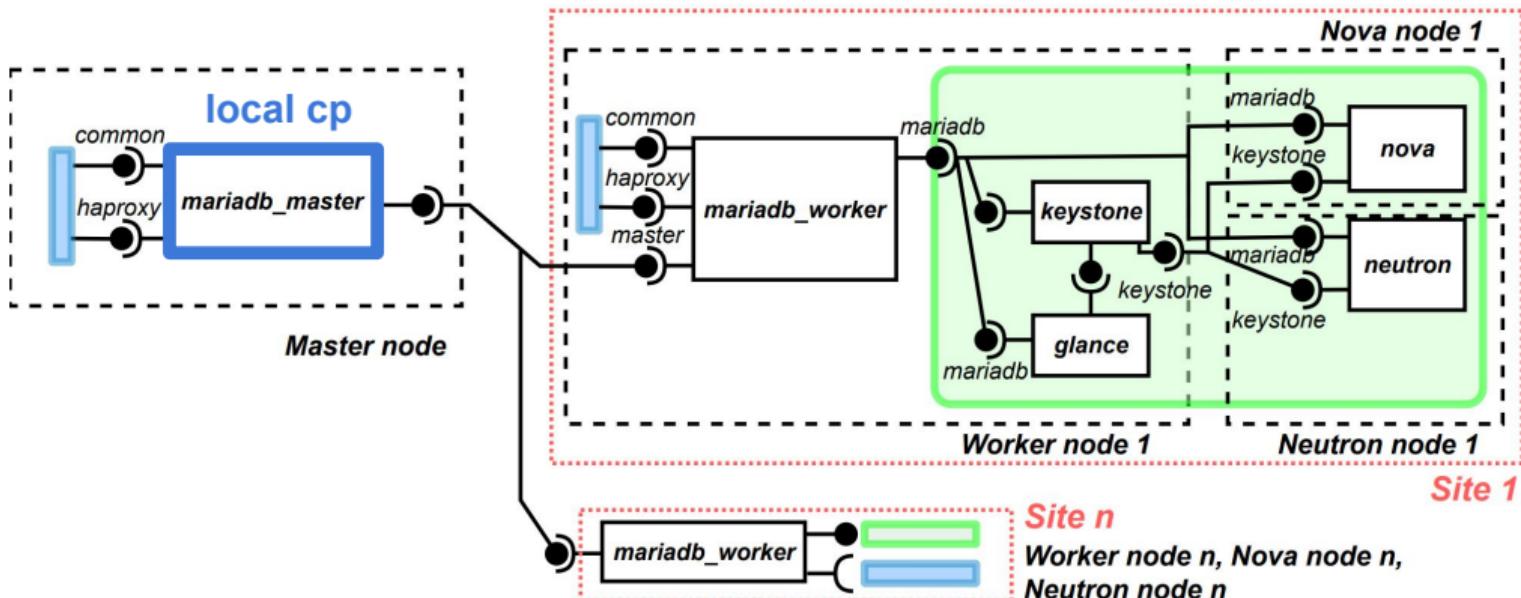
$\Gamma_p \in \{\text{active, inactive}\}$  i.e. { ✓, ✗ }

$b \in \{\text{interrupt, deploy, pause, update, uninstall}\}$

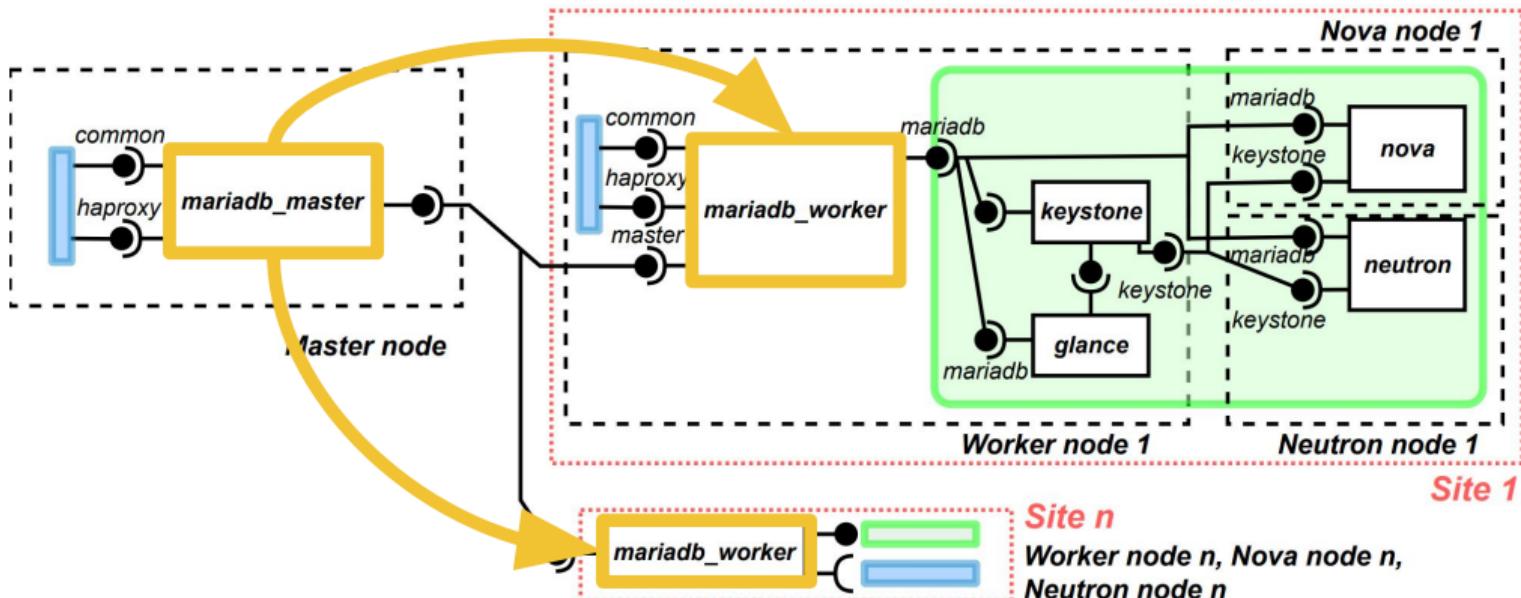
# Communication protocol



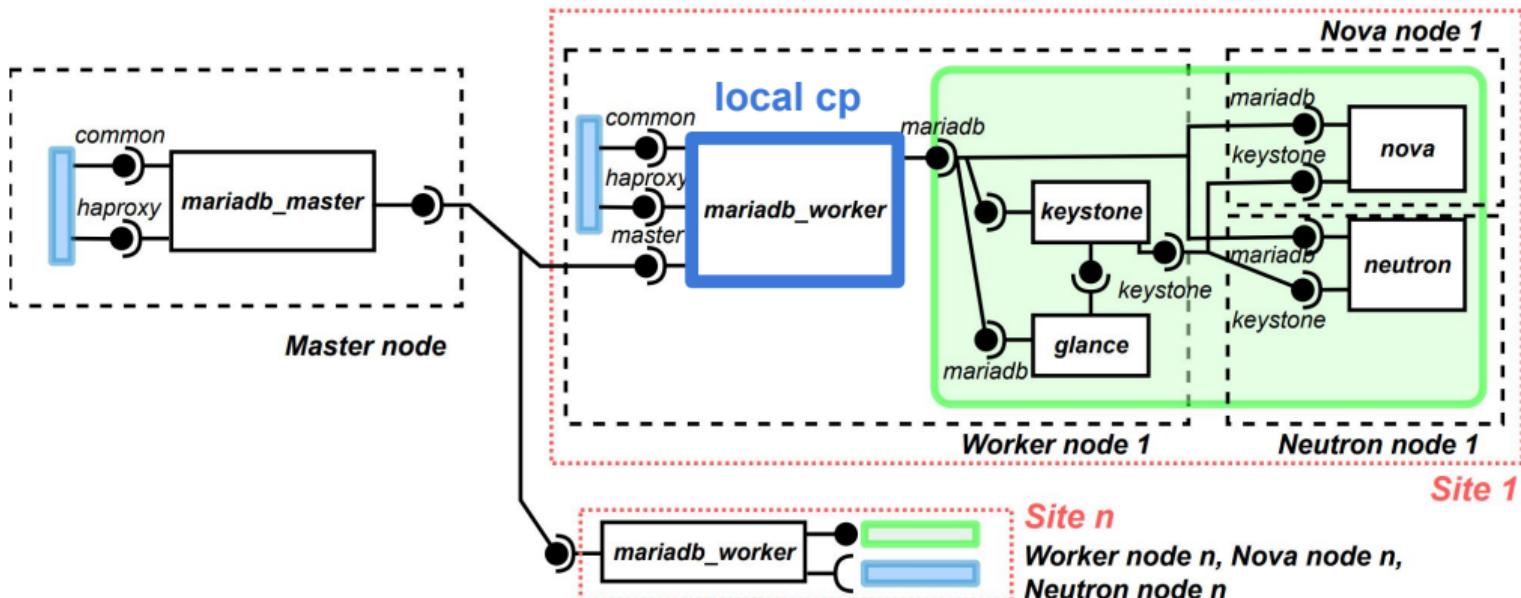
# Communication protocol



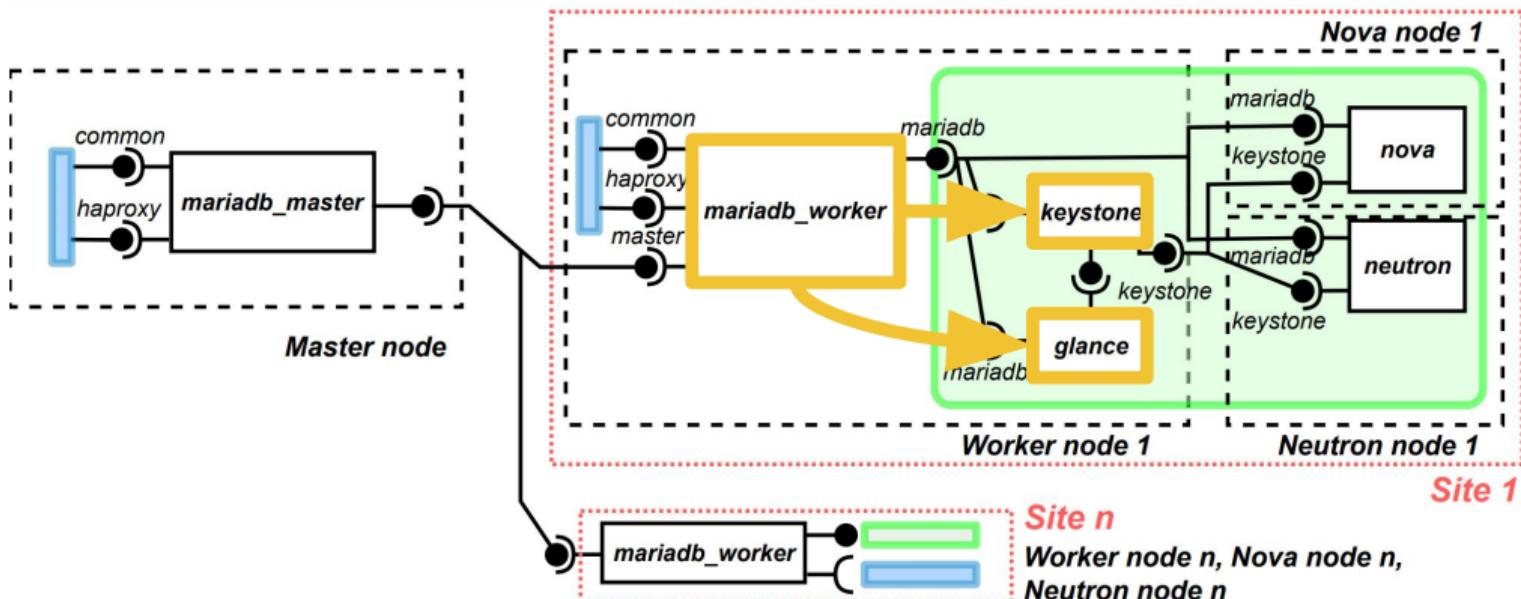
# Communication protocol



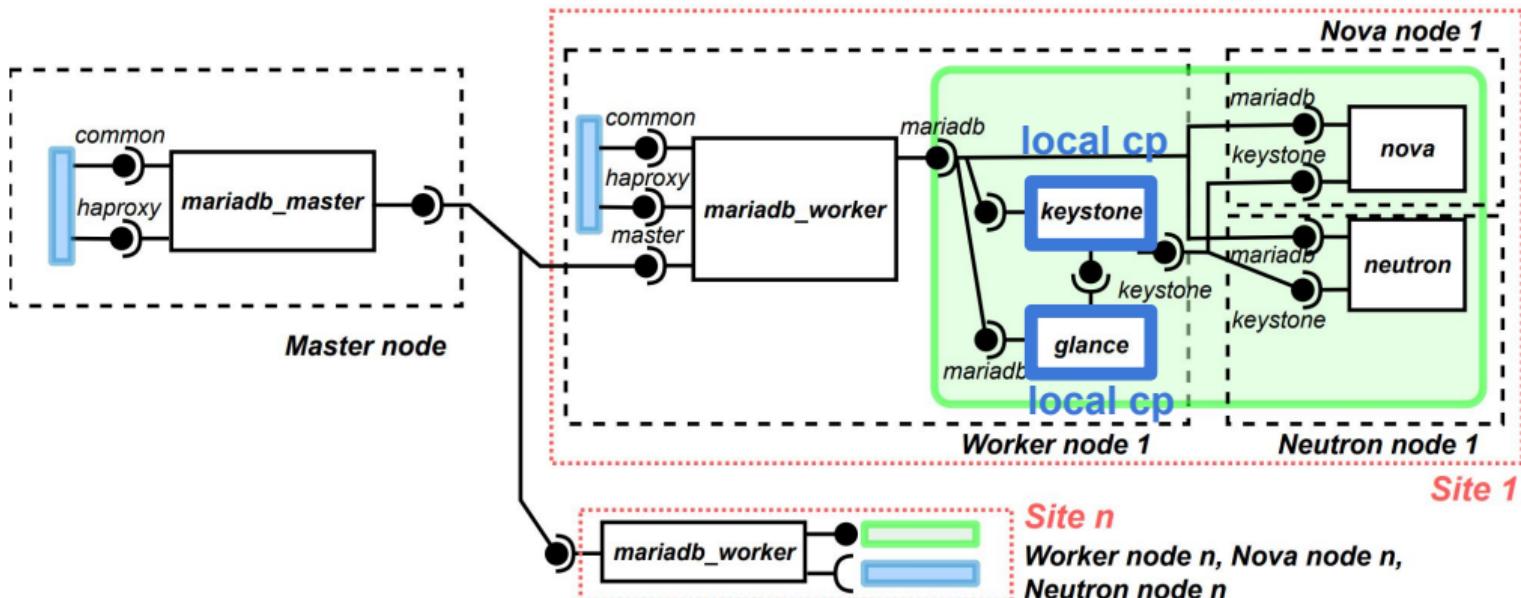
# Communication protocol



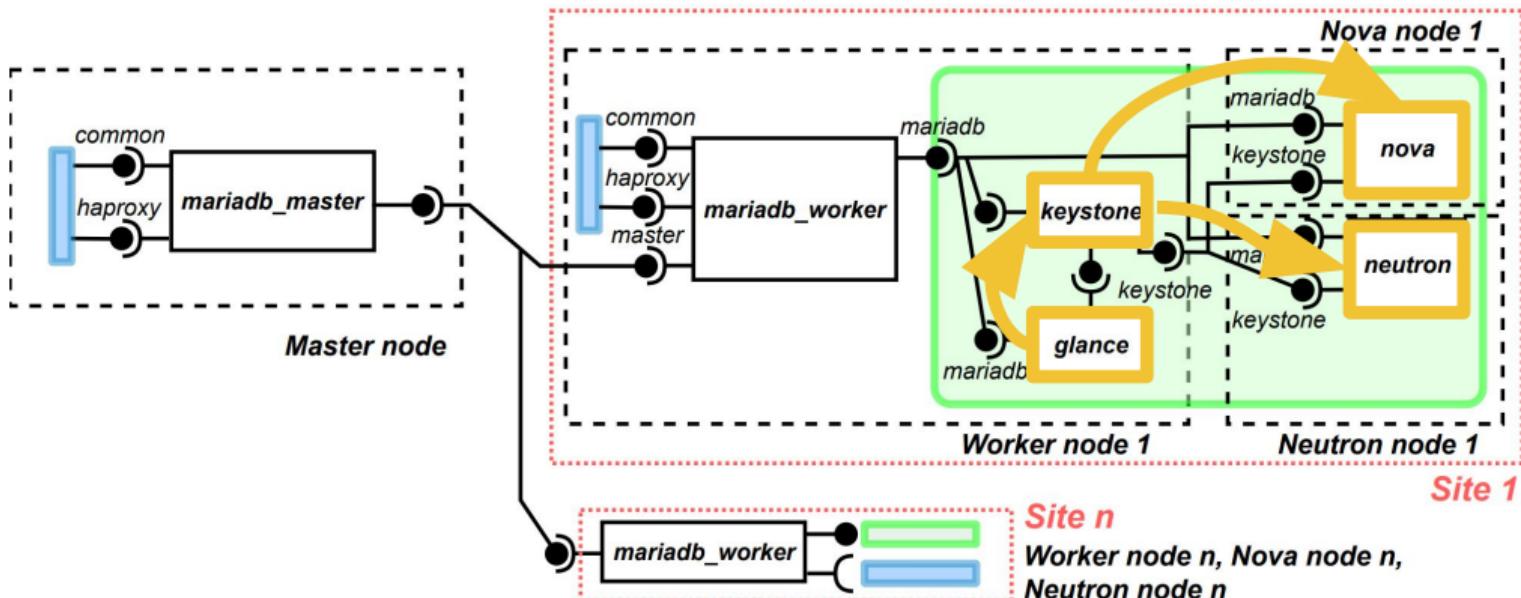
# Communication protocol



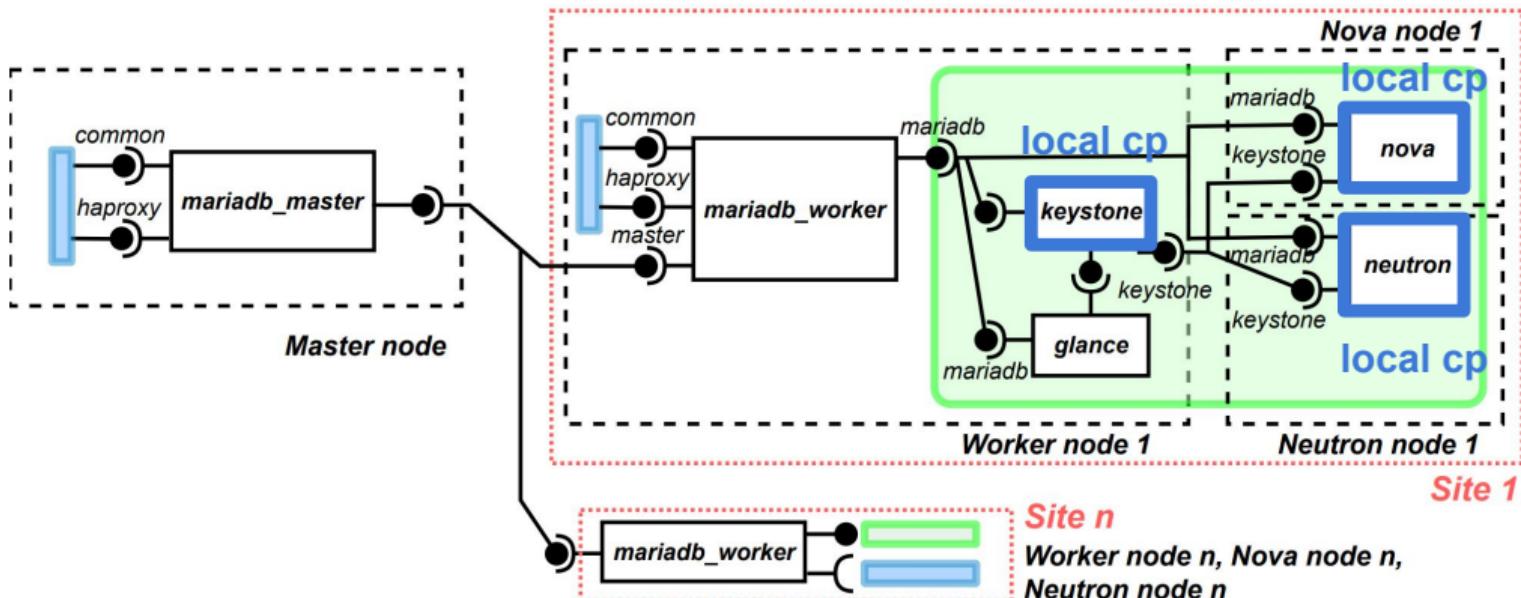
# Communication protocol



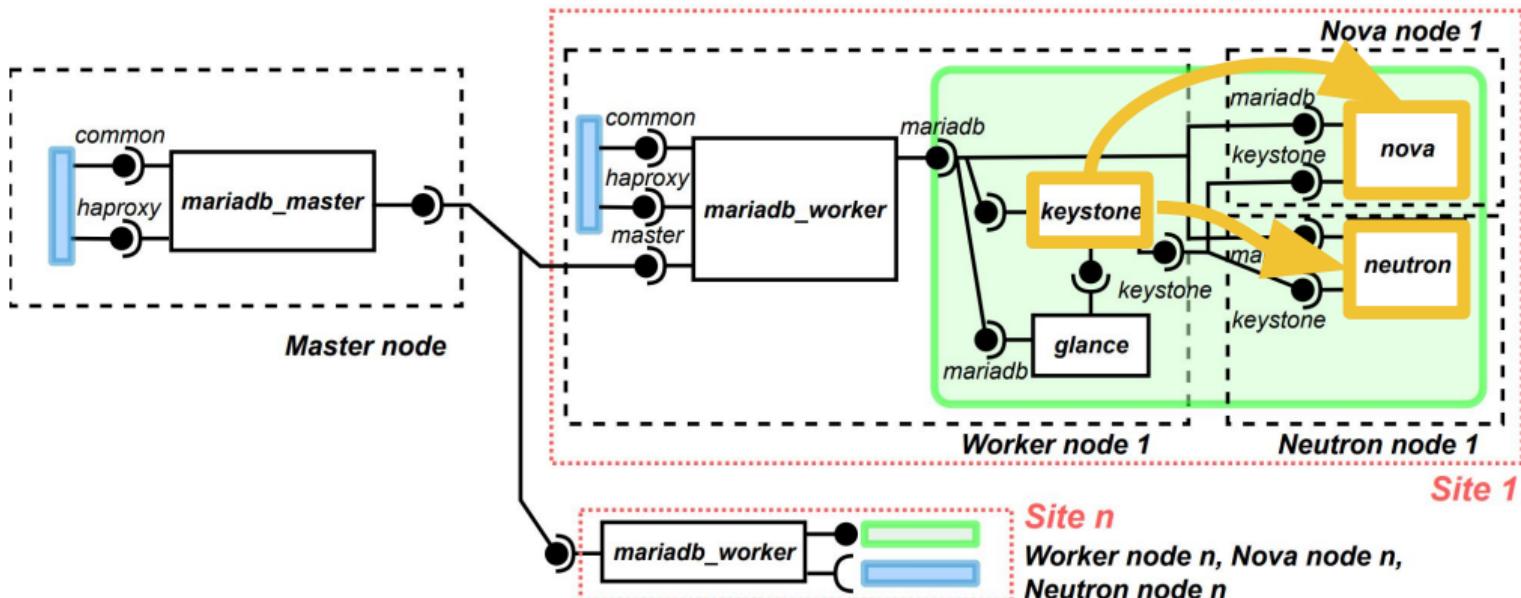
# Communication protocol



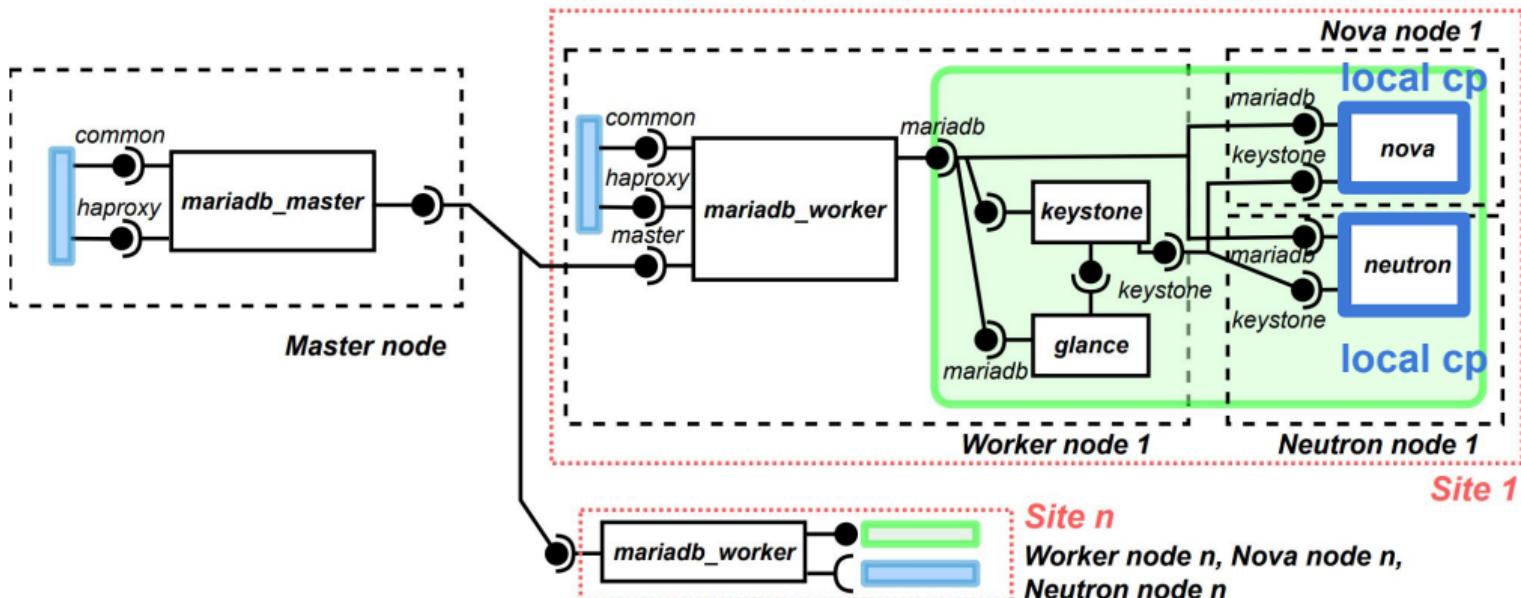
# Communication protocol



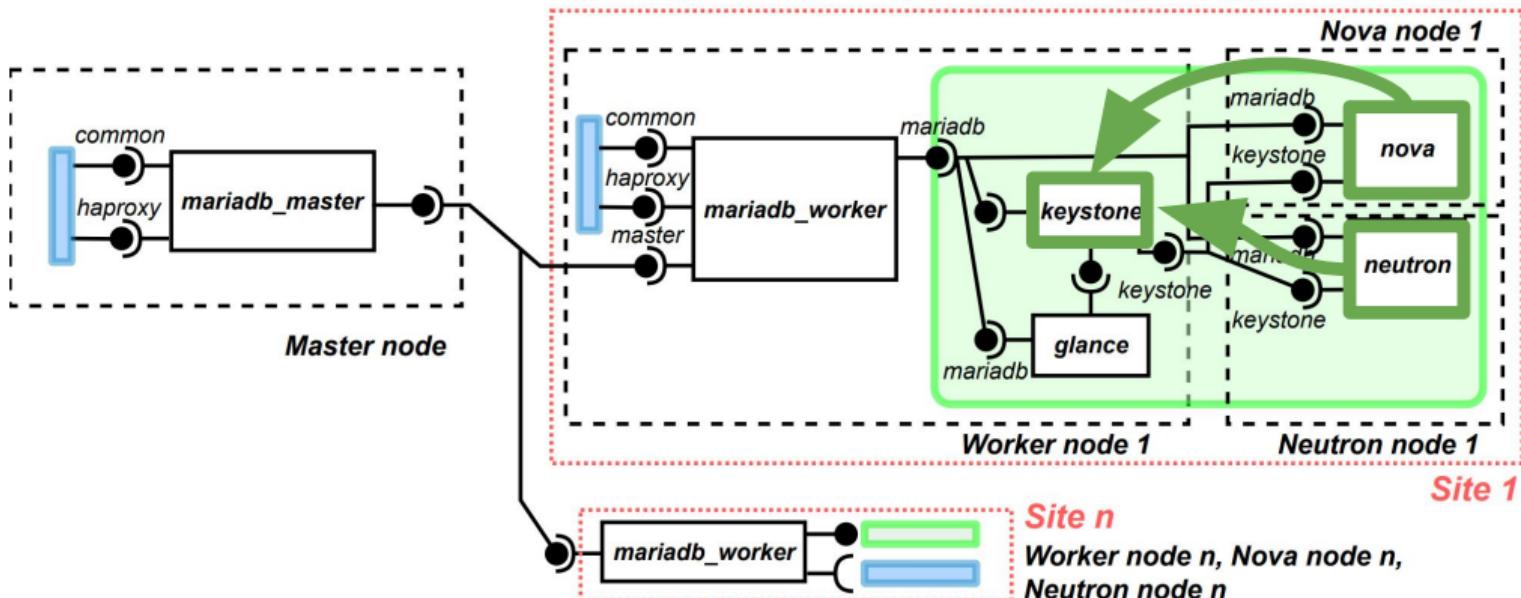
# Communication protocol



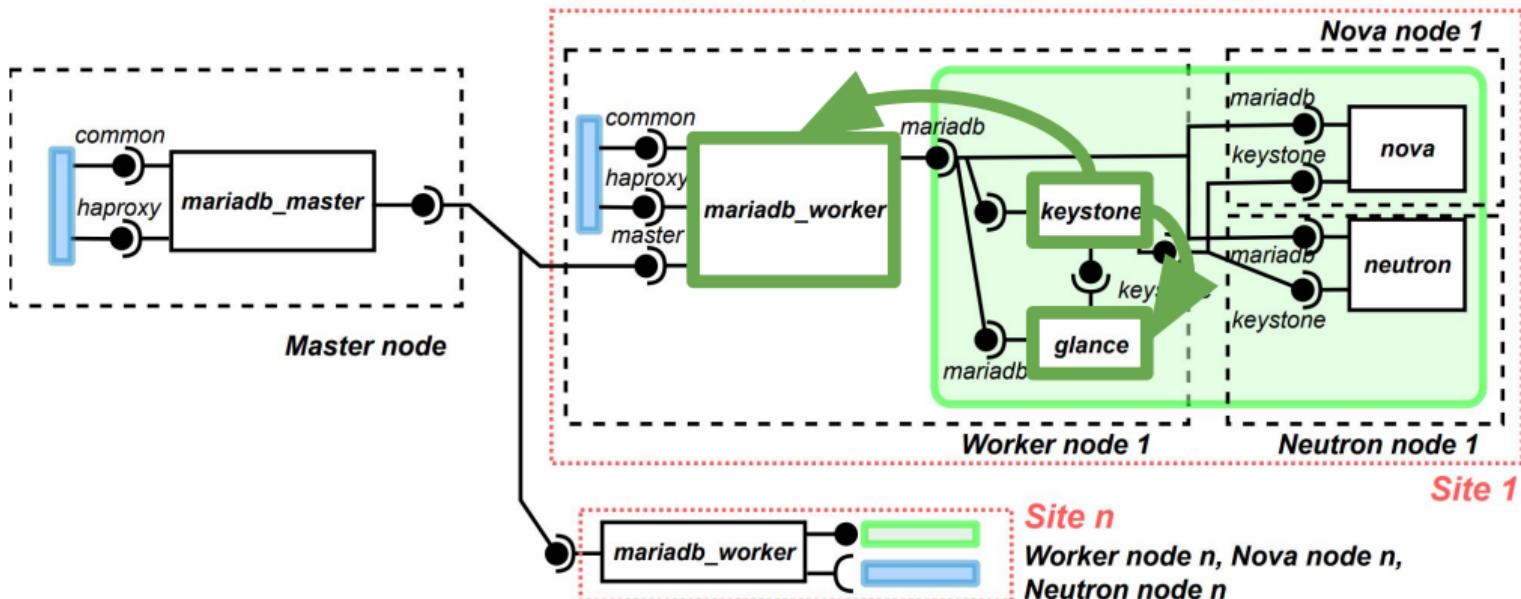
# Communication protocol



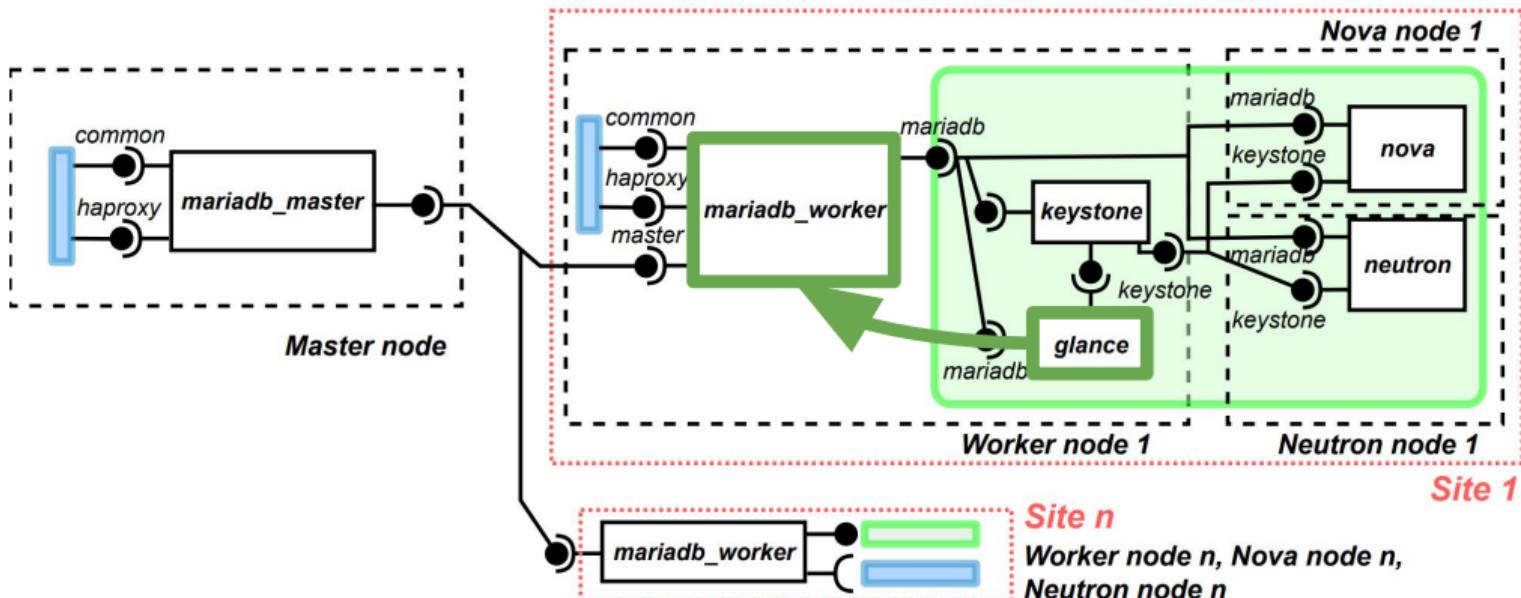
# Communication protocol



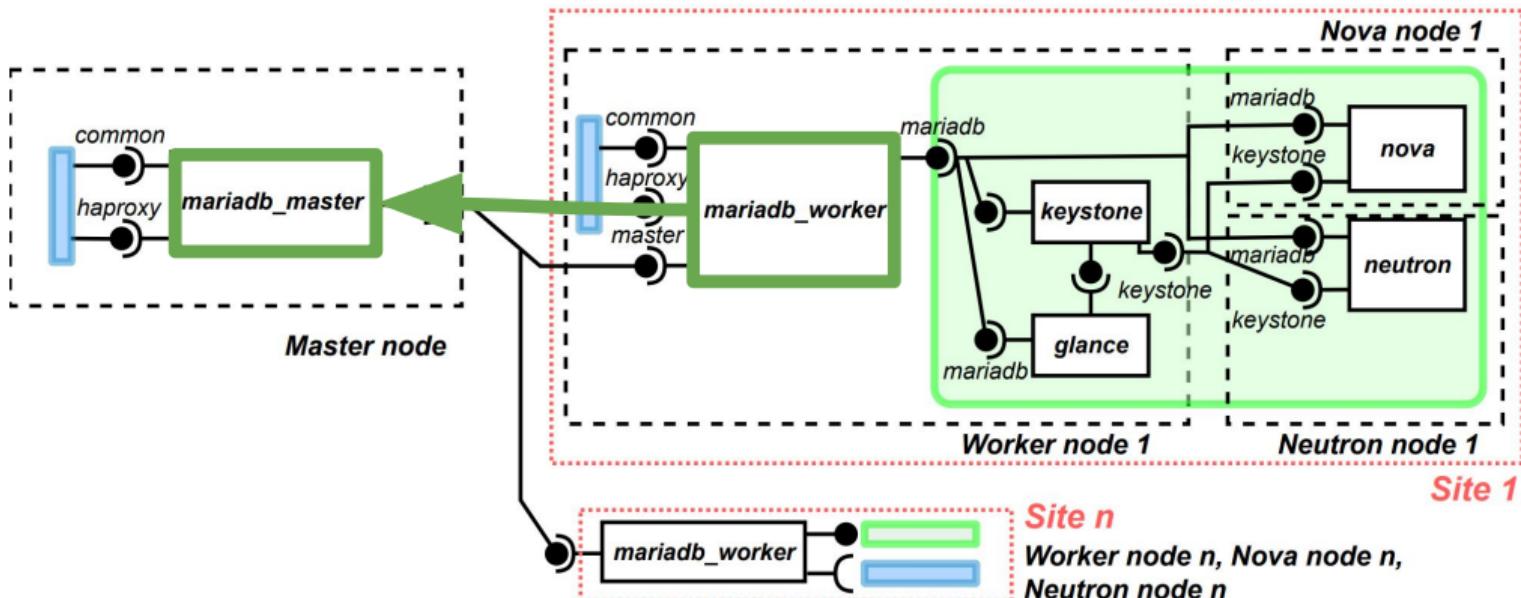
# Communication protocol



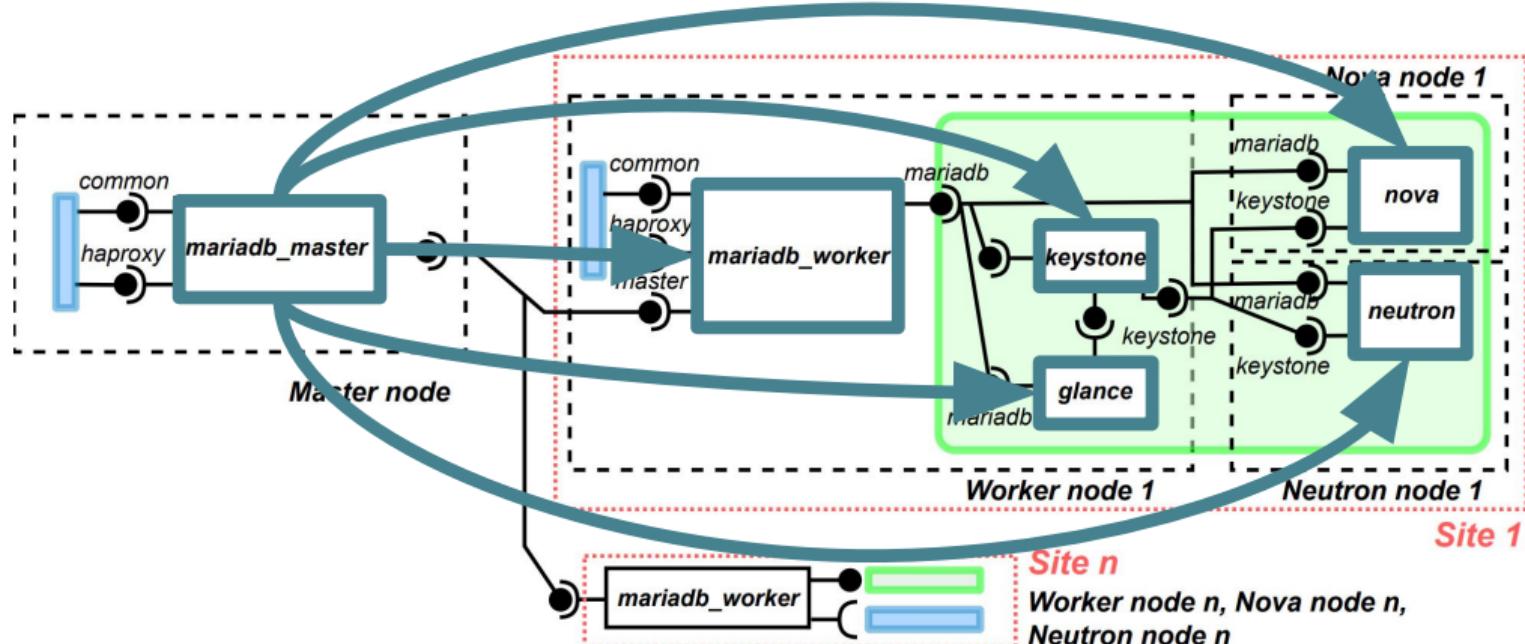
# Communication protocol



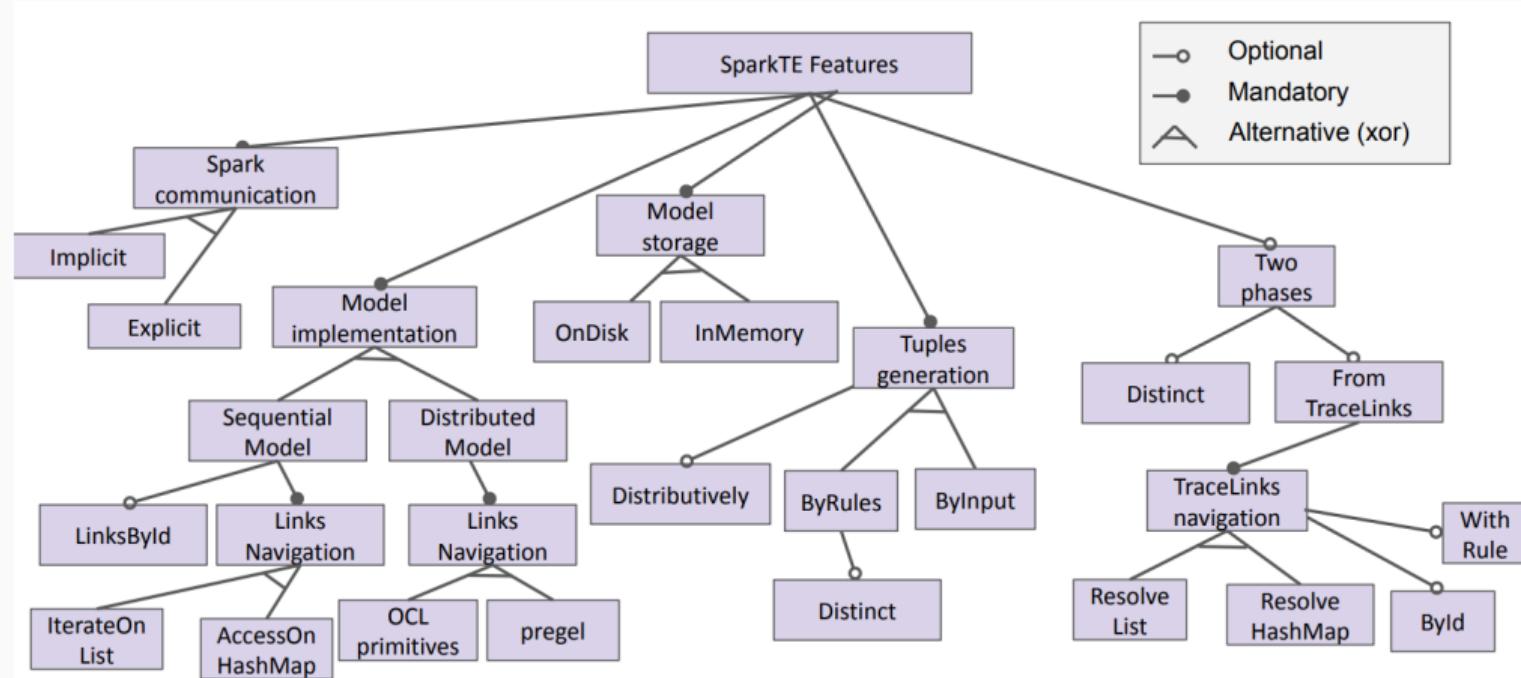
# Communication protocol



## Communication protocol

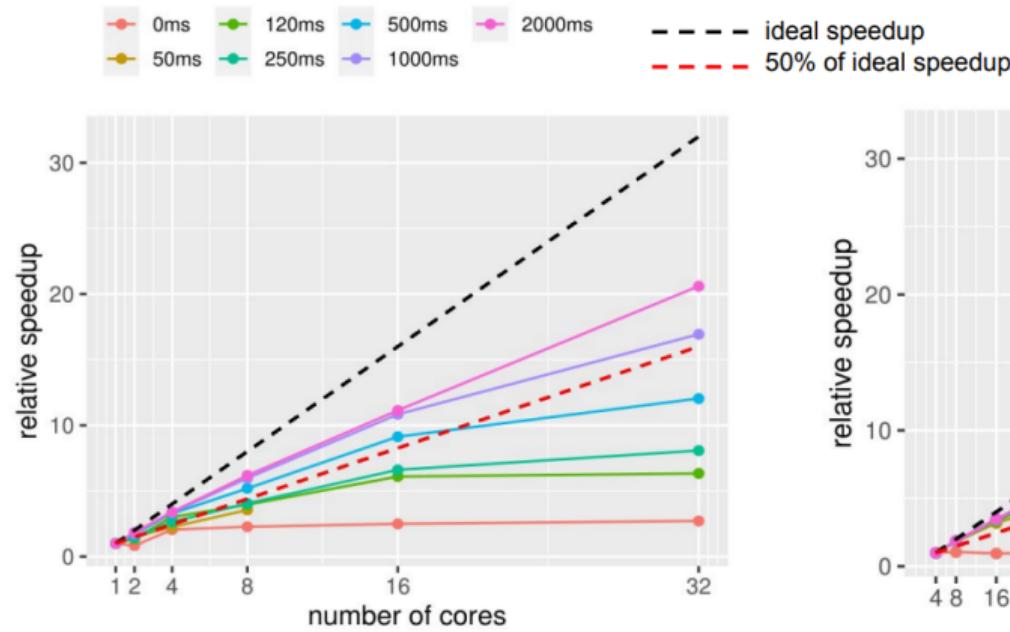


# SparkTE - Configuration space overview

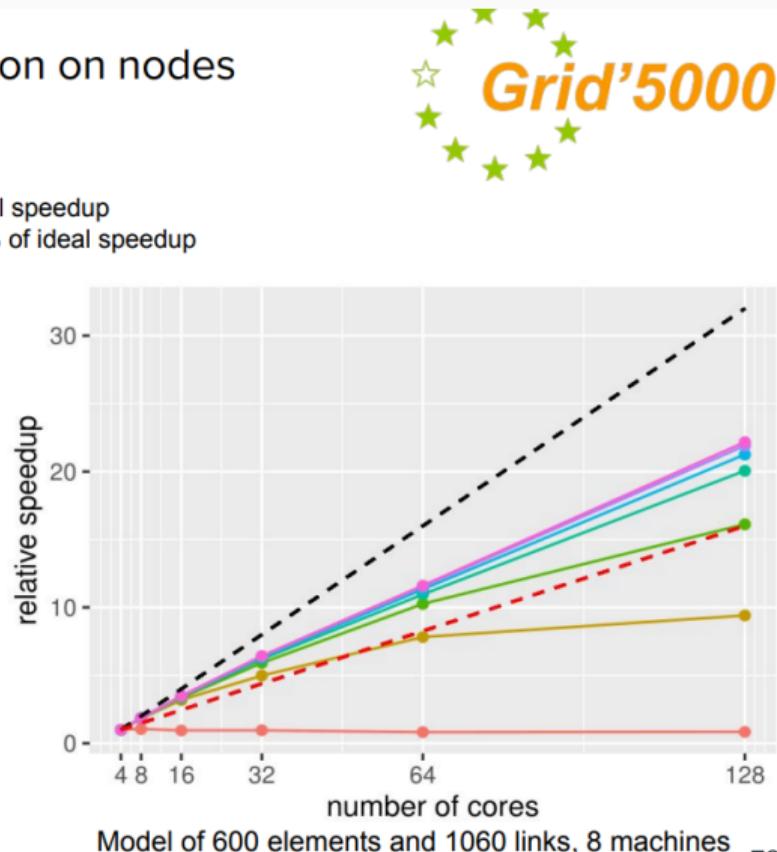


# SparkTE Performances

- Simulate a uniform amount of computation on nodes
  - fixed time for each task



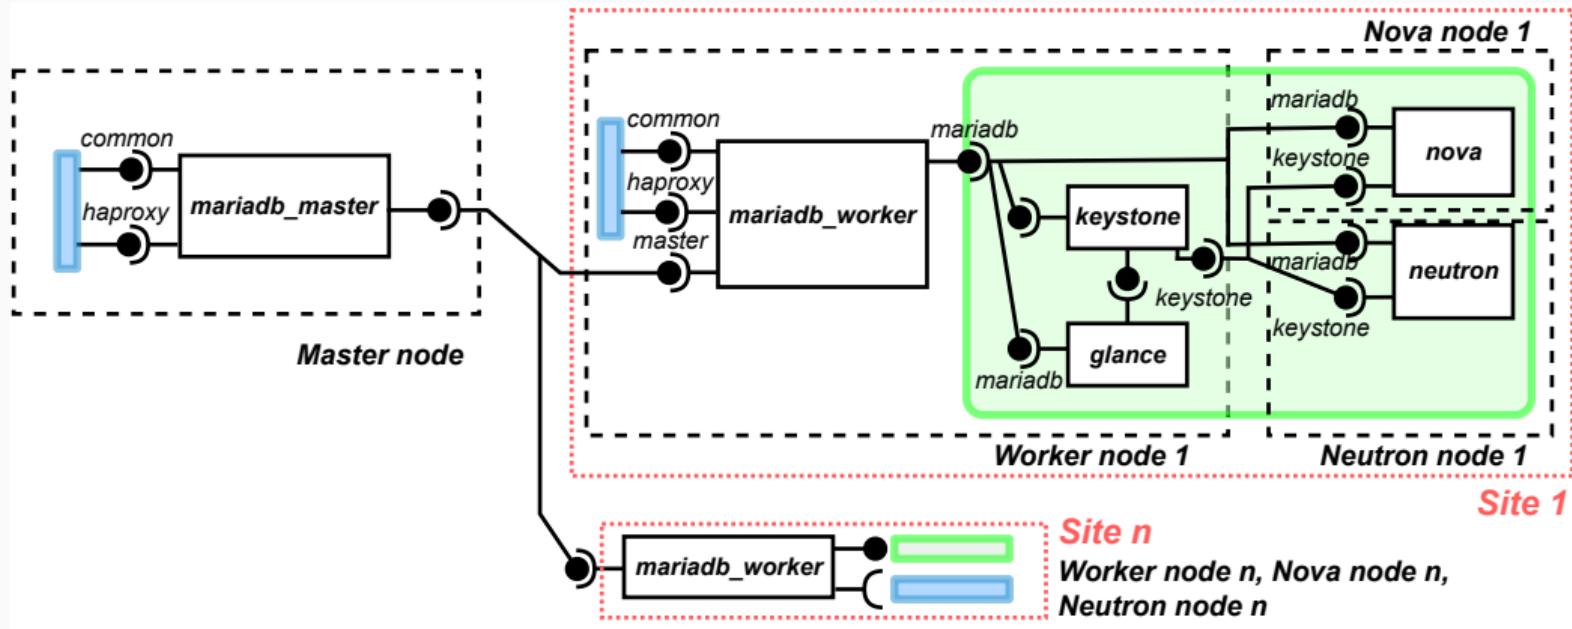
Model of 150 elements and 290 links, on 4 machines



Model of 600 elements and 1060 links, 8 machines



# Ballet performances on real use-case



## Ballet performances on real use-case

Sc.	# Sites	Planning	Ballet Execution	Total	Muse	Gain
Deploy	1	1.69s	306.02s	307.71s	536.57s	42.7%
	2	1.78s	306.09s	307.86s	536.69s	42.6%
	5	1.77s	306.19s	307.97s	537.09s	42.7%
	10	2.02s	306.14s	308.19s	538.13s	42.7%
Update	1	3.36s	416.84s	420.20s	555.56s	24.4%
	2	4.39s	416.92s	421.31s	555.70s	24.2%
	5	6.05s	417.17s	423.22s	556.08s	24.0%
	10	5.97s	417.46s	423.43s	556.77s	24.0%

**Table 2:** Comparison of time for planning and executing a deployment and an update of the MariaDB\_master instance with Ballet and Muse.

# CP Model

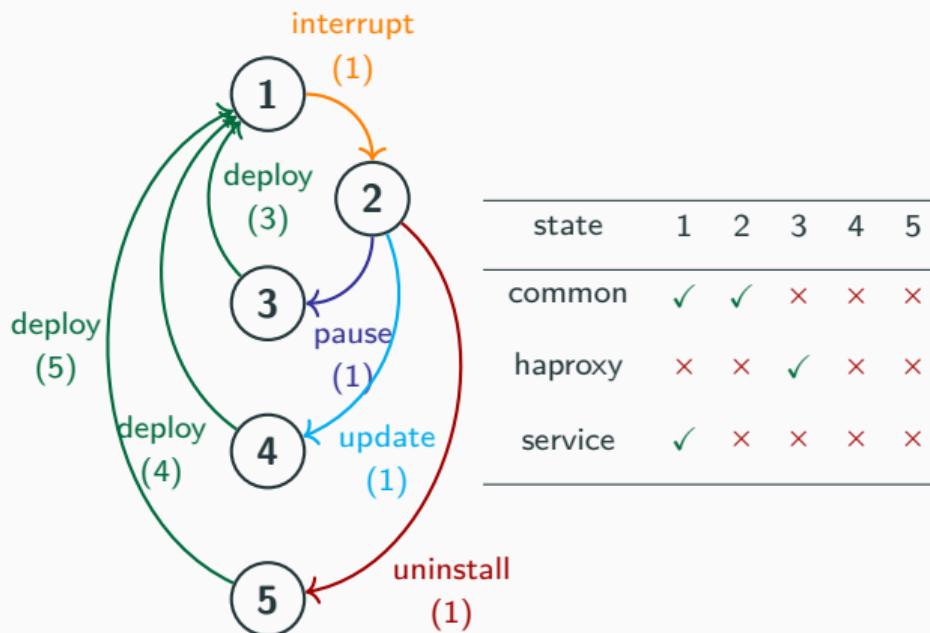
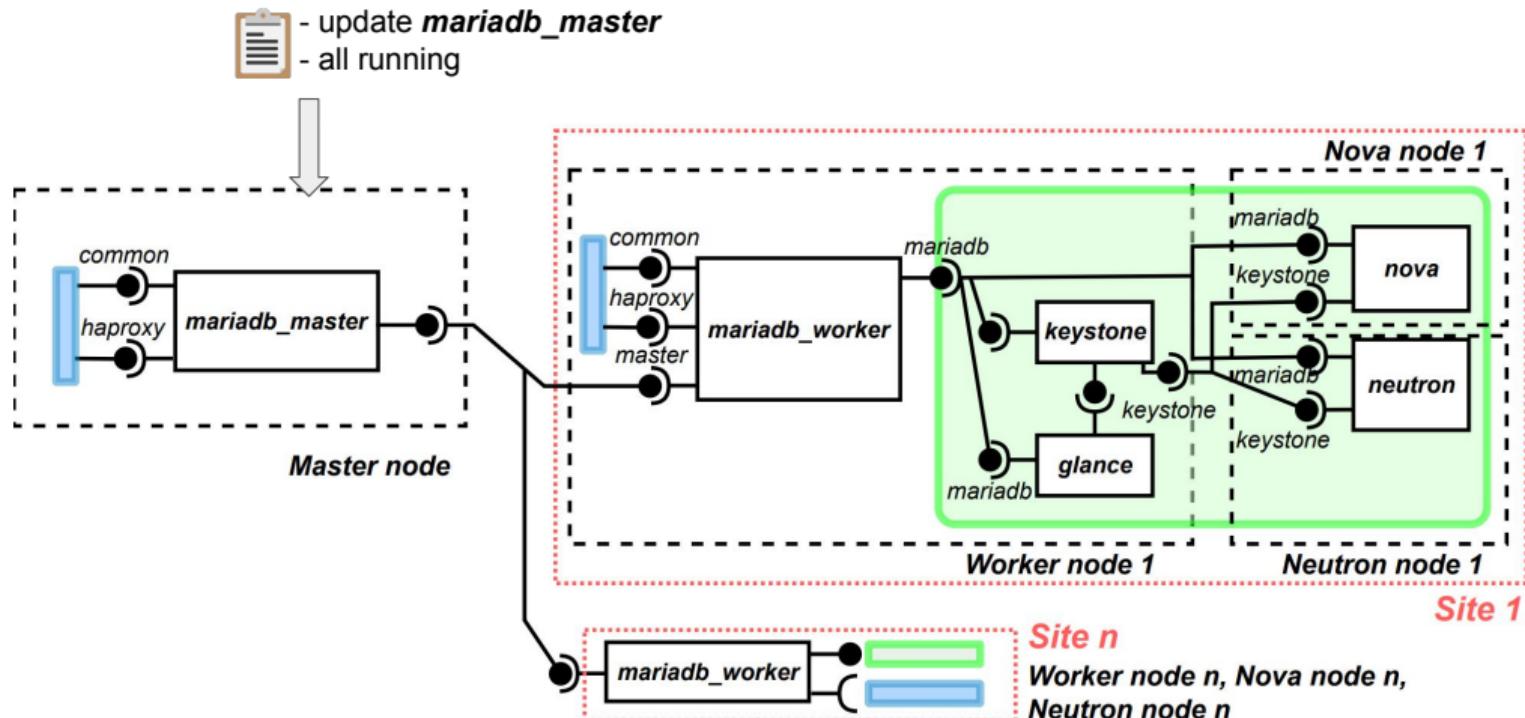


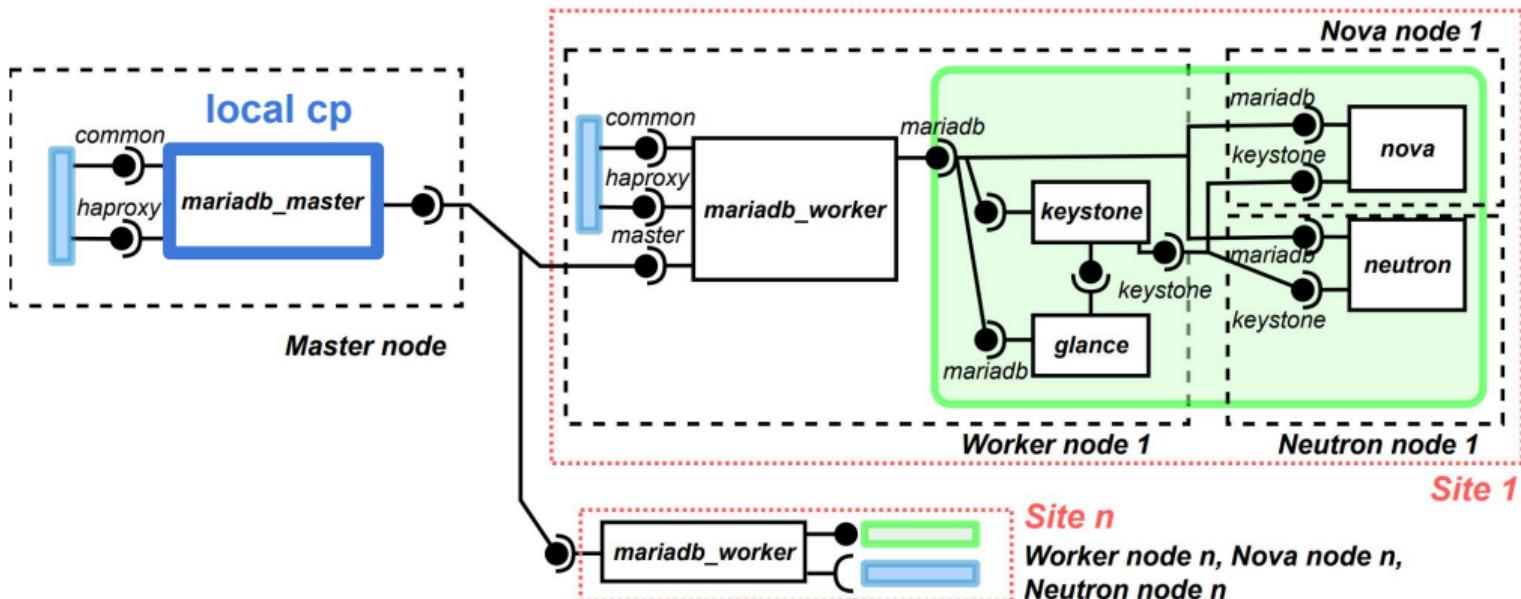
Figure 2: Automaton representation of *Mariadb\_master* component's life cycle with its matrix for ports statuses.

- $(B, \Pi, \mathcal{C}, s_{init}, S_{goal})$
  - $s_{i+1} = inc_{\Pi}[s_i][b_i], \forall i \in 1..m$
  - $(b, B, >, 0)$
  - $status(p, s_{m+1}) = \Gamma_p$
- where**
- $\Pi$  an automaton with  $\mathcal{C}$  costs
  - $B$  a sequence of  $m$  behaviors
  - $\Gamma_p \in \{\text{active, inactive}\}$  i.e. { ✓, ✗ }
  - $b \in \{ \text{interrupt, deploy, pause, update, uninstall} \}$

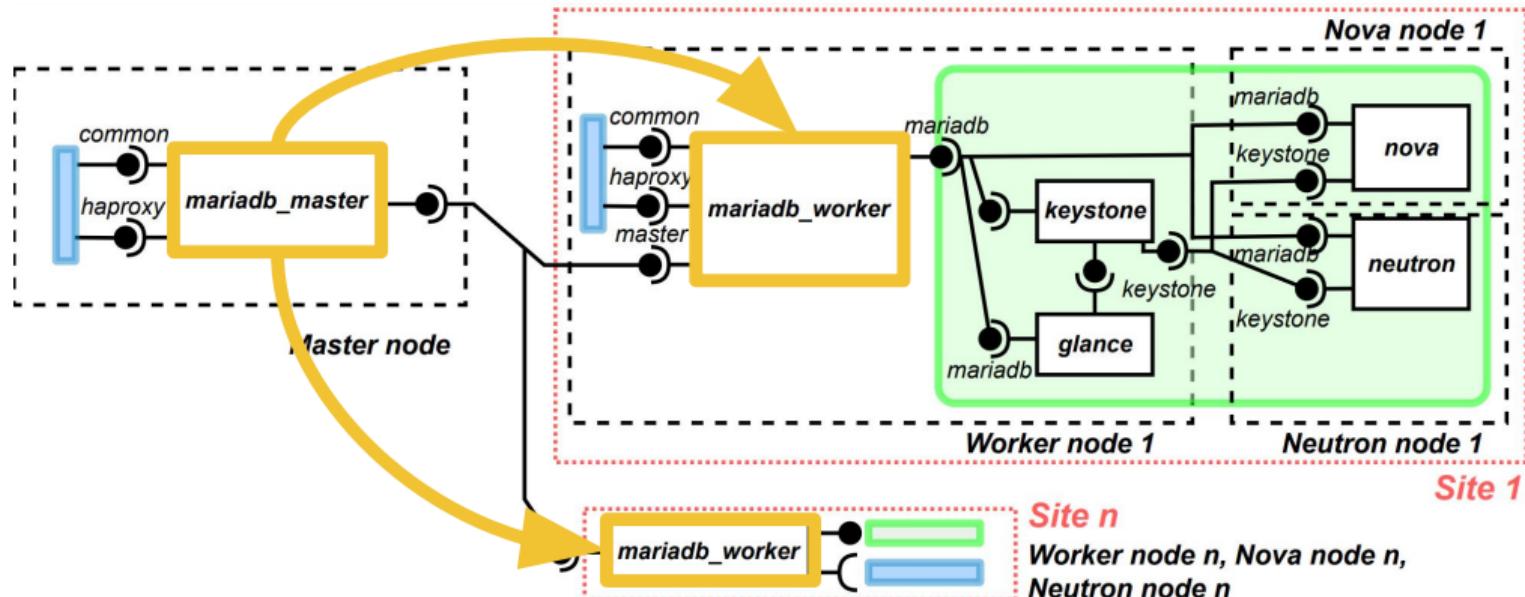
# Communication protocol



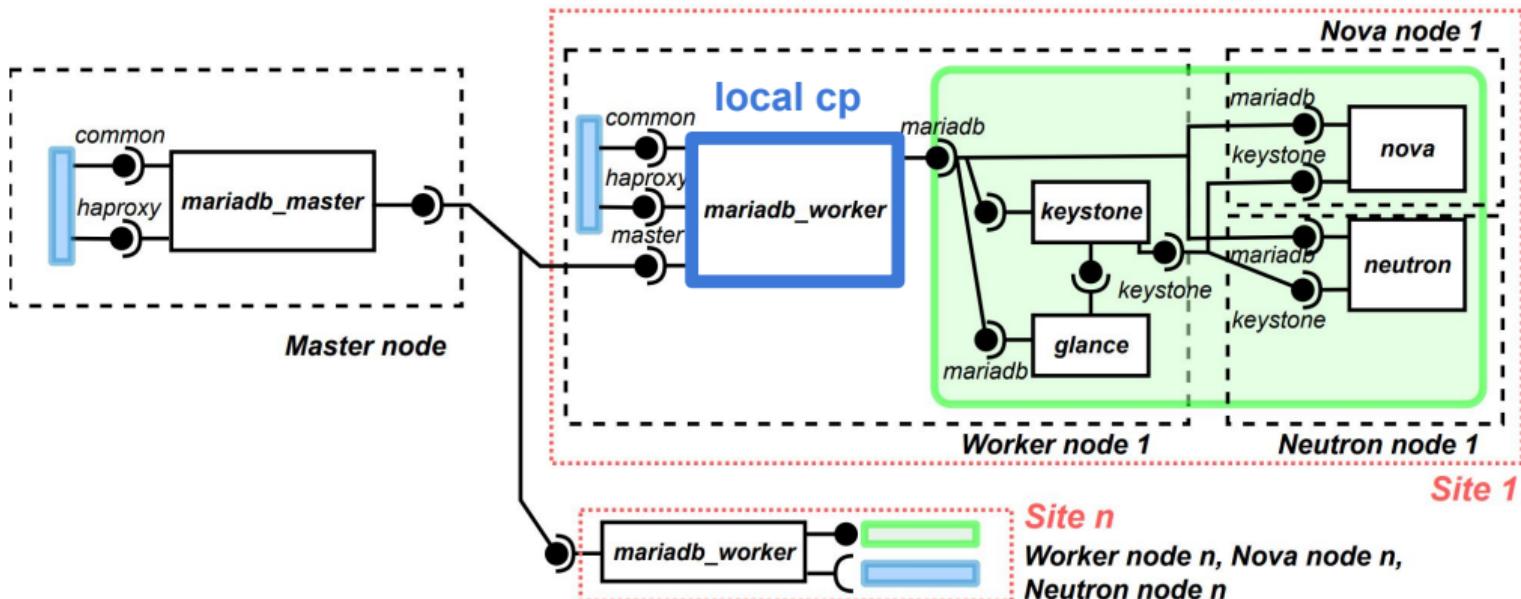
# Communication protocol



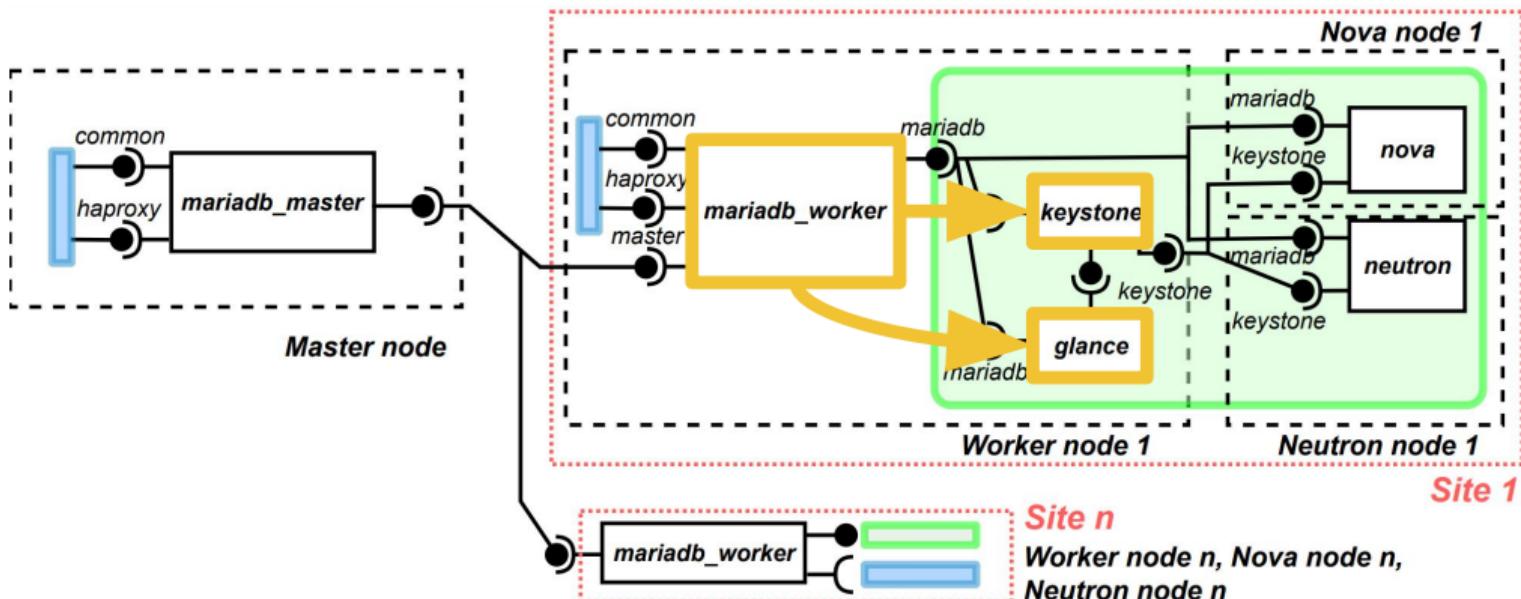
# Communication protocol



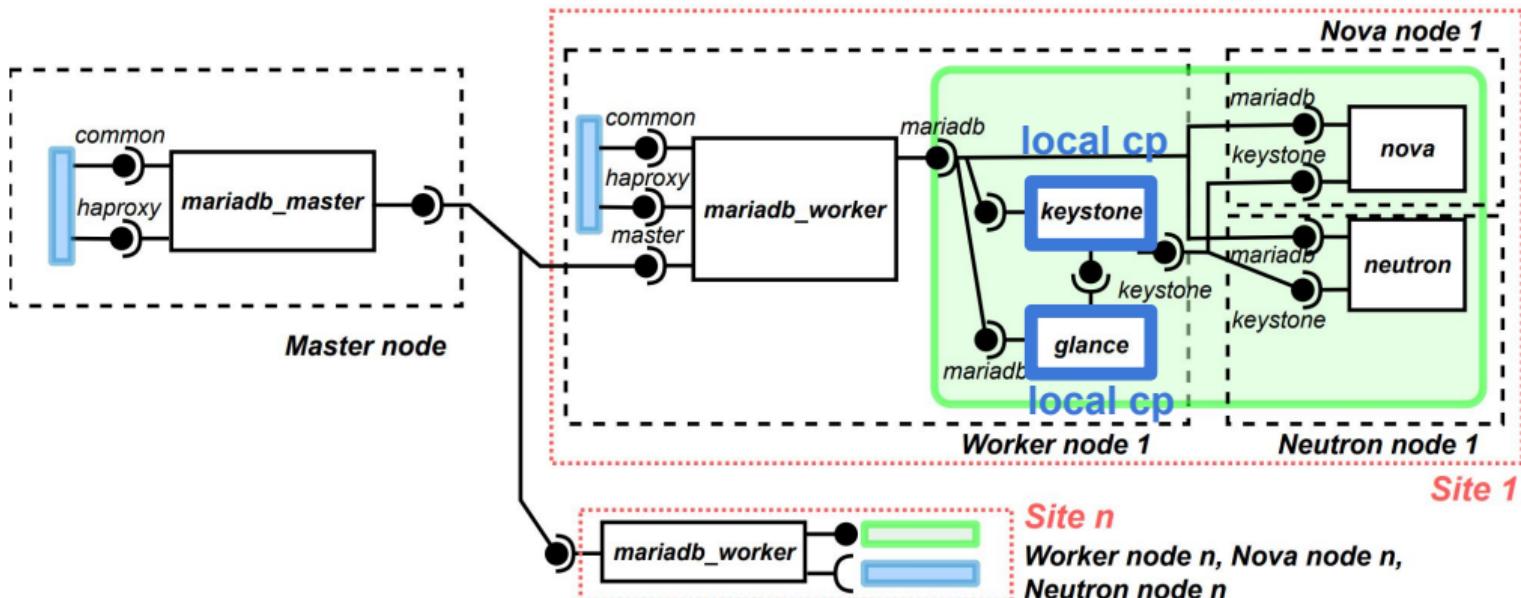
# Communication protocol



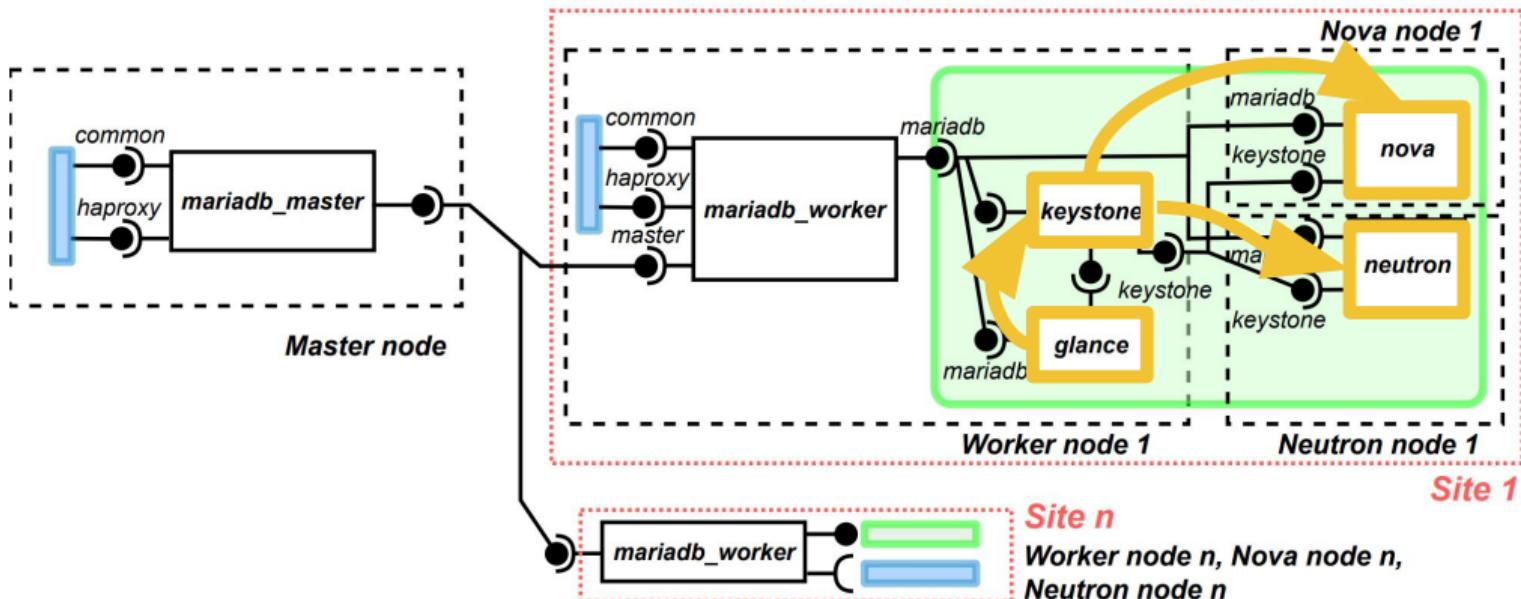
# Communication protocol



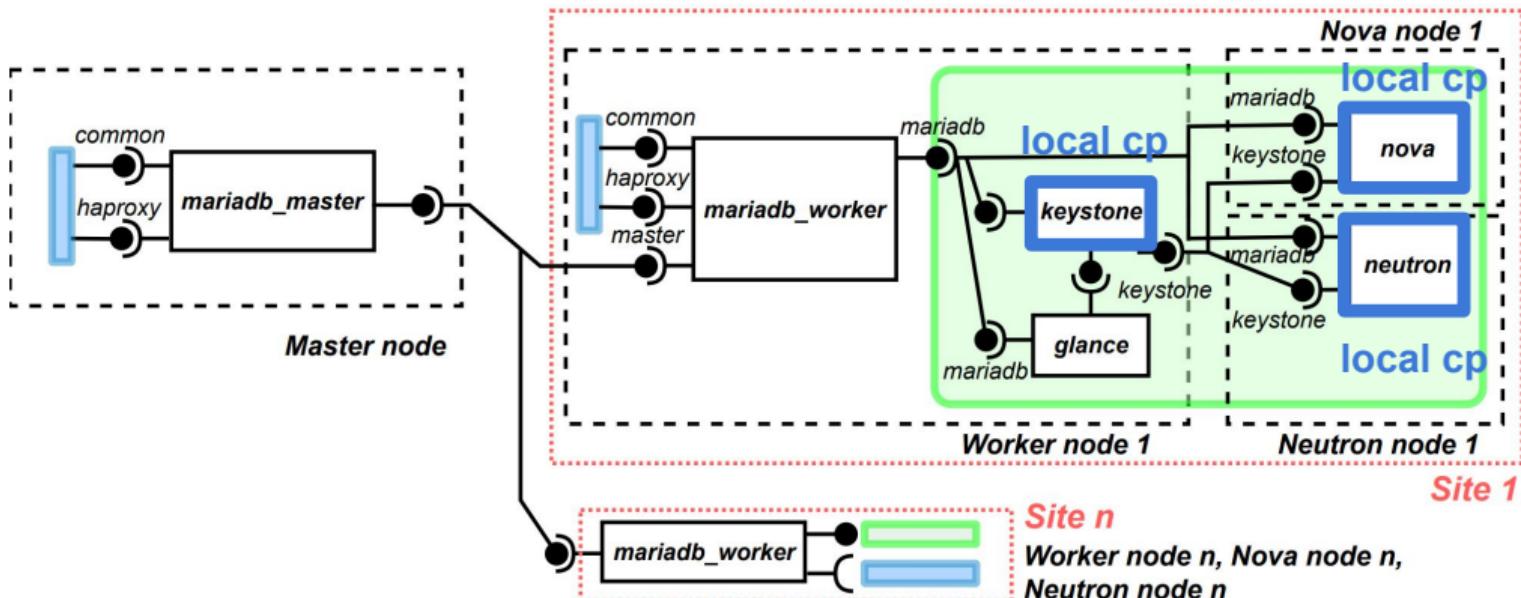
# Communication protocol



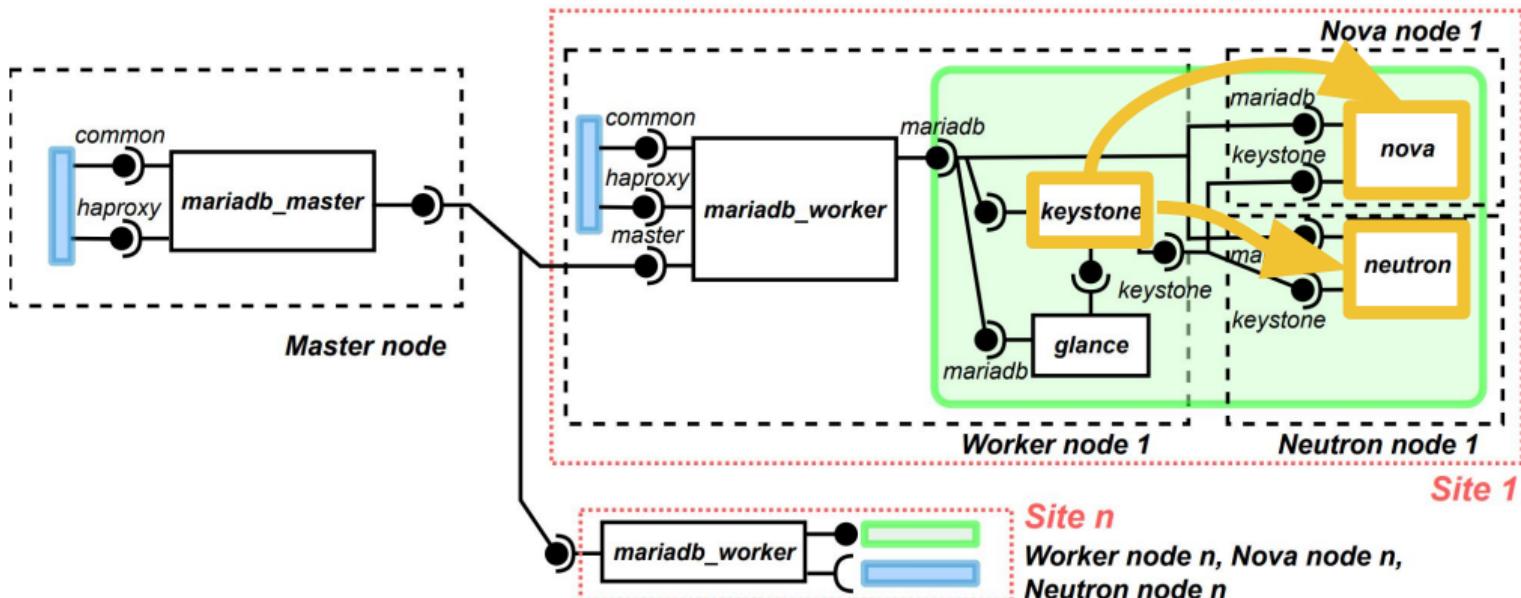
# Communication protocol



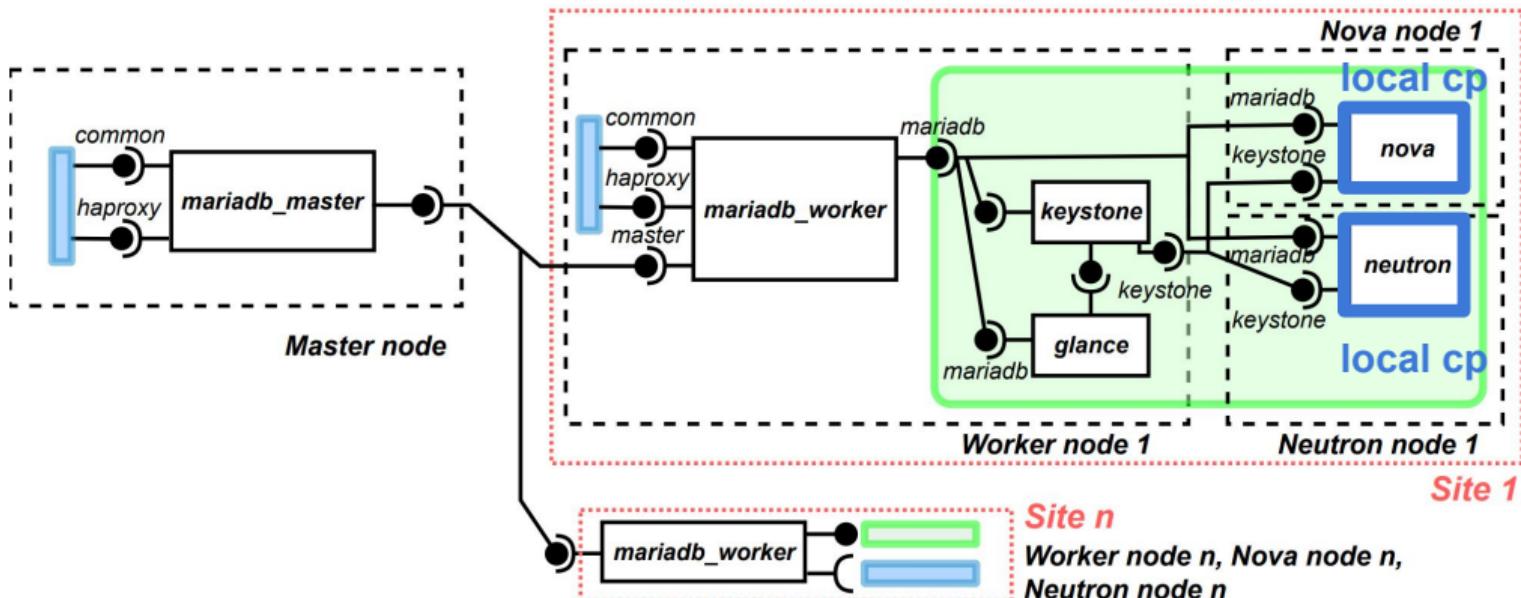
# Communication protocol



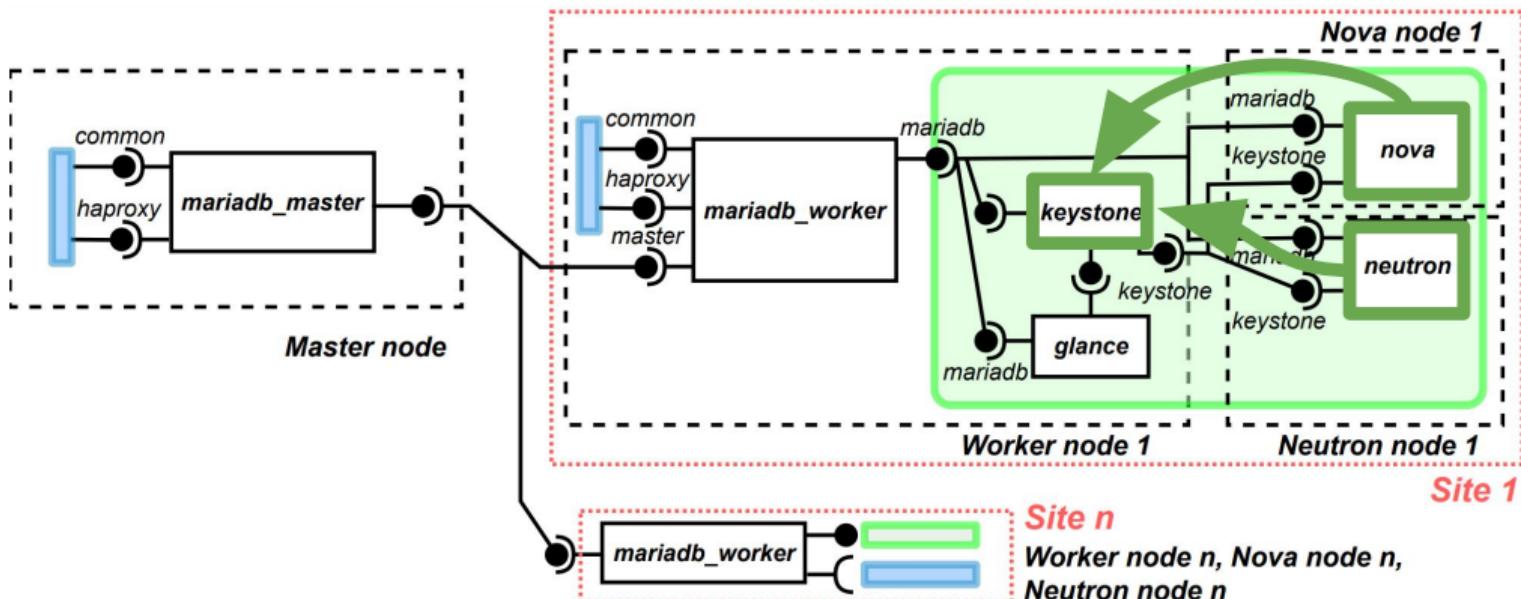
# Communication protocol



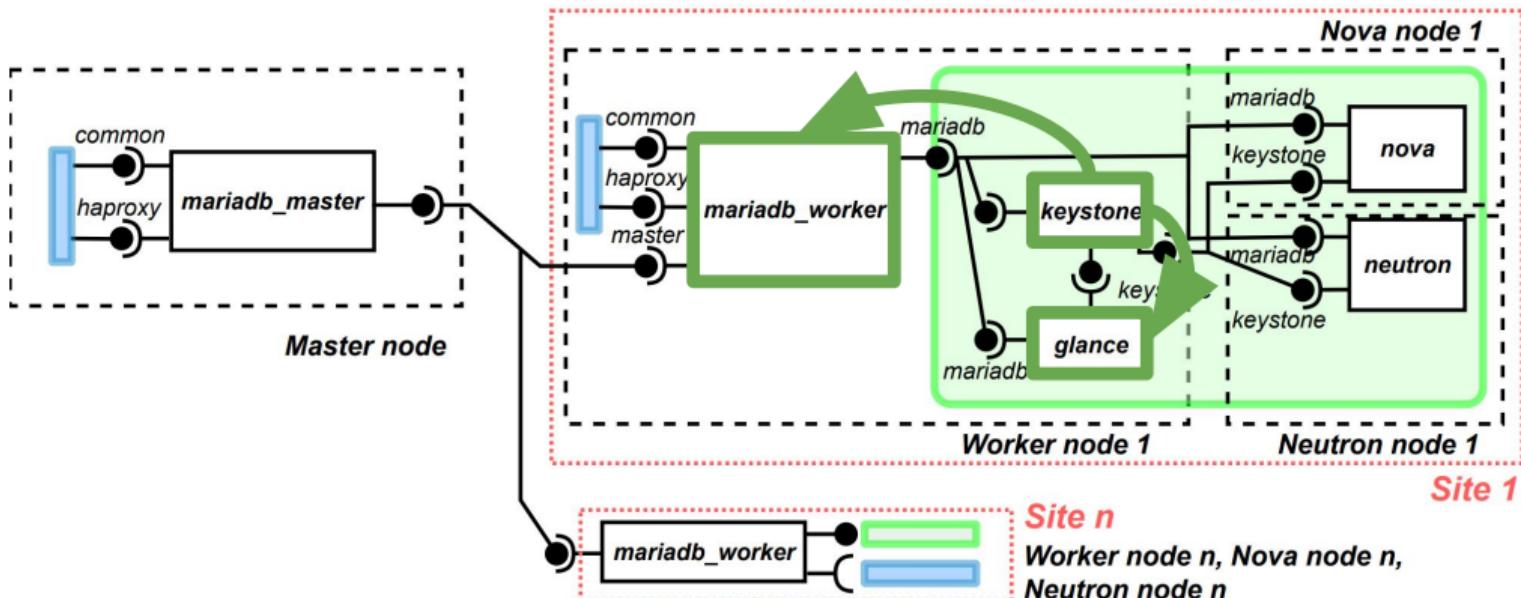
# Communication protocol



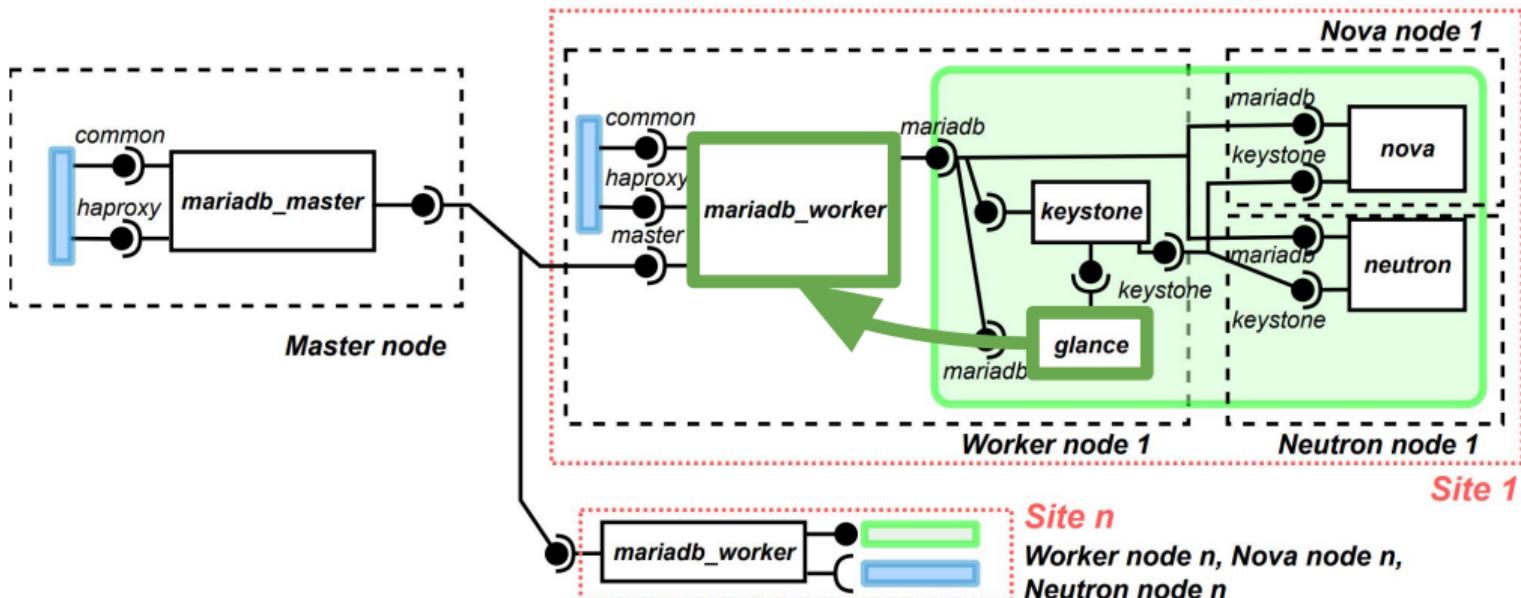
# Communication protocol



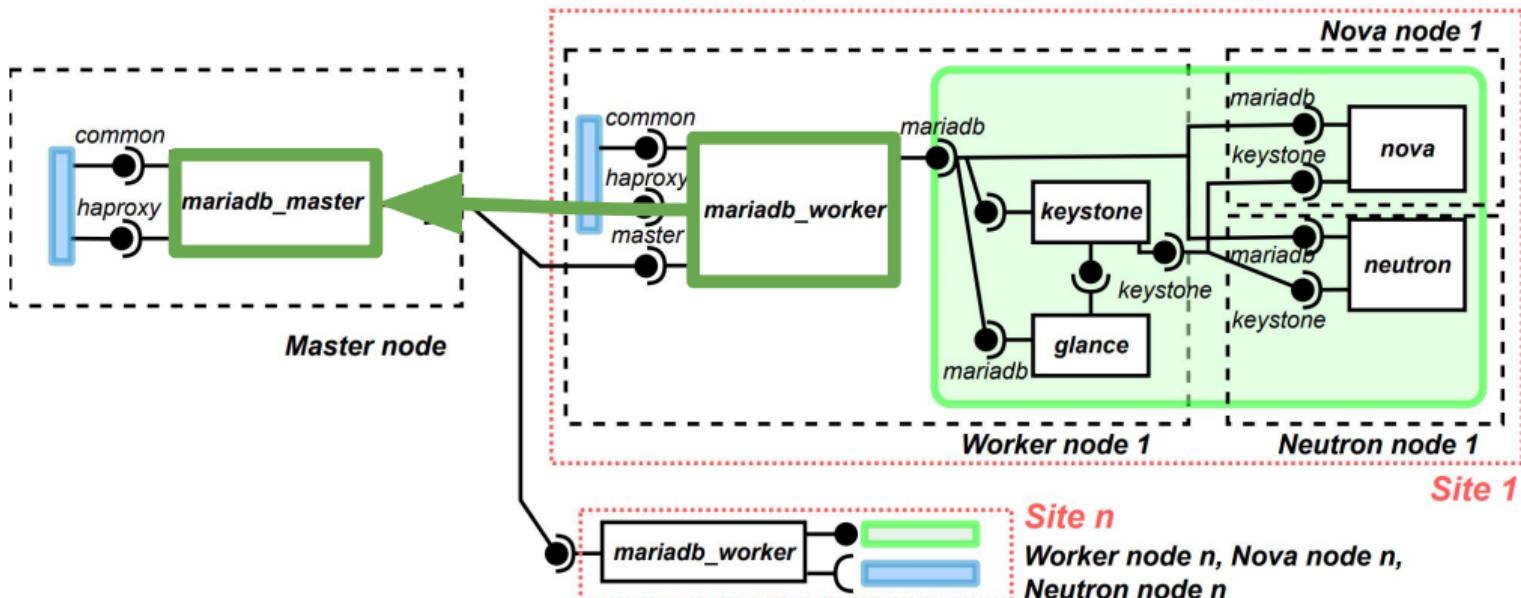
# Communication protocol



# Communication protocol



# Communication protocol



# Communication protocol

